



AUGUST 4-5, 2021

BRIEFINGS

Alcatraz:

A Practical Hypervisor Sandbox to Prevent Escapes
from the KVM/QEMU and KVM-Based MicroVMs

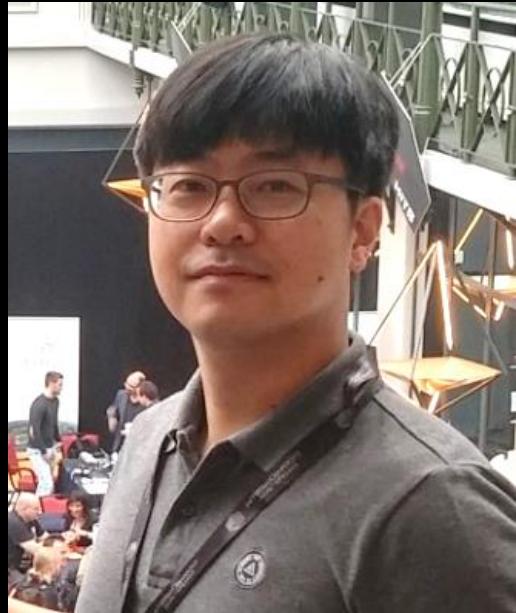
Seunghun Han

hanseunghun@nsr.re.kr

Sungjung Kim, Gaksoo Lim, Byungjoon Kim
[\(sungjungk||daygax||bjkim\)@nsr.re.kr](mailto:(sungjungk||daygax||bjkim)@nsr.re.kr)

A terminal window with a black background and white text. The text consists of various binary and ASCII characters, including brackets, slashes, and symbols like '^' and '#'. The window has rounded corners and a thin black border.

Who Am I?



- **Senior security researcher** at the Affiliated Institute of ETRI
- **Review board member** of **Black Hat Asia** and **KimchiCon**
- **Speaker** at **USENIX Security**, **Black Hat Asia/Europe**,
HITBSecConf, **BlueHat Shanghai**, **TyphoonCon**, **KimmchiCon**,
BECS, etc.
- **Author** of “64-bit multi-core OS principles and structure, Vol.1&2”
- **Debian Linux maintainer** and **Linux kernel contributor**
- a.k.a kkamogui,  [@kkamogui1](https://twitter.com/kkamogui1)



Goal of This Talk

- I present escape paths of KVM/QEMU and KVM-based microVMs
 - KVM-based hypervisors are widely used from VMs to containers
 - KVM has higher privilege than the kernel and can be used for escapes
- I introduce a new tool, Alcatraz, to prevent escapes
 - It downgrades KVM's privilege (Ring -1) to the guest hypervisor (Ring 0)
 - It makes a sandbox for the KVM and monitors system calls to prevent escapes
 - Process creation, kernel code modification, and privilege escalation

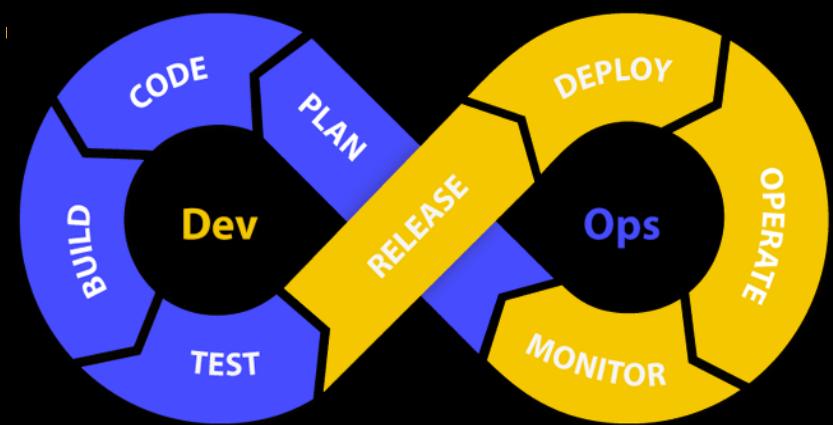


Once upon a time, there was ...

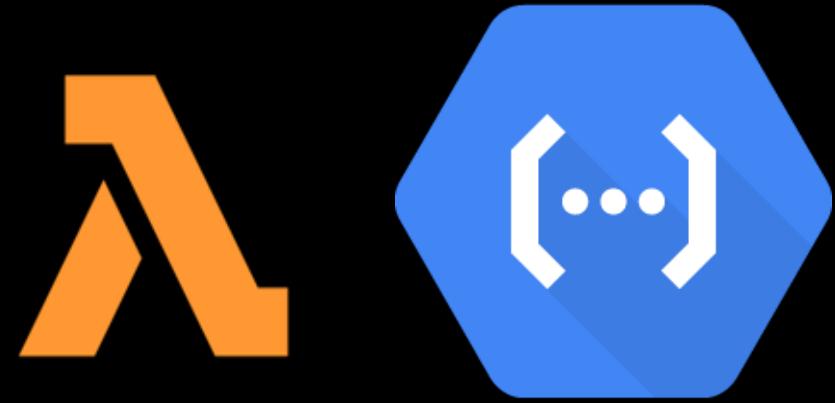


There was JUST a security researcher
(not a CEO or CTO, not even a new CSO or manager)

A New Era of Cloud Platforms has come



DevOps

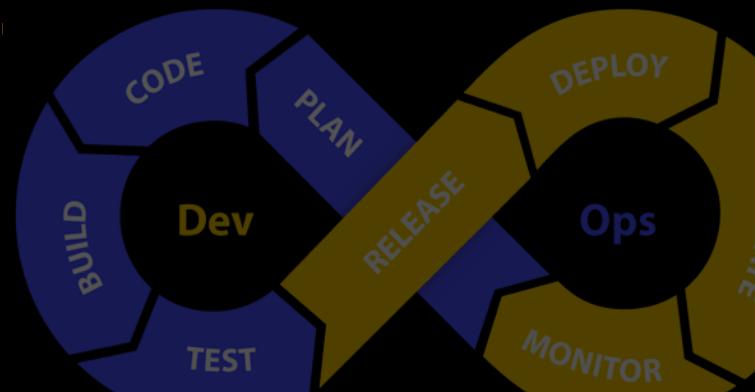


Serverless

VM-based
Technology

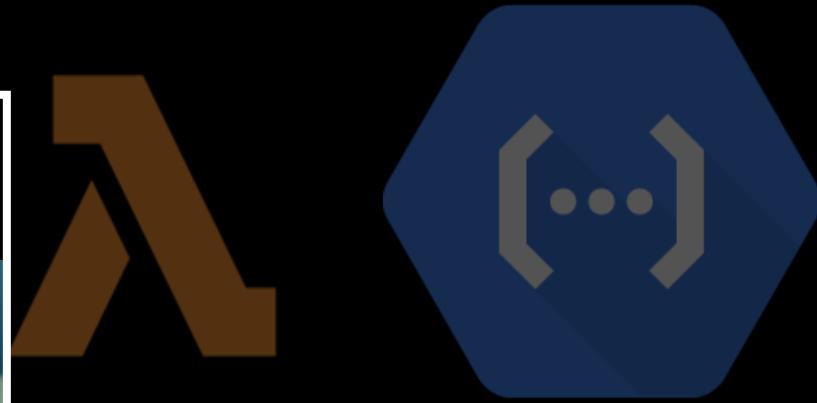
Container-based
Technology

A New Era of Cloud Platforms has come



DevOps

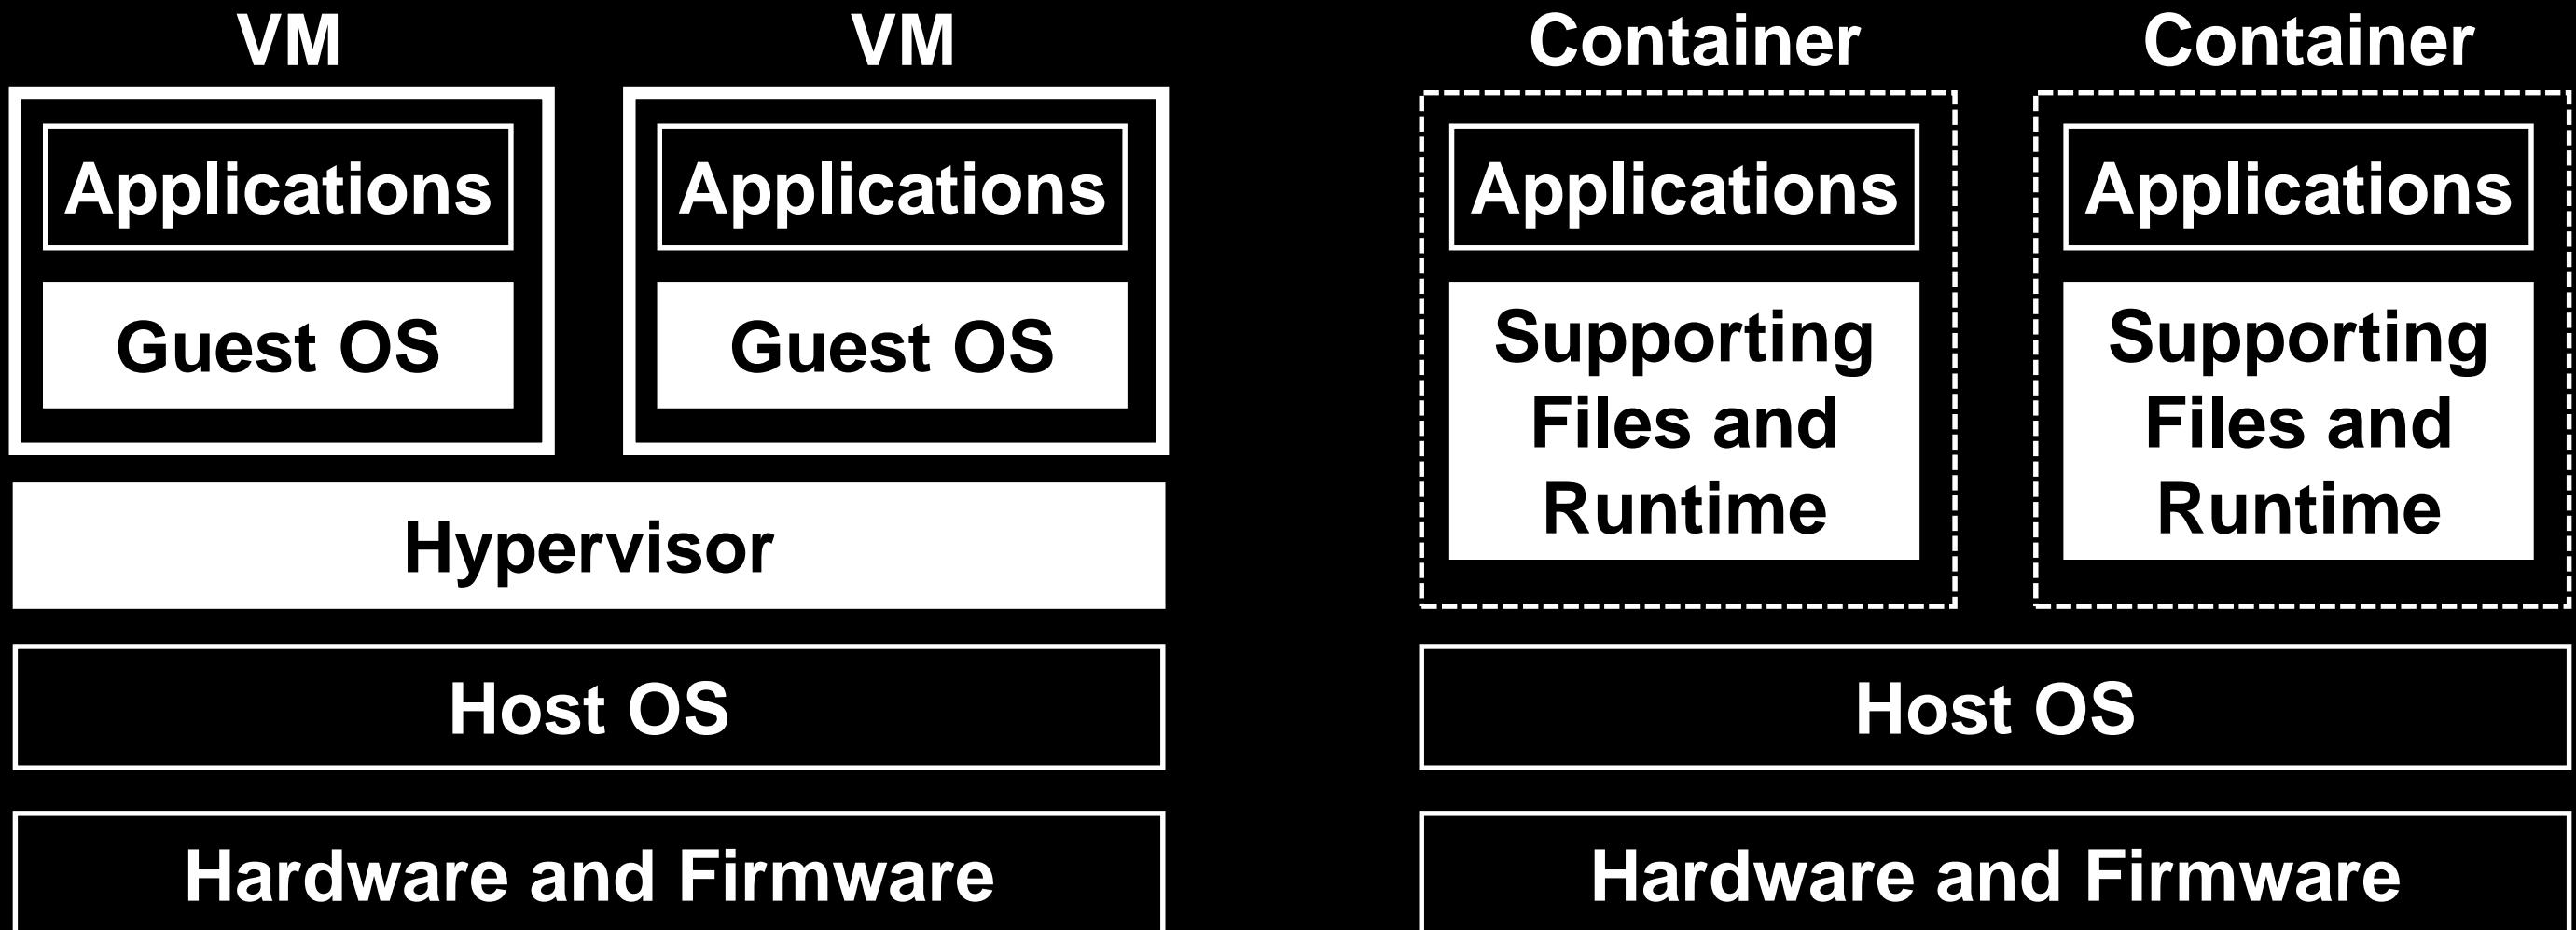
VM-based
Technology



Serverless

Container-based
Technology

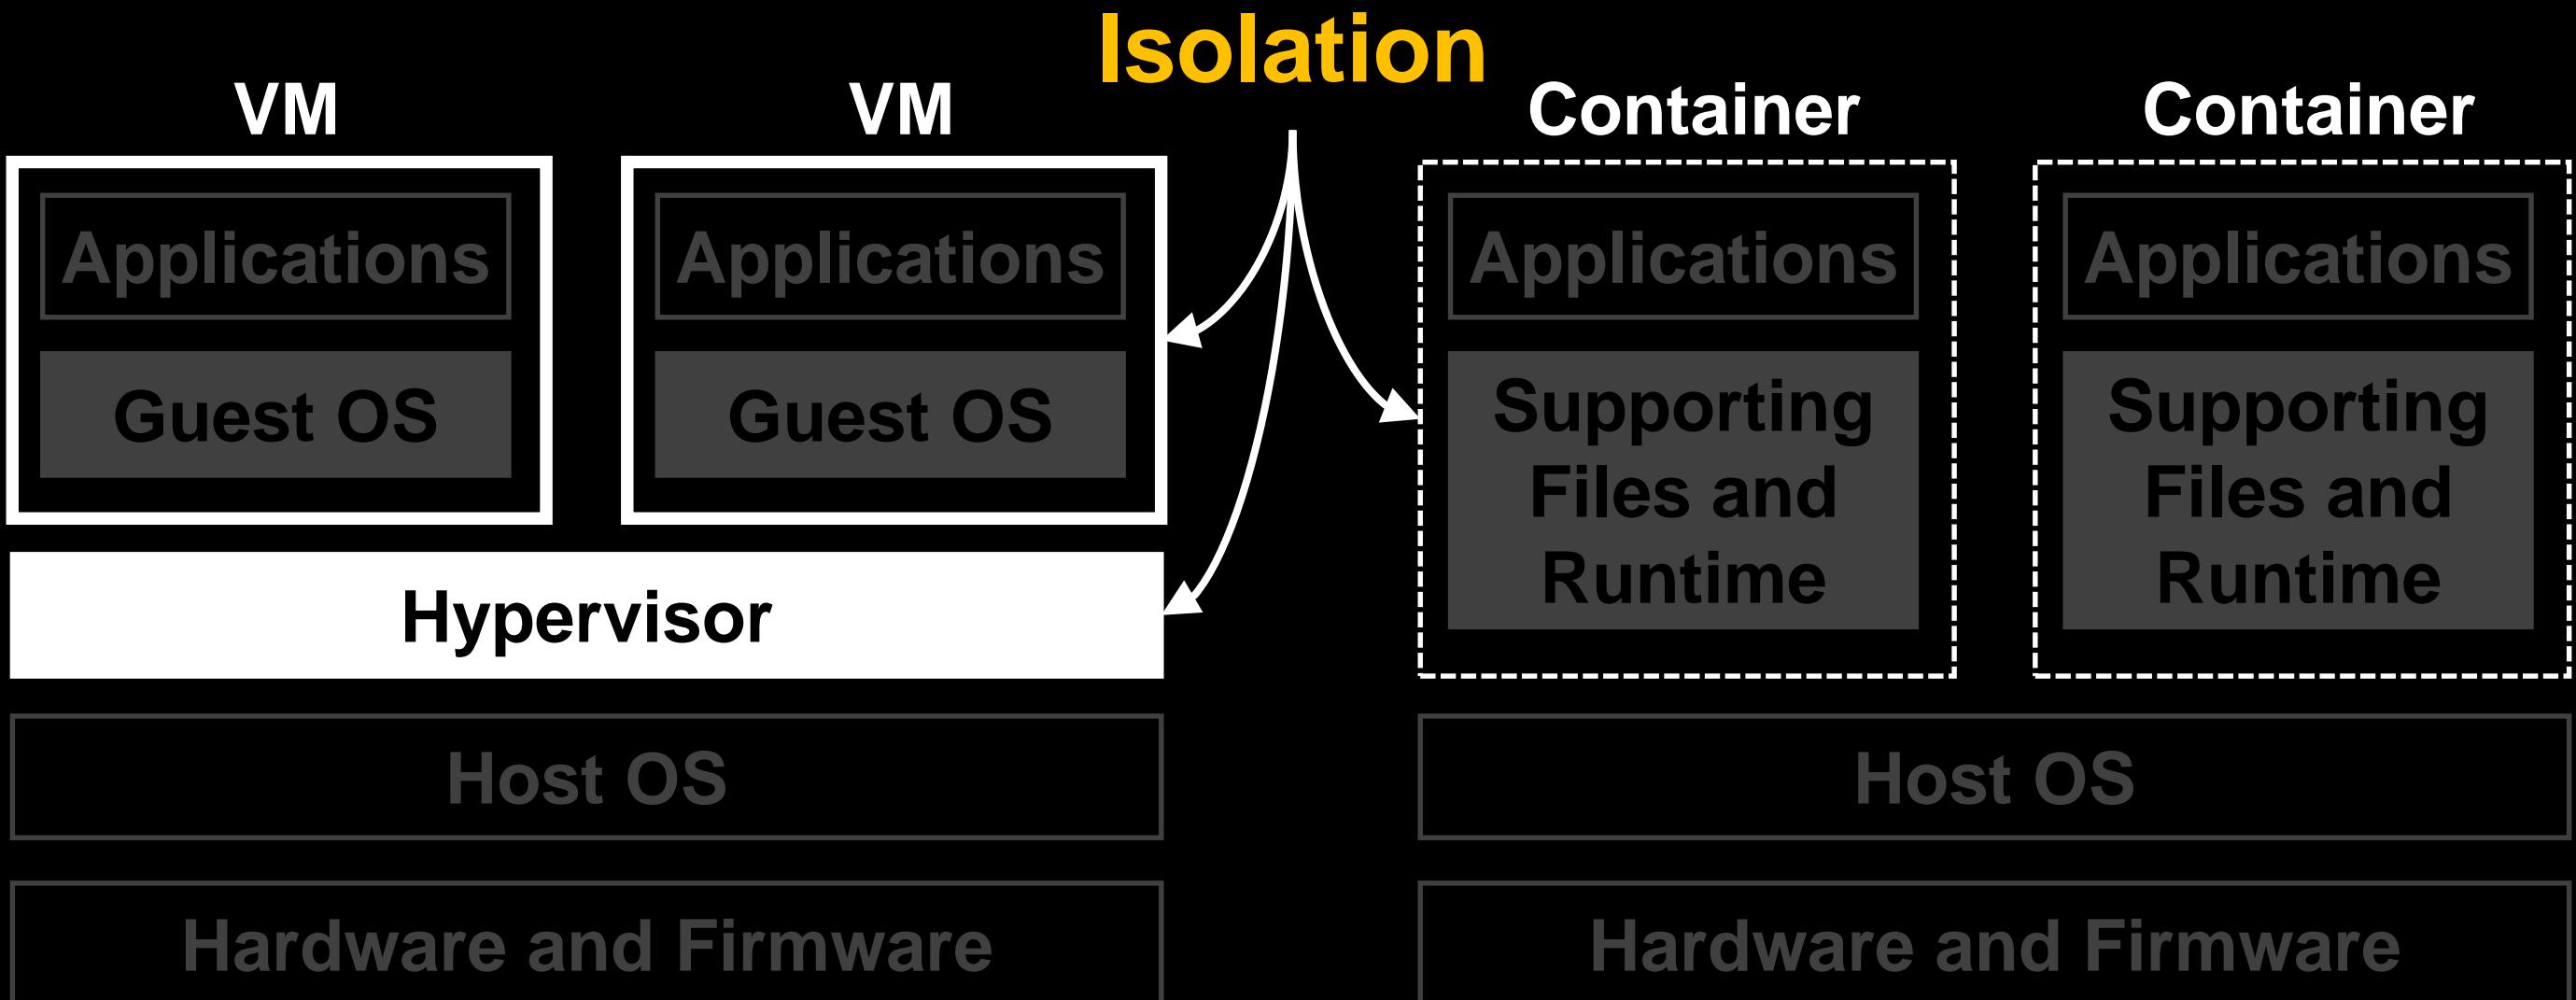
Virtual Machine (VM) vs Container (1)



Virtual Machine Architecture

Container Architecture

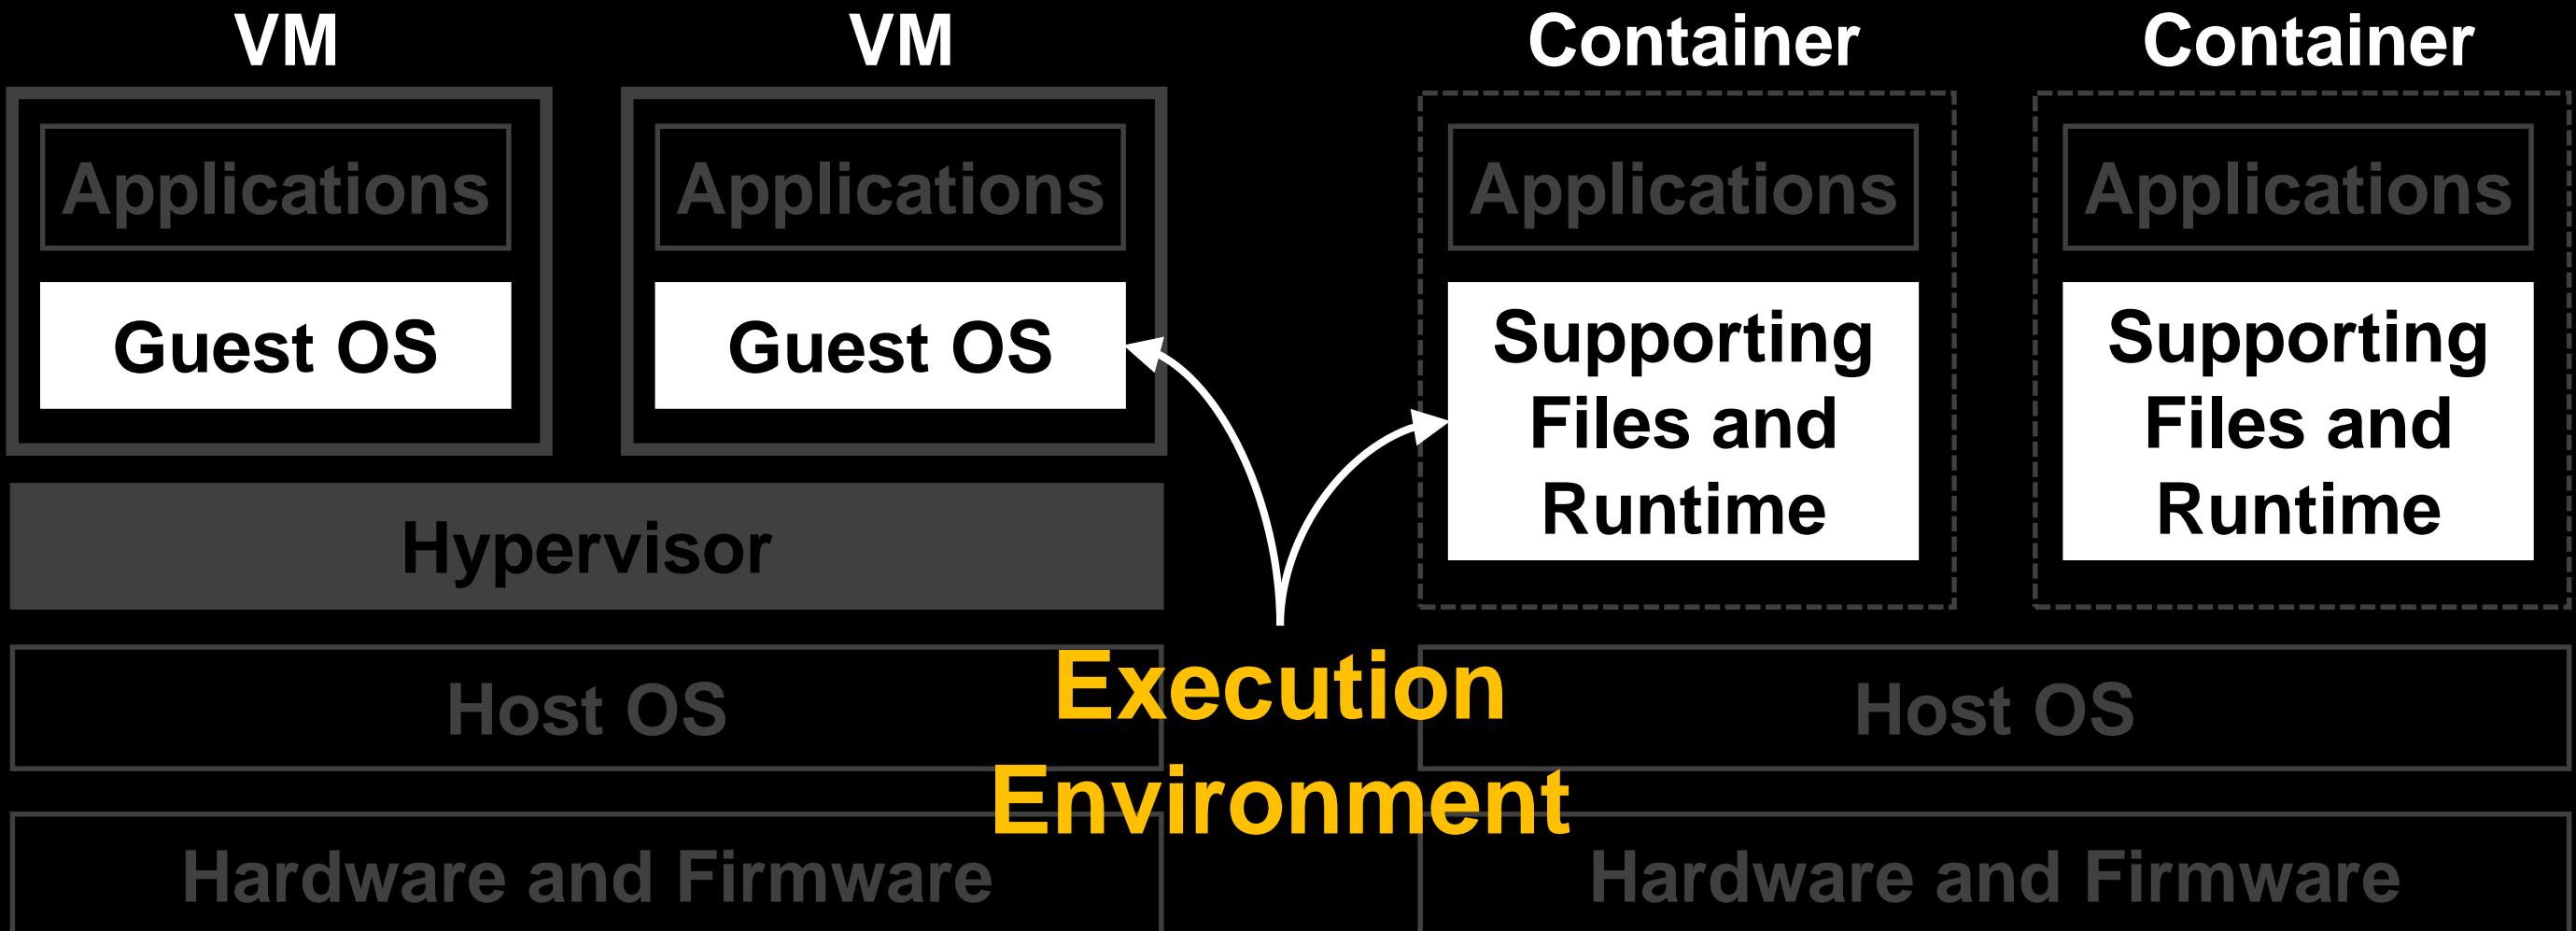
Virtual Machine (VM) vs Container (1)



Virtual Machine Architecture

Container Architecture

Virtual Machine (VM) vs Container (1)



Virtual Machine Architecture

Container Architecture

Virtual Machine (VM) vs Container (2)

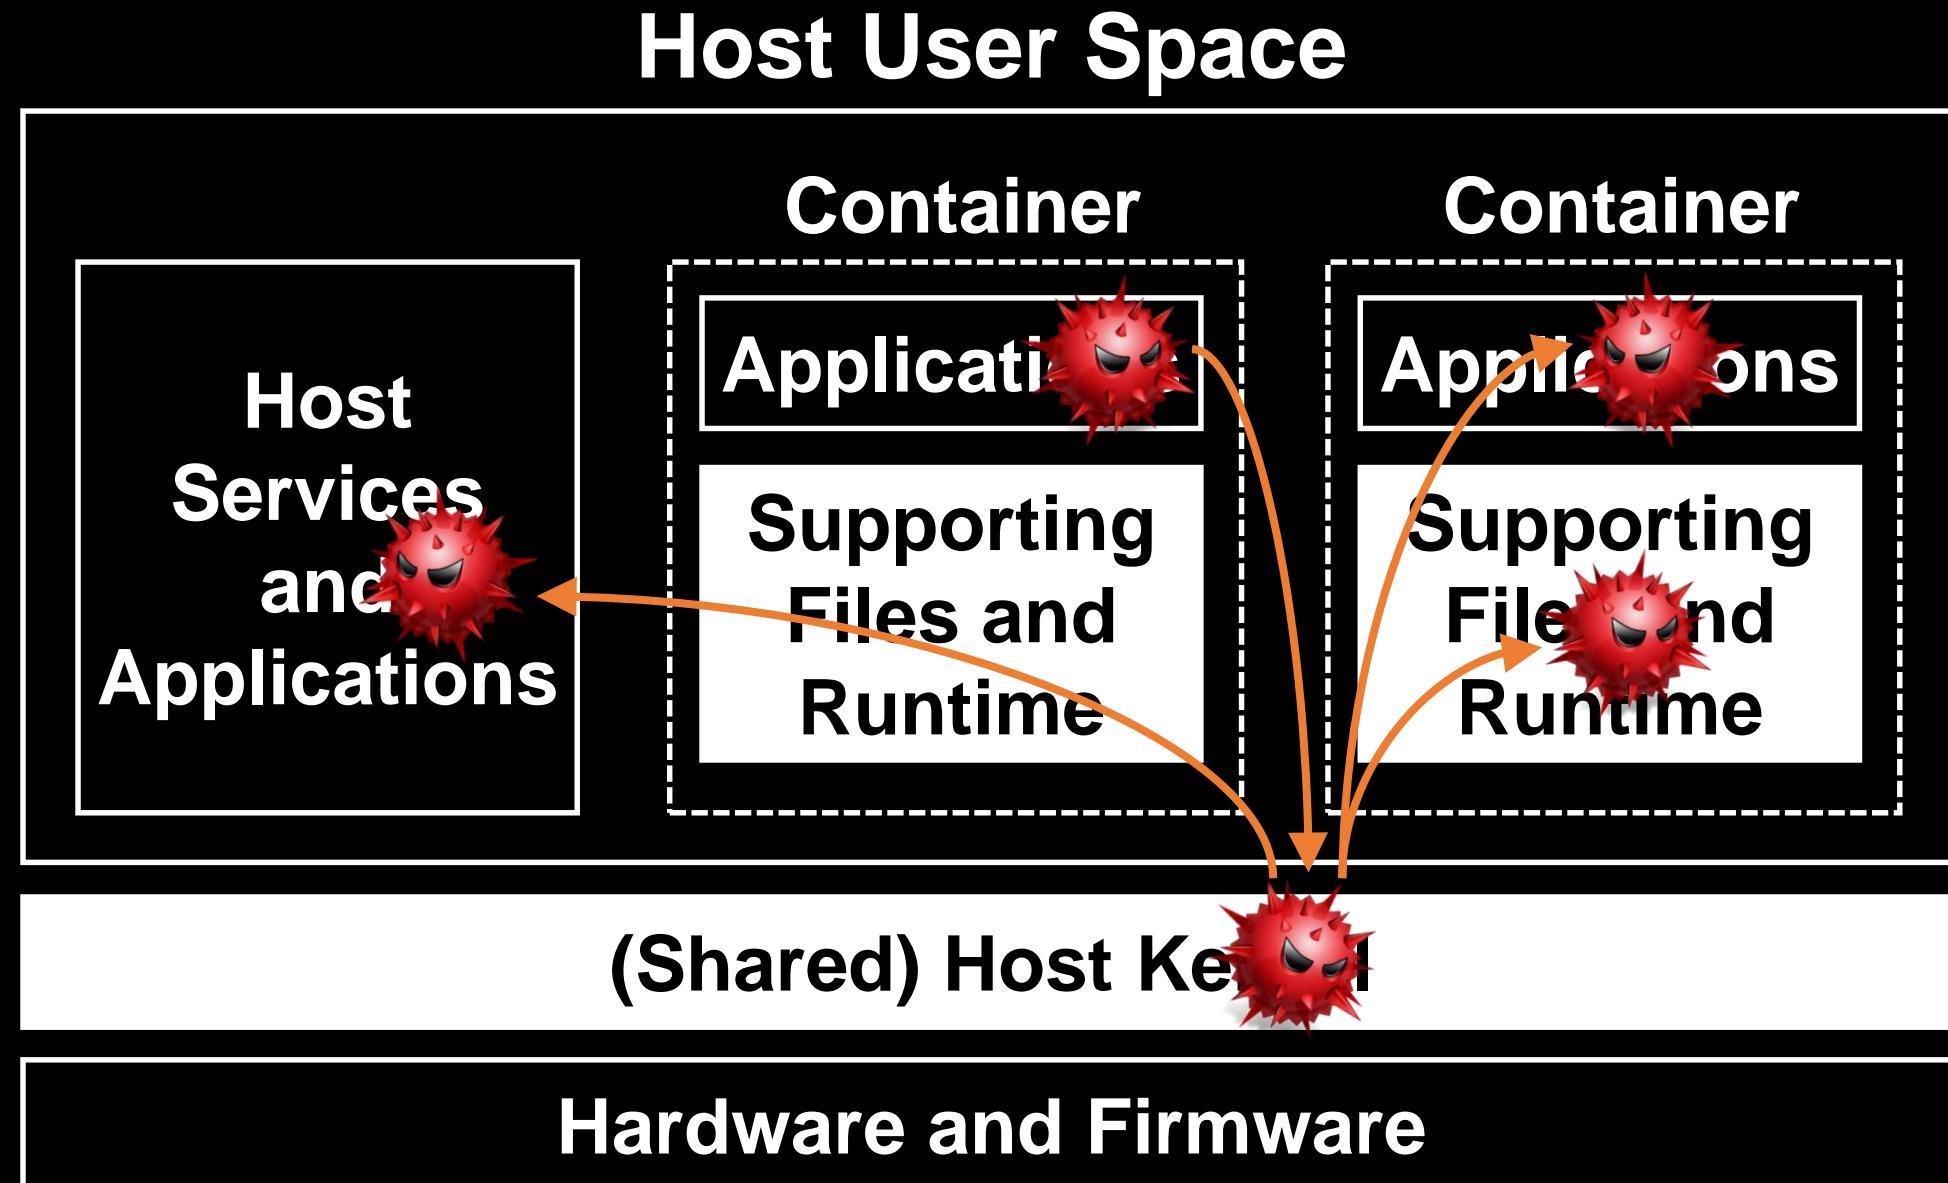
- Virtual machine architecture

- uses **Virtual Machine Monitor (VMM)**, a.k.a. hypervisor
- provides strong isolation with virtualized hardware
 - A guest OS is needed, and performance overhead is high
- has each execution environment, such as kernel and user space

- Container architecture

- uses **namespace and cgroup features of Linux kernel**
- provides weak isolation with the kernel
 - Supporting files are needed, and performance overhead is low
- shares the kernel and has its own user space only

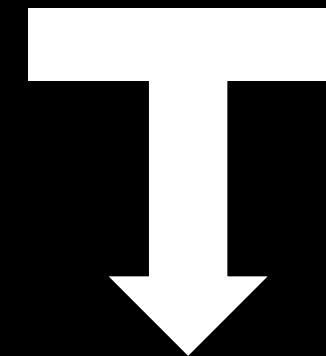
The container is good, but the problem is ...



Containers are indispensable!

But, **strong isolation** is needed
to prevent the **ESCAPES**

Container



Hypervisor



kata
containers



Firecracker

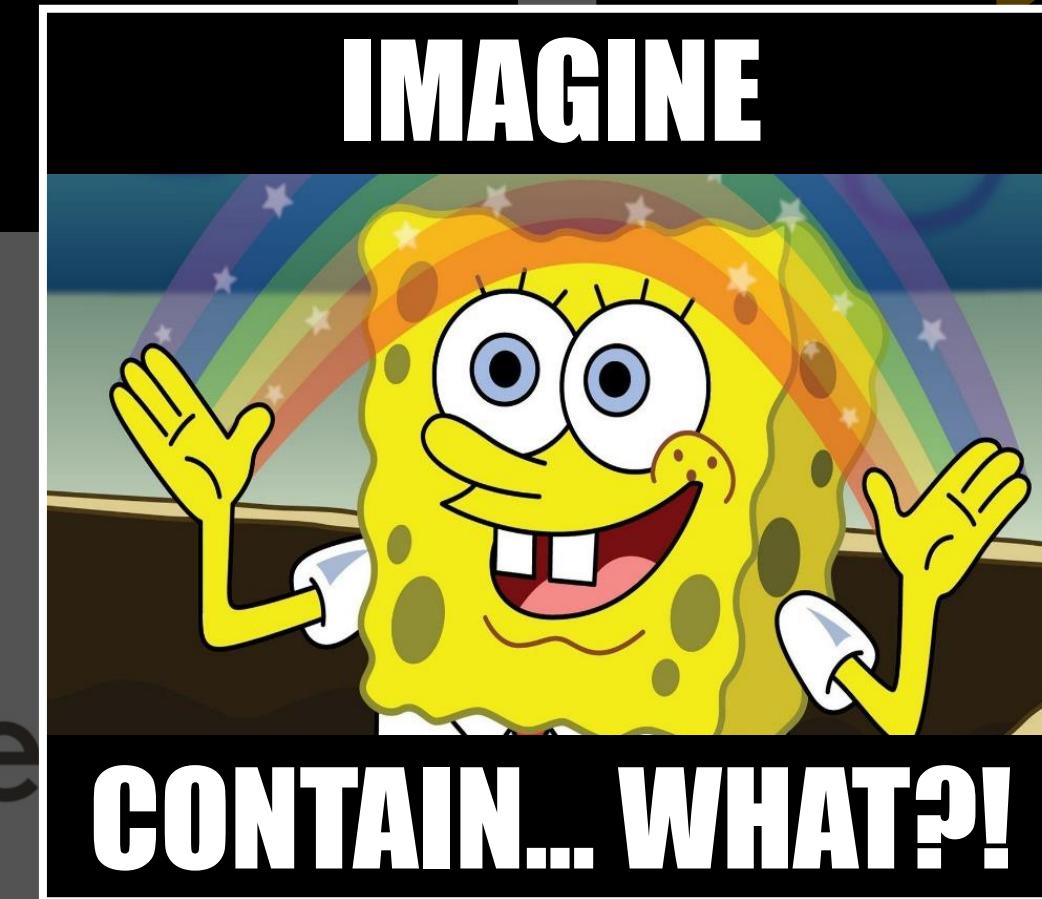


gVisor

Container



Hypervisor



Fire

Visor

To be or not to be,
that is the question



- William Shakespeare

So, the security researcher realized ...

To be a container or not to be a container,
that is **NOT** the question



- Anonymous Researcher

The **REAL** question is ...

How can we prevent escapes?

ALCATRAZ



It's me!

FOR



- Background
- Analysis of Critical Paths of Escapes
- Design and Implementation of Alcatraz
- Evaluation and Demo
- Conclusion and Black Hat Sound Bytes

- Background



- Analysis of Critical Paths of Escapes

- Design and Implementation of Alcatraz

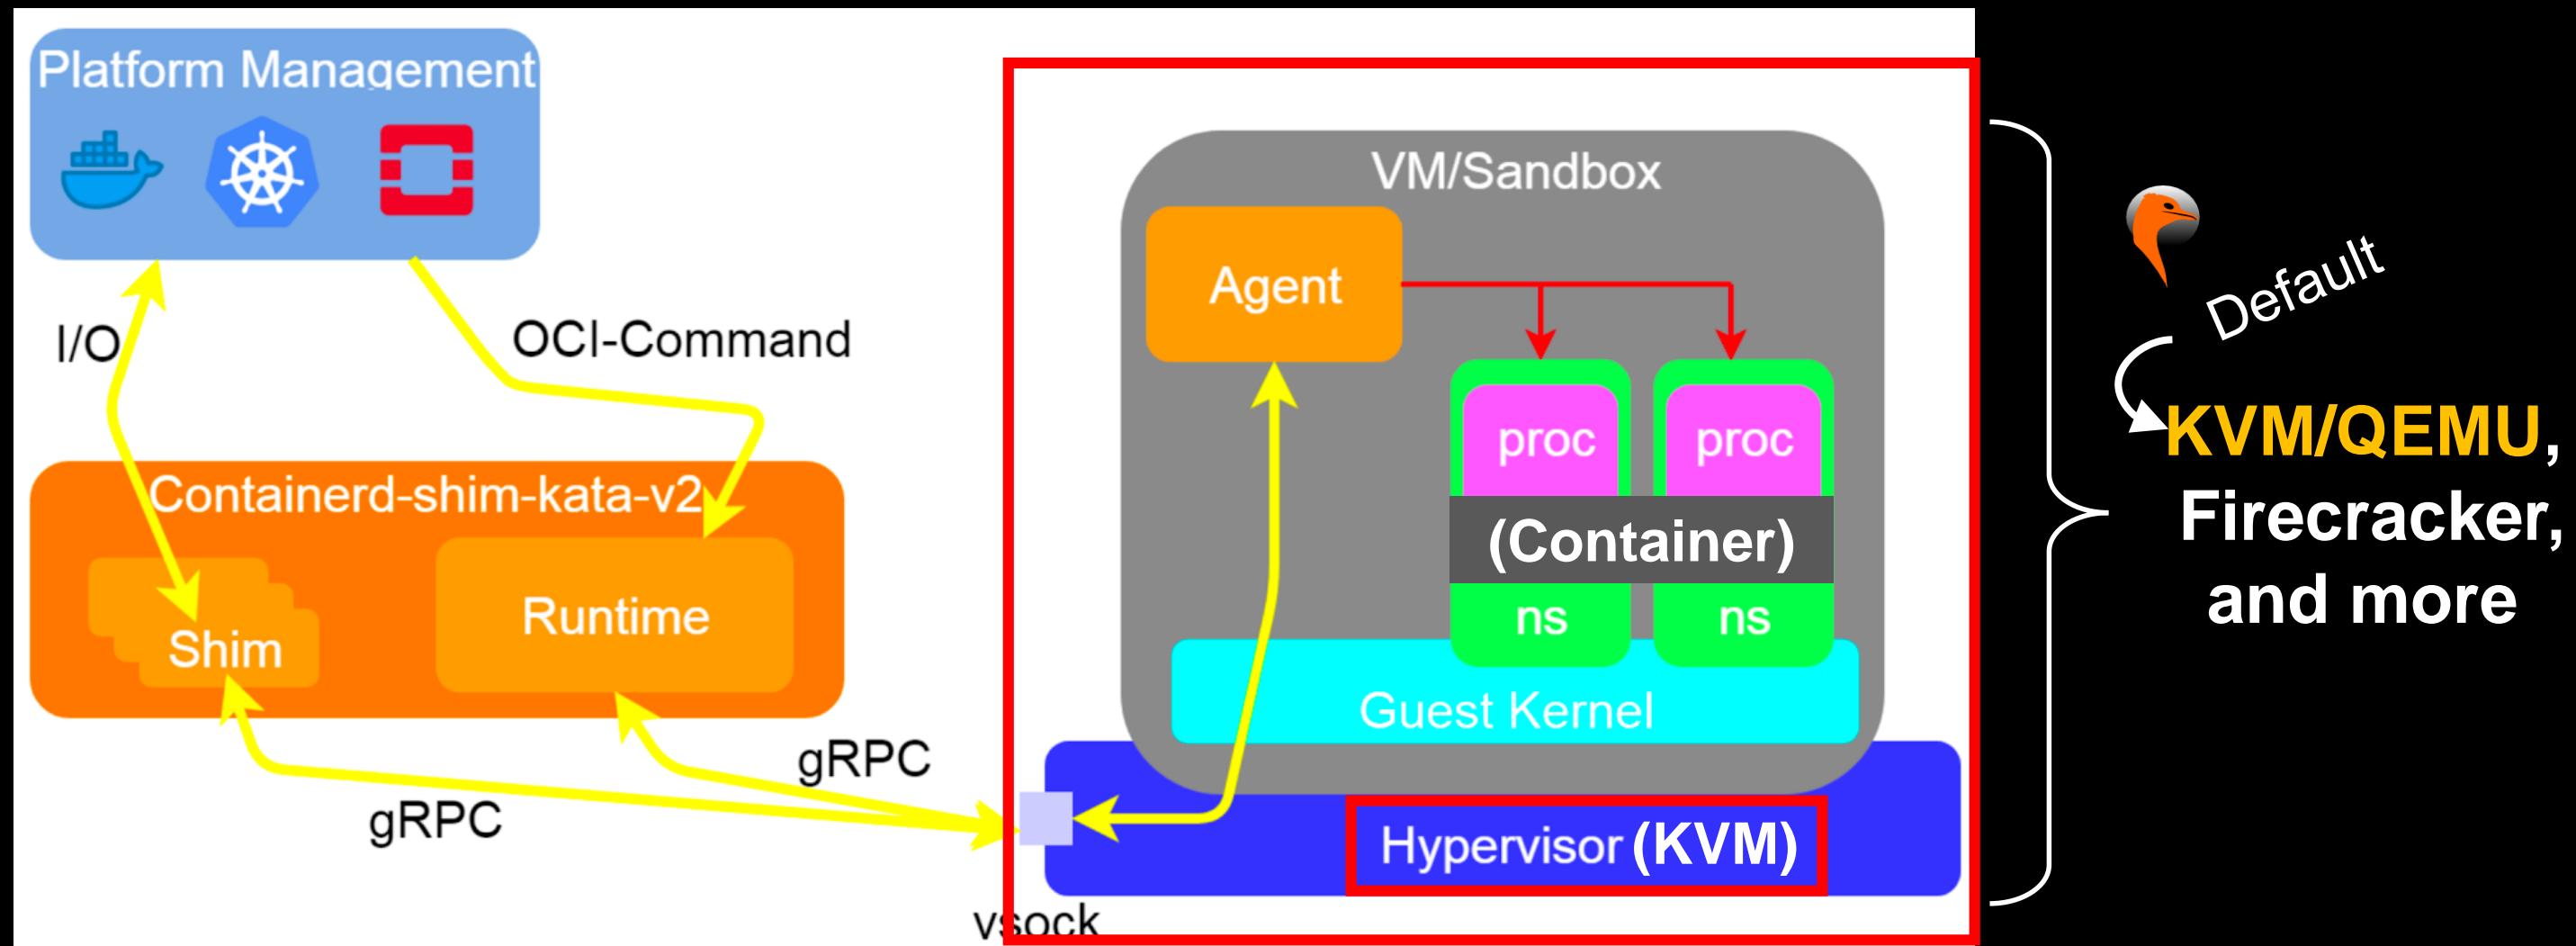
- Evaluation and Demo

- Conclusion and Black Hat Sound Bytes



Kata Container

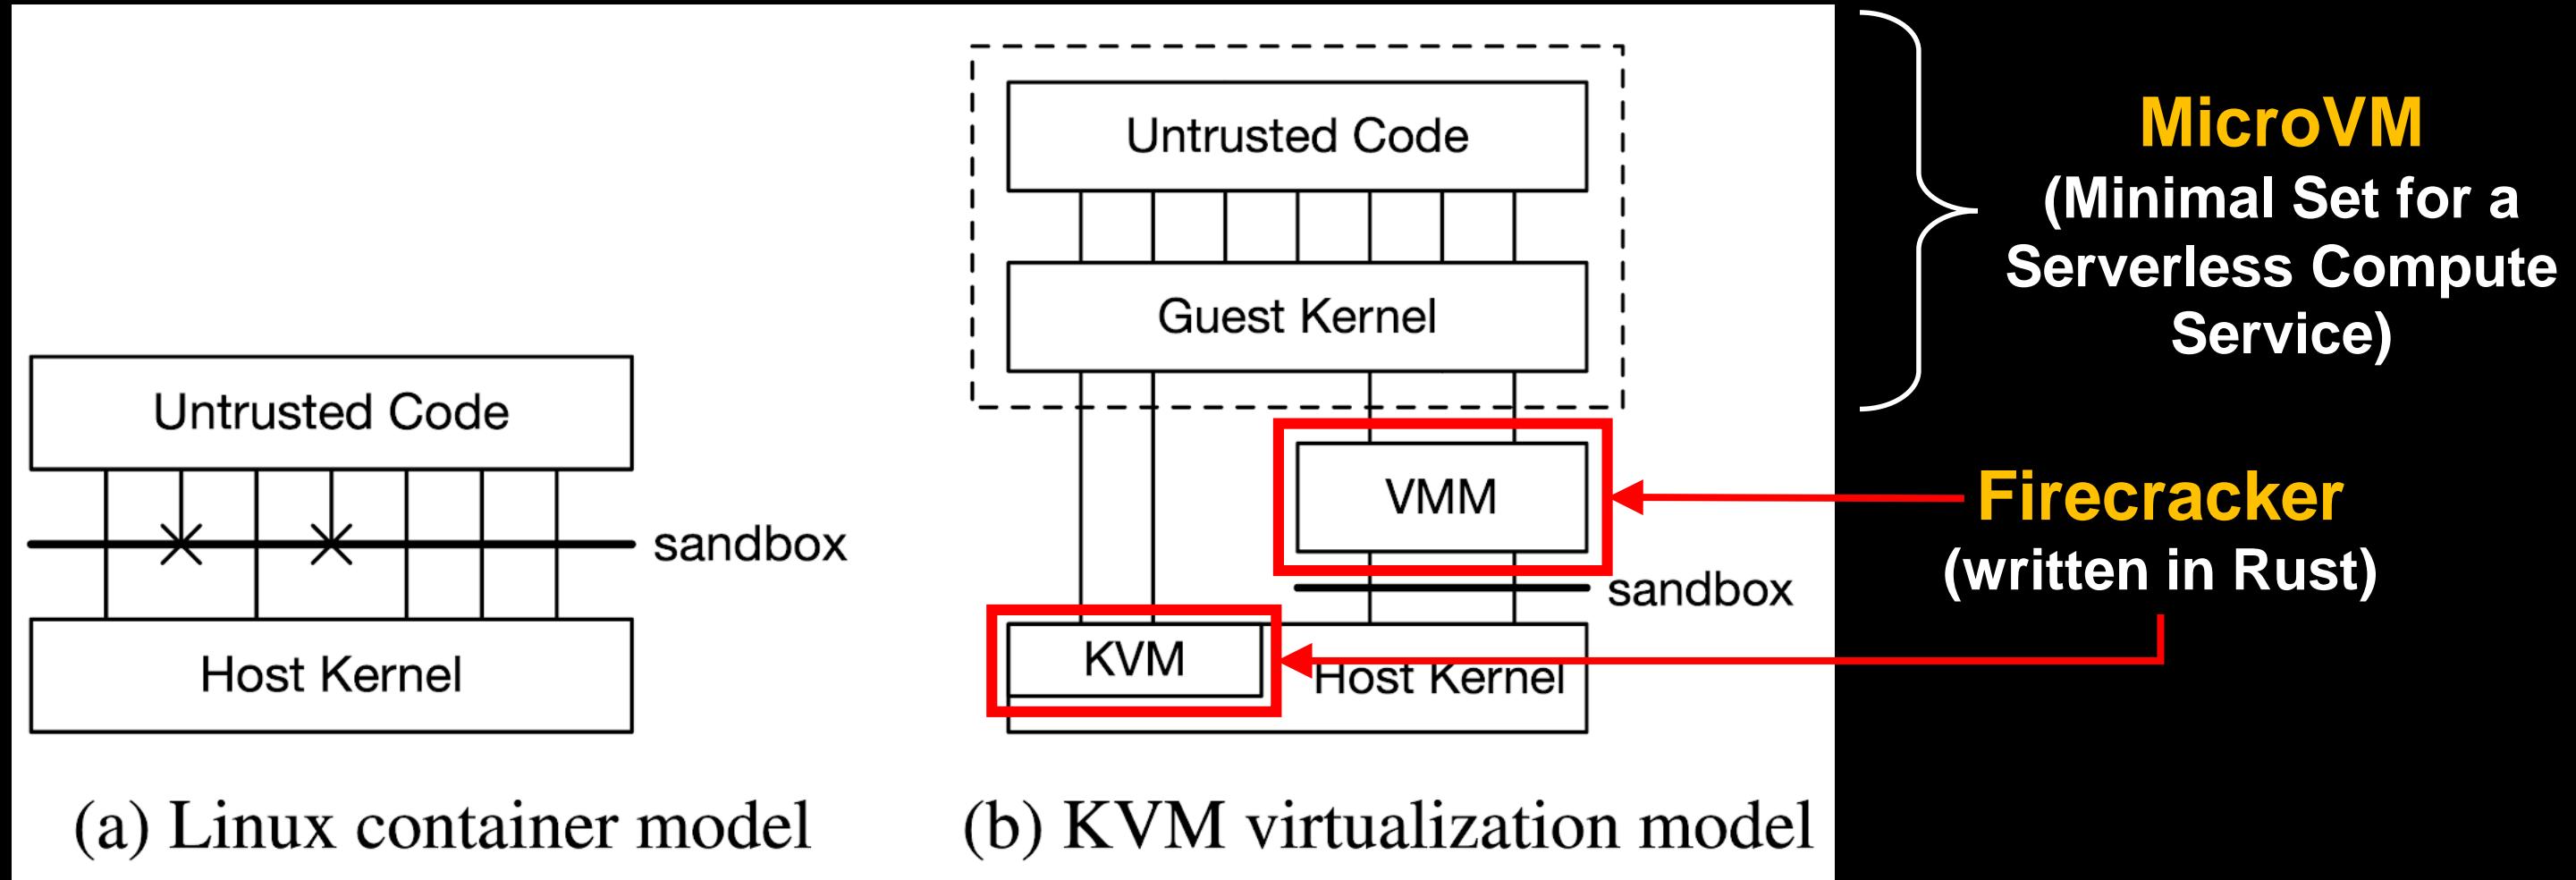
Kata Container



**Kata Architecture - Kata Containers: An Emerging Architecture
for Enabling MEC Services in Fast and Secure Way**



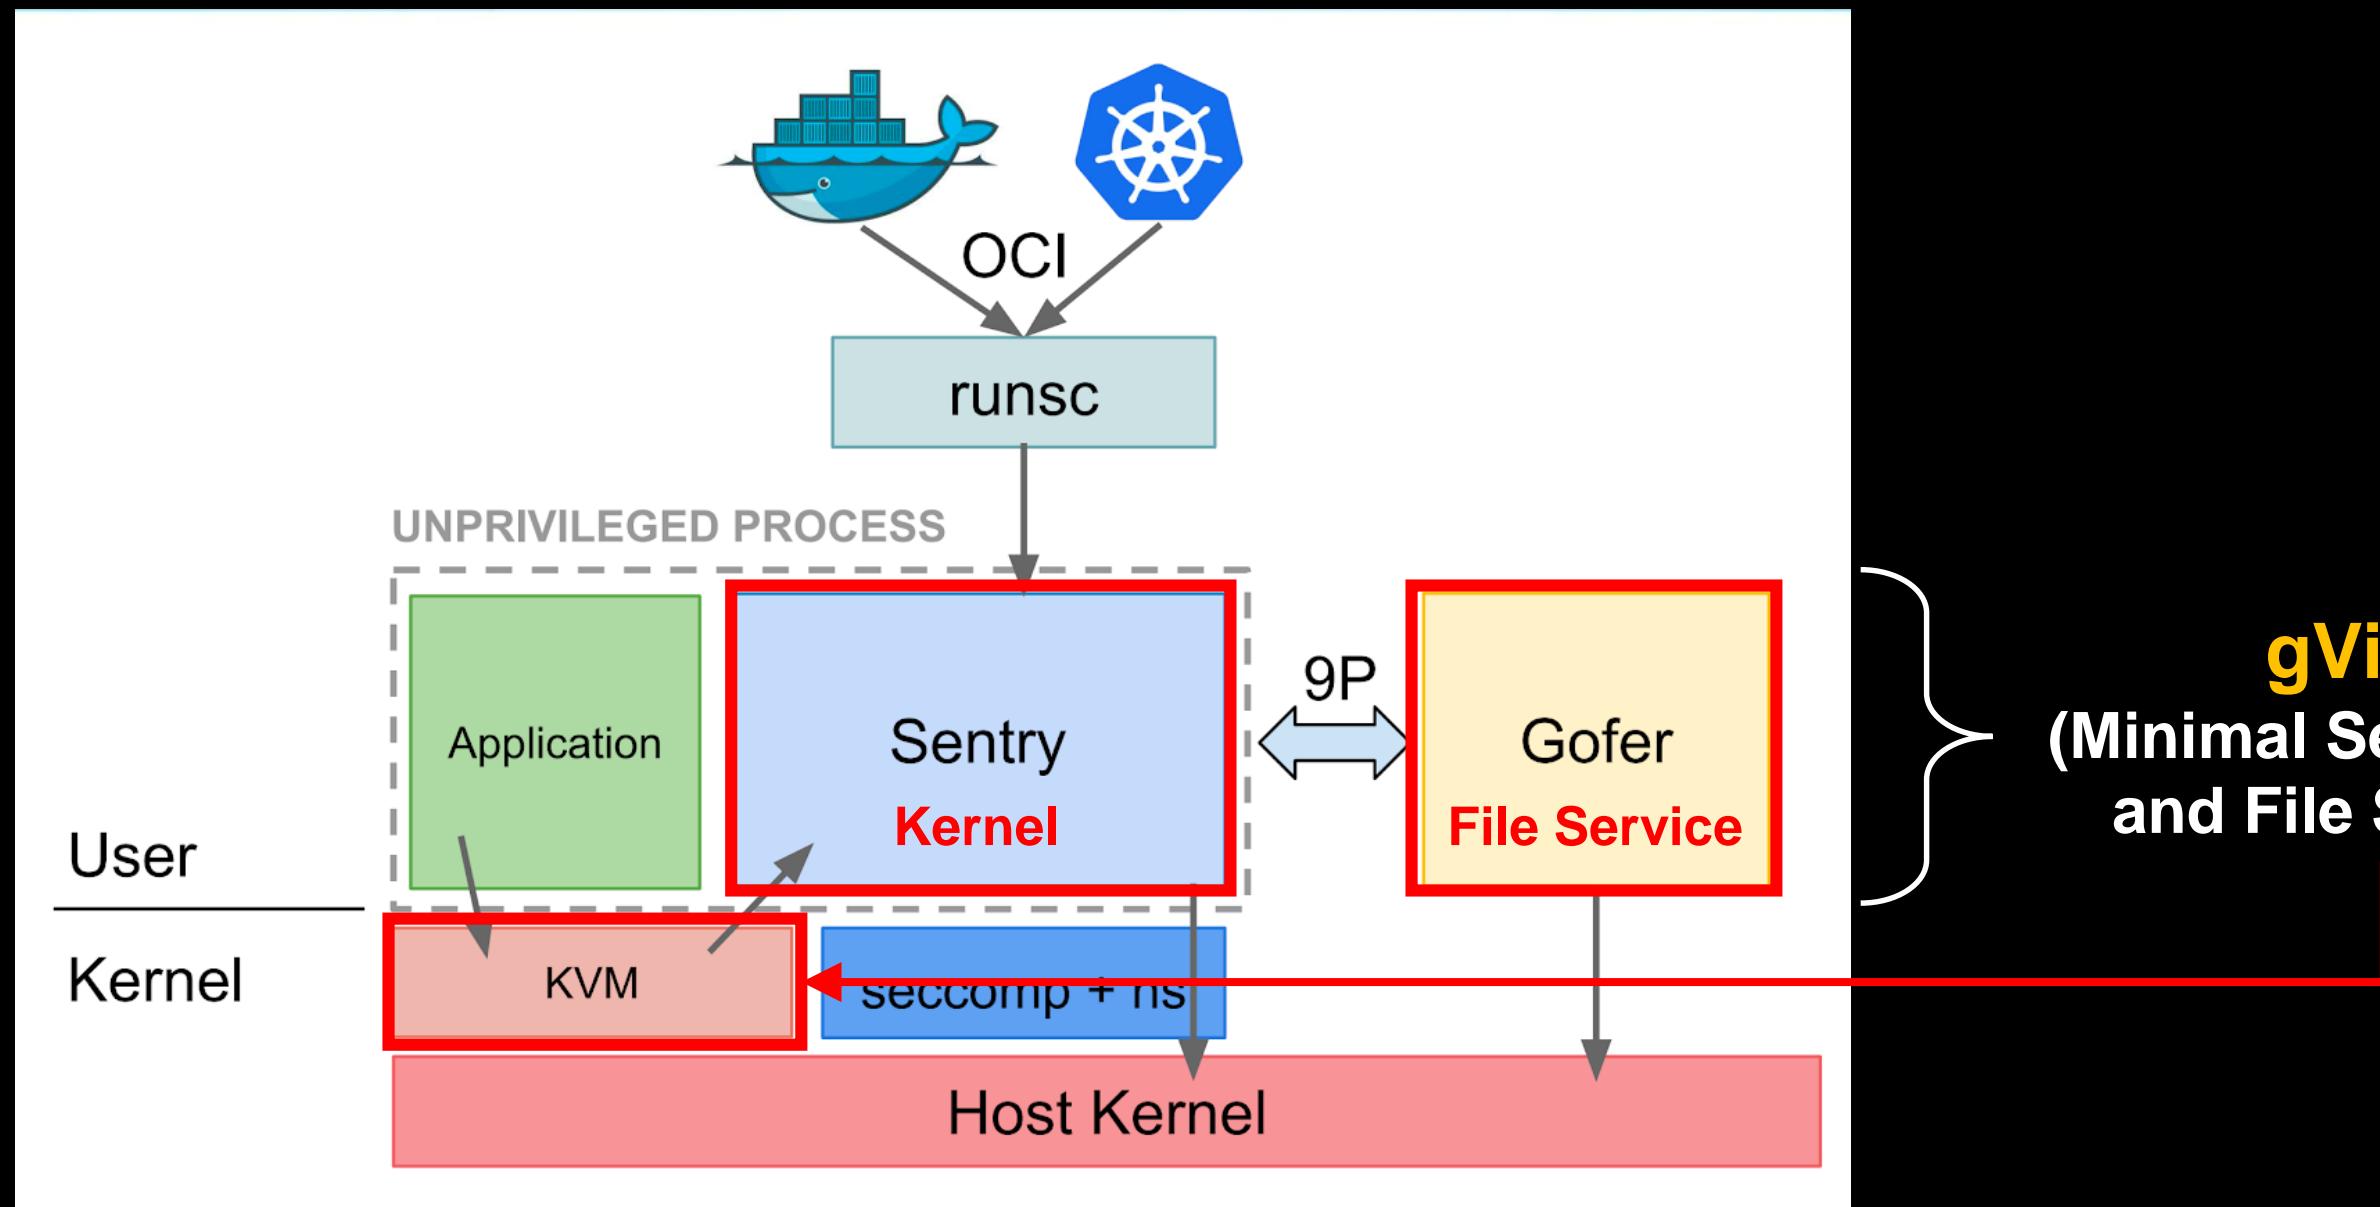
Amazon Firecracker



**The Security Model of Linux Containers - Firecracker:
Lightweight Virtualization for Serverless Applications**

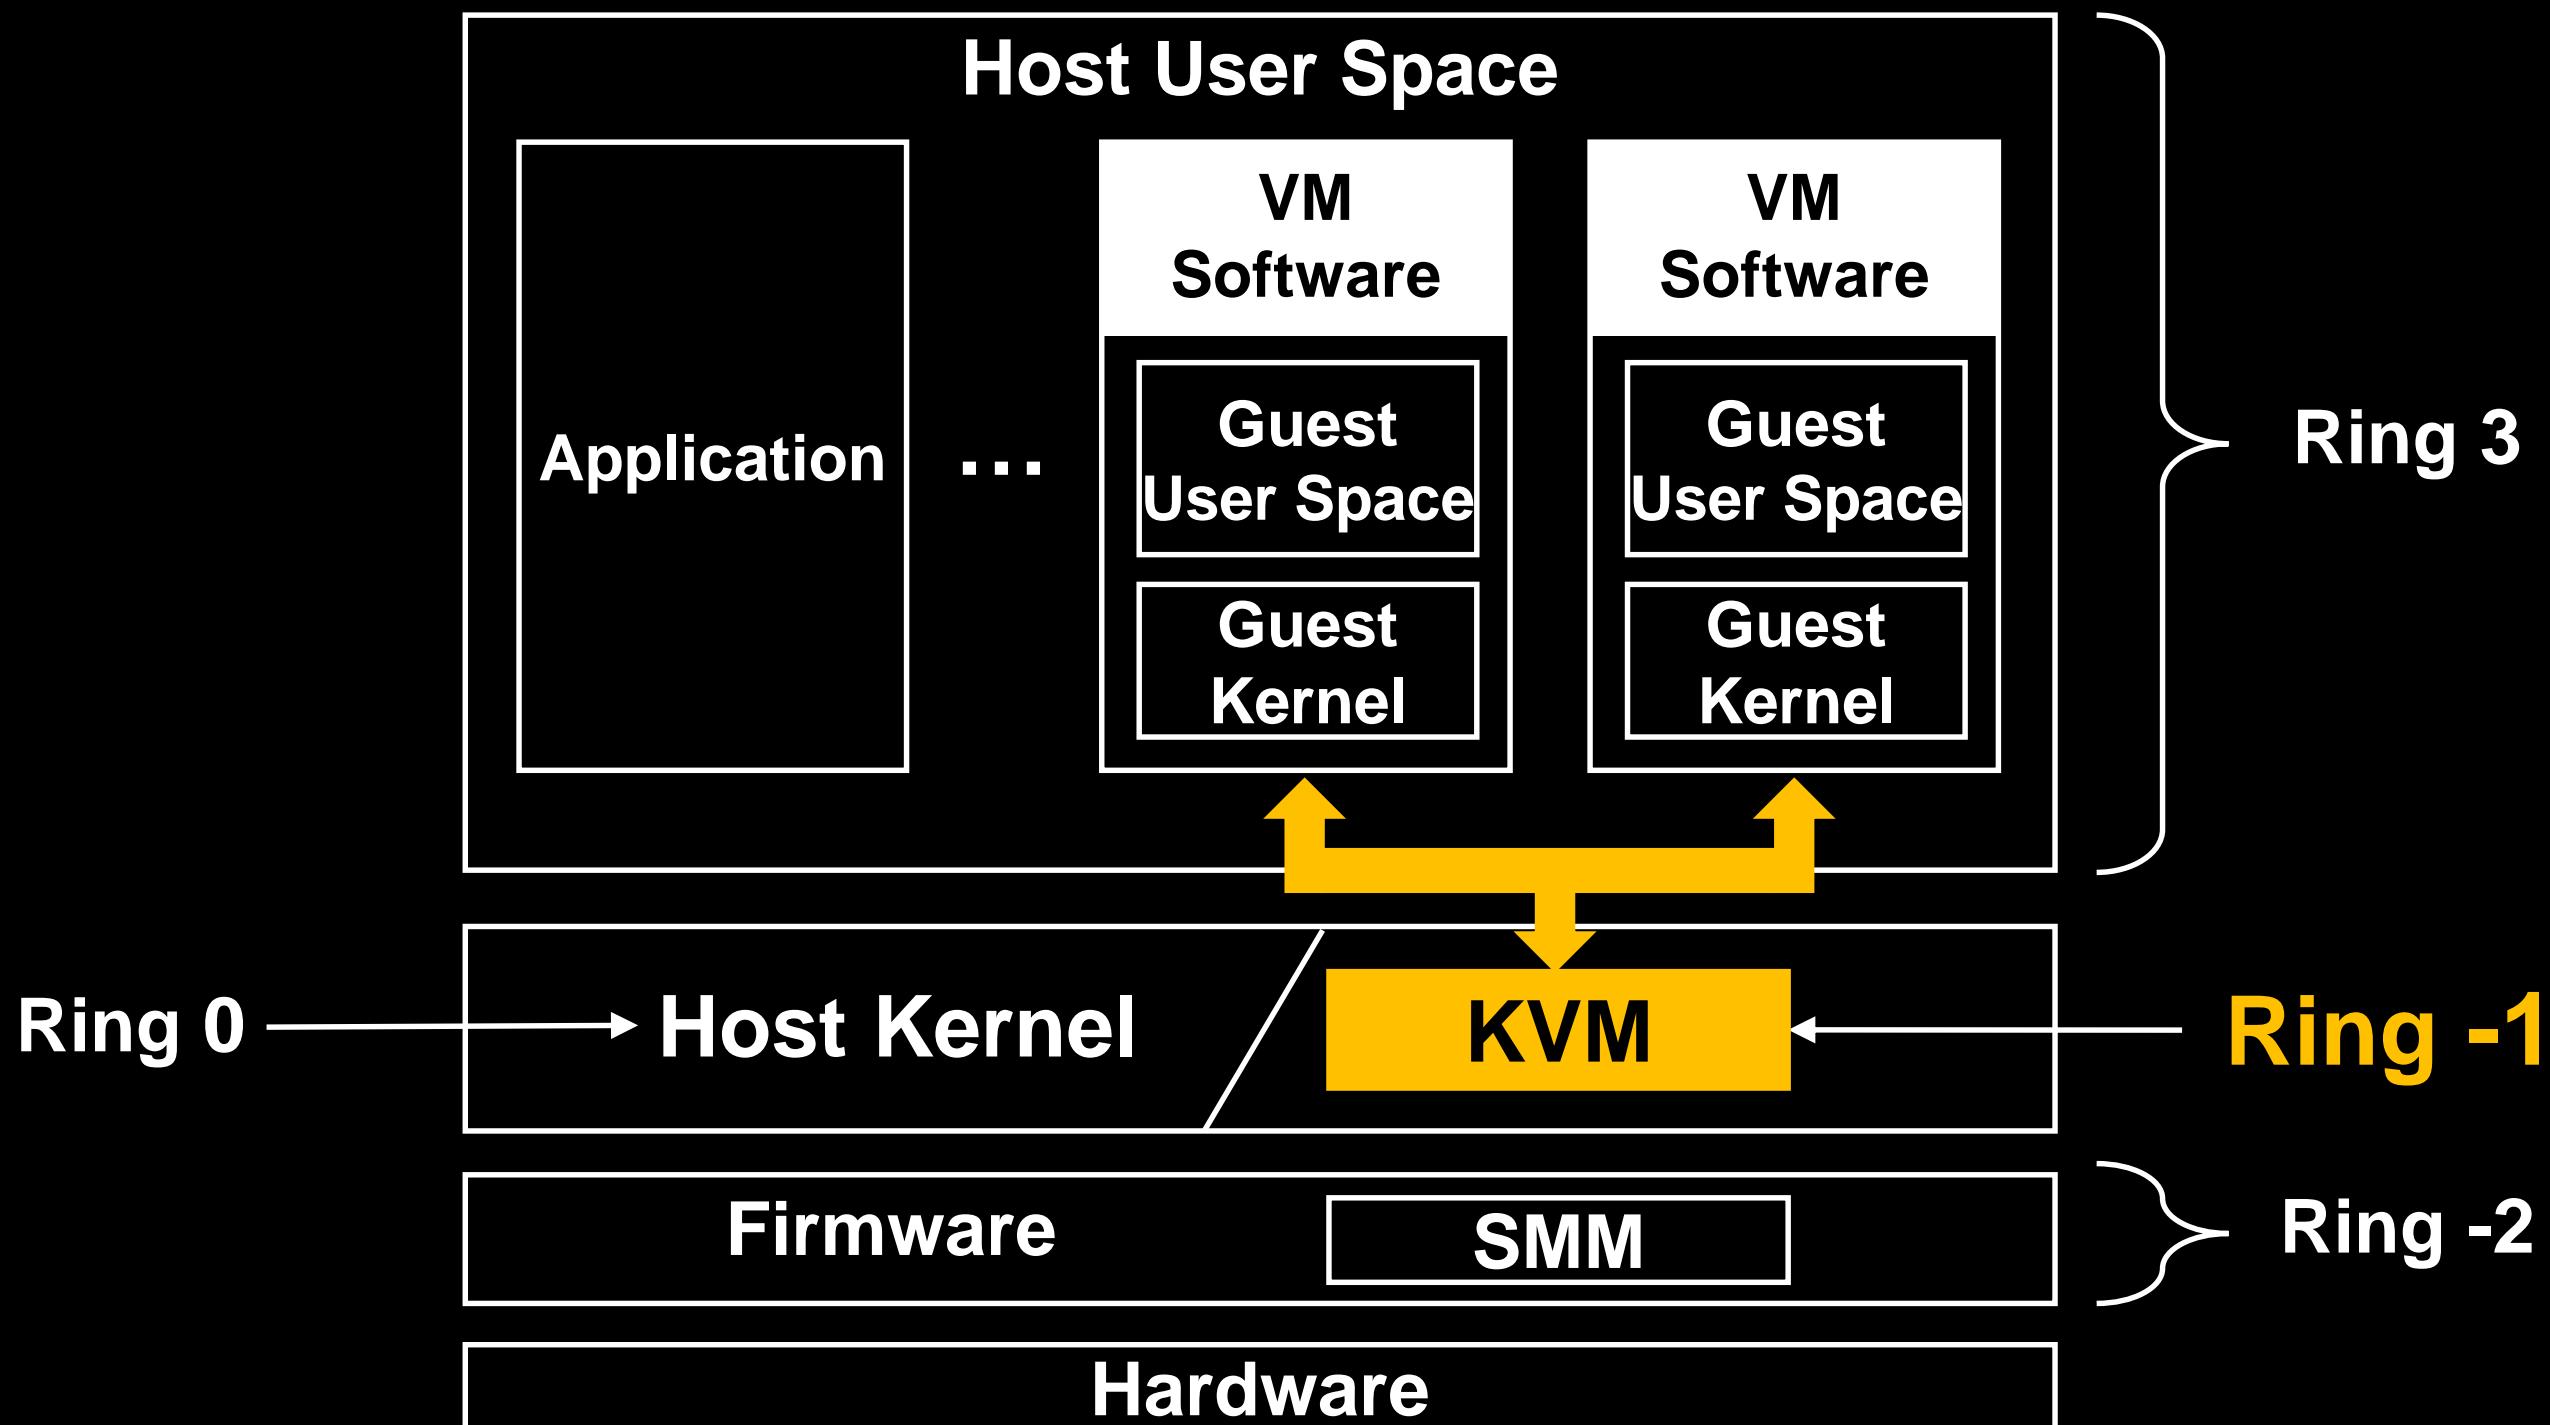


Google gVisor

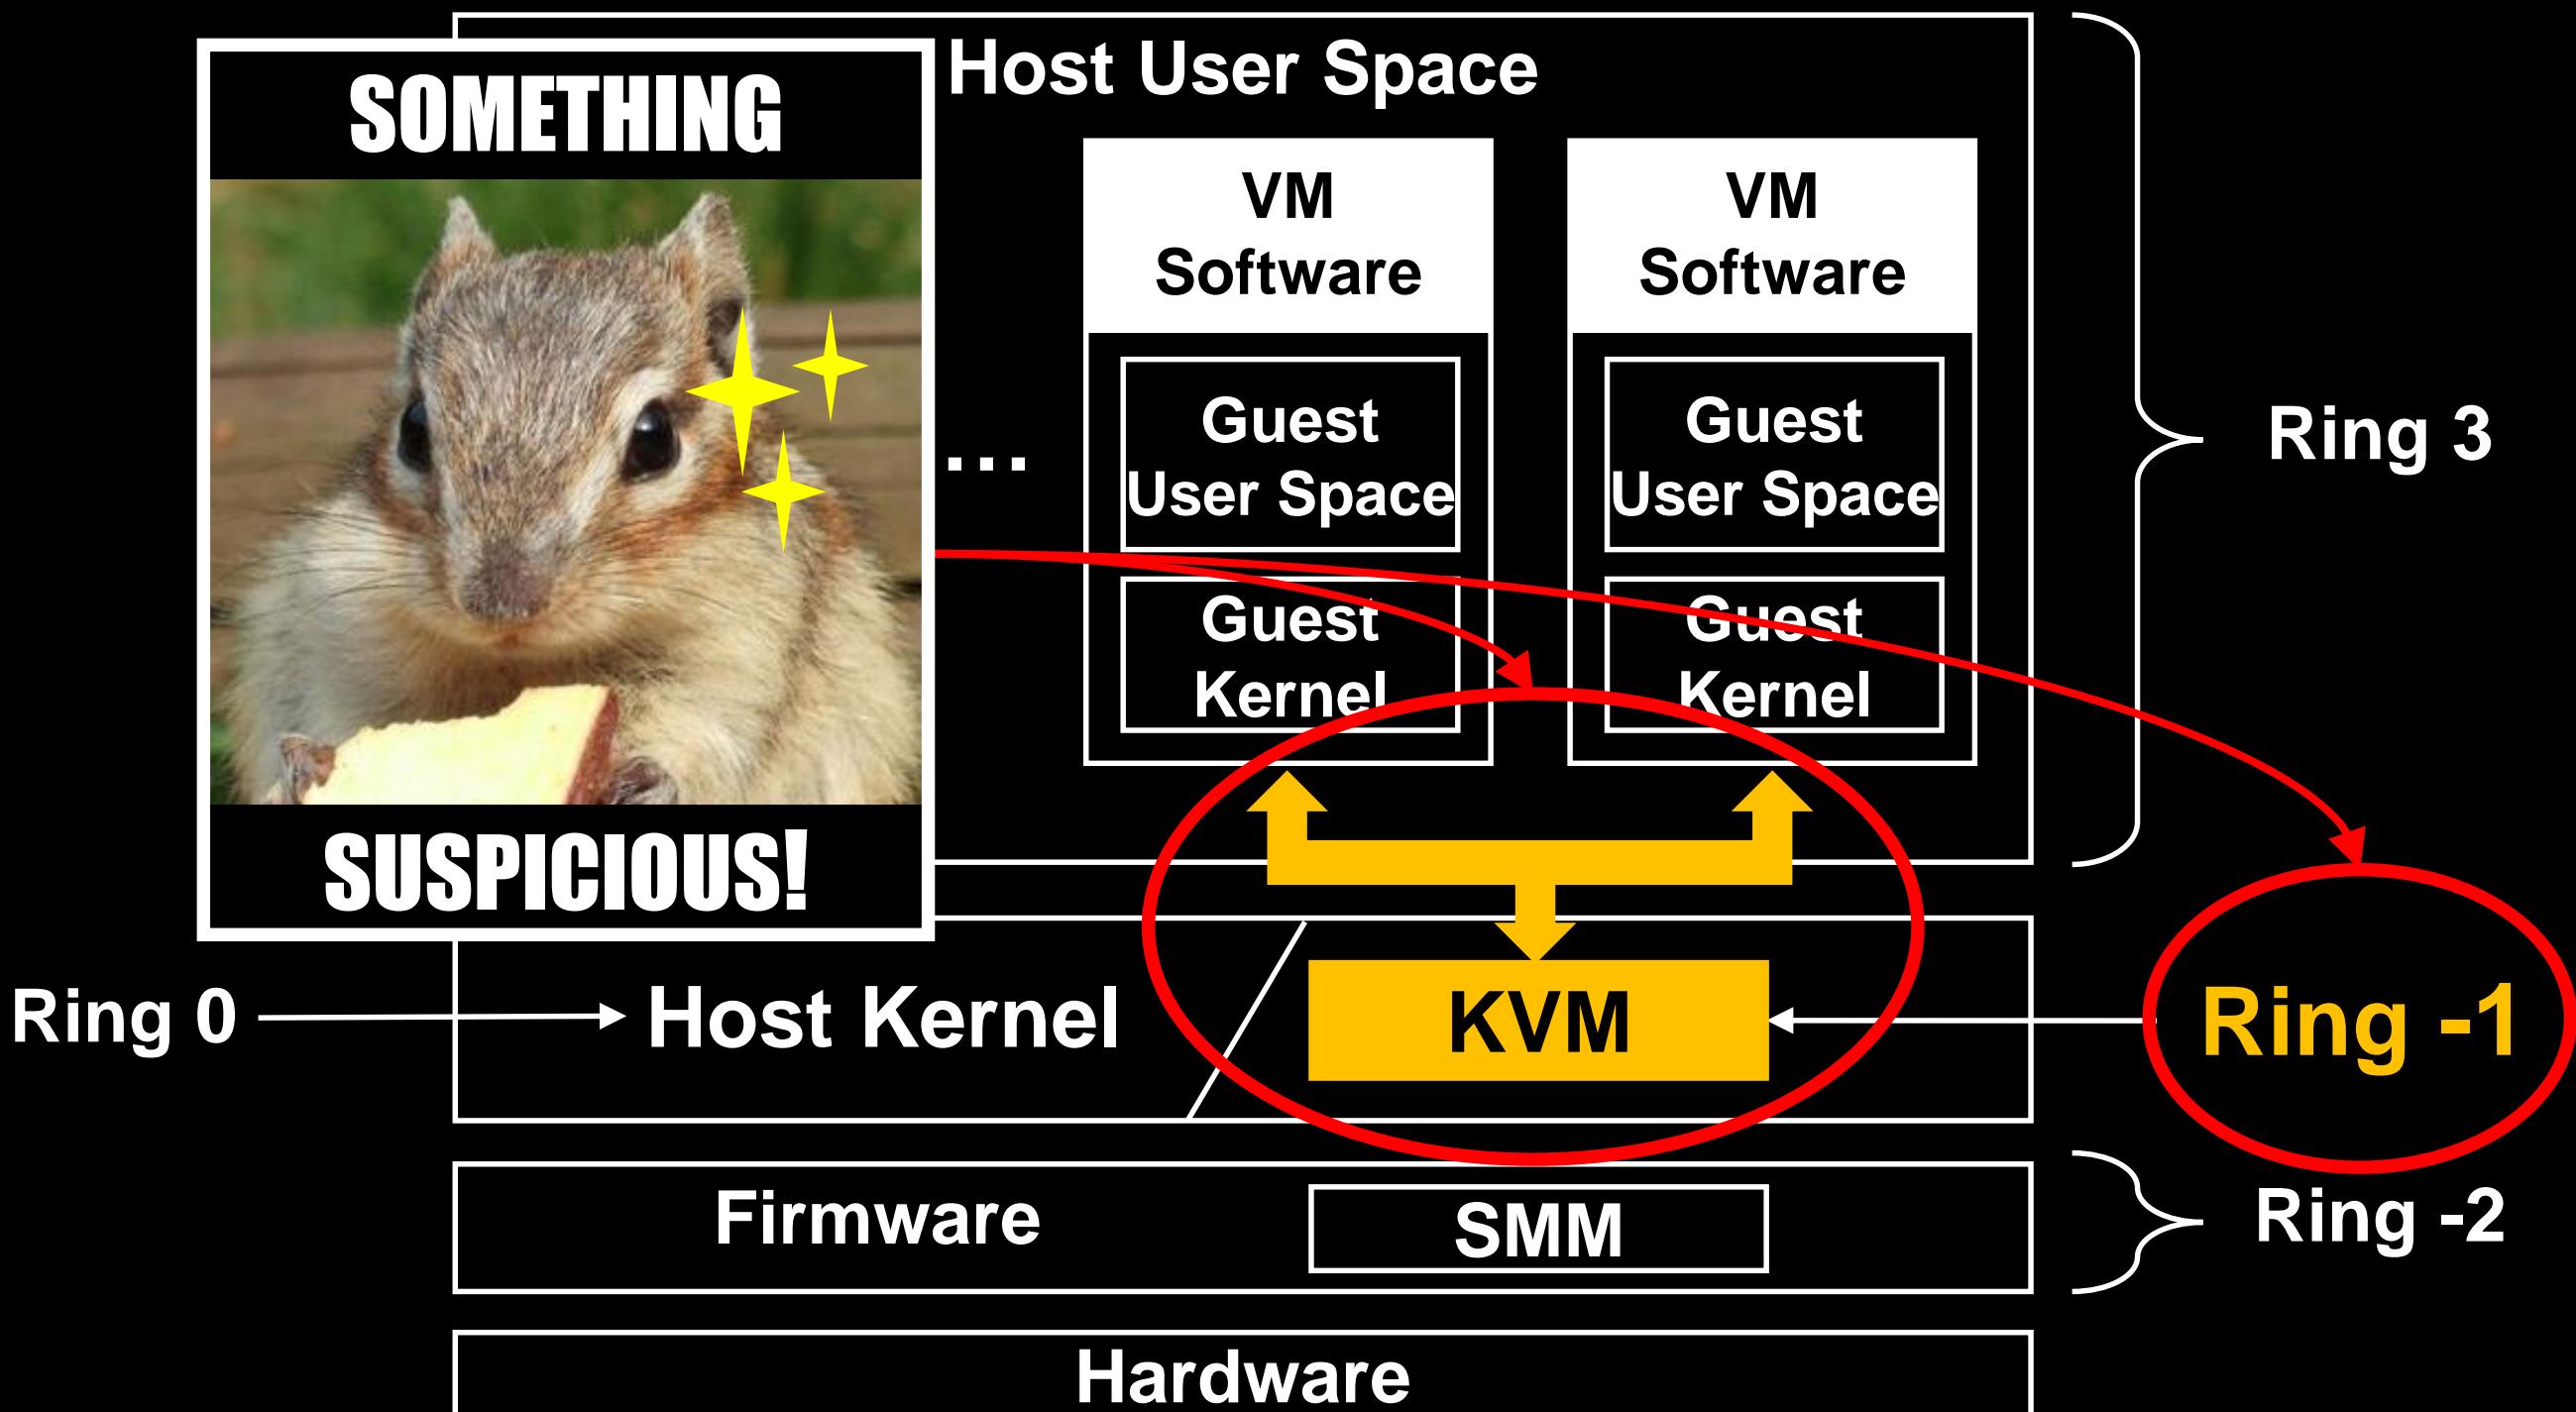


An OCI runtime powered by gVisor-Container Isolation
at Scale (... and introducing gVisor)

Kernel-based Virtual Machine (KVM)



Kernel-based Virtual Machine (KVM)



A hypervisor attaches to VMs **tightly**
It has **Ring -1** privilege!

The attacker can **escape** from the VM
with a **KVM vulnerability**
(and subvert the system)

Fortunately,
I was not the only one who was worried about
the excessive privilege of a hypervisor!



Previous Works

HyperSentry: Enabling Stealthy In-context Measurement of Hypervisor Integrity *

Ahmed M. Azab Peng Ning Zhi Wang Xuxian Jiang
Department of Computer Science, North Carolina State University
{amazab, pning, zhi_wang, xuxian_jiang}@ncsu.edu

Xiaolan Zhang Nathan C. Skalsky
IBM T.J. Watson Research Center IBM Systems & Technology Group
cxzhang@us.ibm.com nskalsky@us.ibm.com

ABSTRACT

This paper presents *HyperSentry*, a novel framework to enable integrity measurement of a running hypervisor (or any other highest privileged software layer on a system). Unlike existing solutions for protecting privileged software, HyperSentry does not introduce a higher privileged software layer below the integrity measurement target, which could start another race with malicious attackers in obtaining the highest privilege in the system. Instead, HyperSentry introduces a software component that is properly isolated from the hypervisor to enable *stealthy* and *in-context* measurement of the runtime integrity of the hypervisor. While stealthiness is necessary to ensure that a compromised hypervisor does not have a chance to hide the attack traces upon detecting an up-coming measurement, in-context measurement is necessary to allow the developer to perform useful

we implement a prototype of the framework along with an integrity measurement agent for the Xen hypervisor. Our experimental evaluation shows that HyperSentry is a low-overhead practical solution for real world systems.

Categories and Subject Descriptors

D.4.6 [Operating Systems]: Security and Protection—*invasive software*

General Terms

Security

Keywords

Virtualization, Hypervisor Integrity, Integrity Measurement

1. INTRODUCTION



CHIPSET BASED APPROACH TO DETECT VIRTUALIZATION MALWARE a.k.a. DeepWatch

Yuriy Bulygin

Joint work with David Samyde

Security Center of Excellence / PSIRT @ Intel Corporation

Preventing and Detecting Xen Hypervisor Subversions

Joanna Rutkowska & Rafał Wojtczuk
Invisible Things Lab

Black Hat USA 2008, August 7th, Las Vegas, NV

Protecting Cloud Virtual Machines from Commodity Hypervisor and Host Operating System Exploits

Shih-Wei Li John S. Koh Jason Nieh
Department of Computer Science
Columbia University
{shihwei, koh, nieh}@cs.columbia.edu

Abstract

Hypervisors are widely deployed by cloud computing providers to support virtual machines, but their growing complexity poses a security risk as large codebases contain many vulnerabilities. We have created HypSec, a new hypervisor design for retrofitting an existing commodity hypervisor using microkernel principles to reduce its trusted computing base while protecting the confidentiality and integrity of virtual machines. HypSec partitions the hypervisor into an untrusted host that performs most complex hypervisor functionality without access to virtual machine data, and a trusted core that provides access control to virtual machine data and performs basic CPU and memory virtualization. Hardware virtualization support is used to isolate and protect the trusted core and execute it

maintenance effort. For example, KVM [44] is integrated with Linux and Hyper-V [56] is integrated with Windows. The result is a huge potential attack surface with access to VM data in CPU registers, memory, I/O data, and boot images. The surge in outsourcing of computational resources to the cloud and away from privately-owned data centers further exacerbates this security risk of relying on the trustworthiness of complex and potentially vulnerable hypervisor and host OS infrastructure. Attackers that successfully exploit hypervisor vulnerabilities can gain unfettered access to VM data, and compromise the privacy and integrity of all VMs—an undesirable outcome for both cloud providers and users.

Recent trends in application design and hardware virtualization support provide an opportunity to revisit hypervisor design requirements to address this crucial security problem.

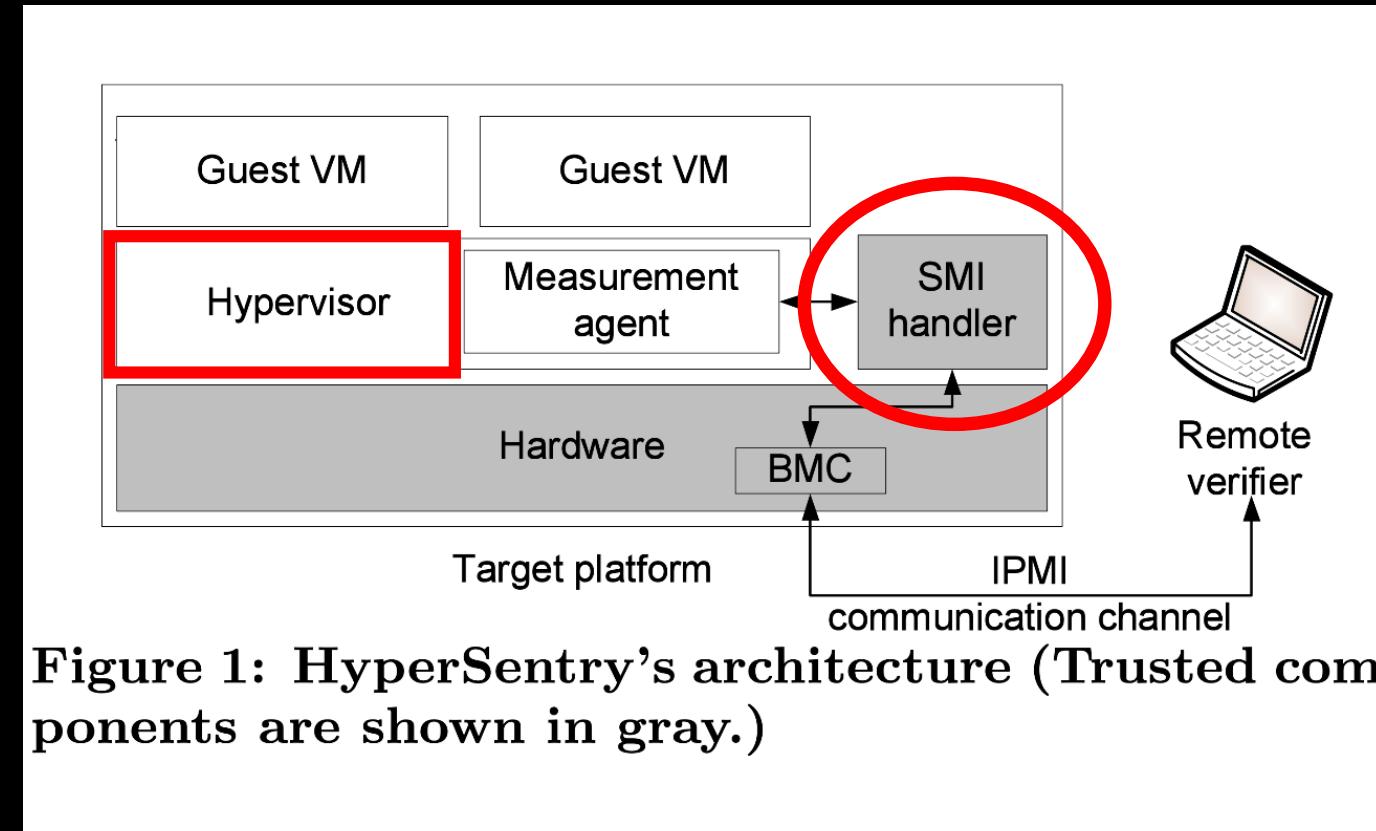


Figure 1: HyperSentry's architecture (Trusted components are shown in gray.)

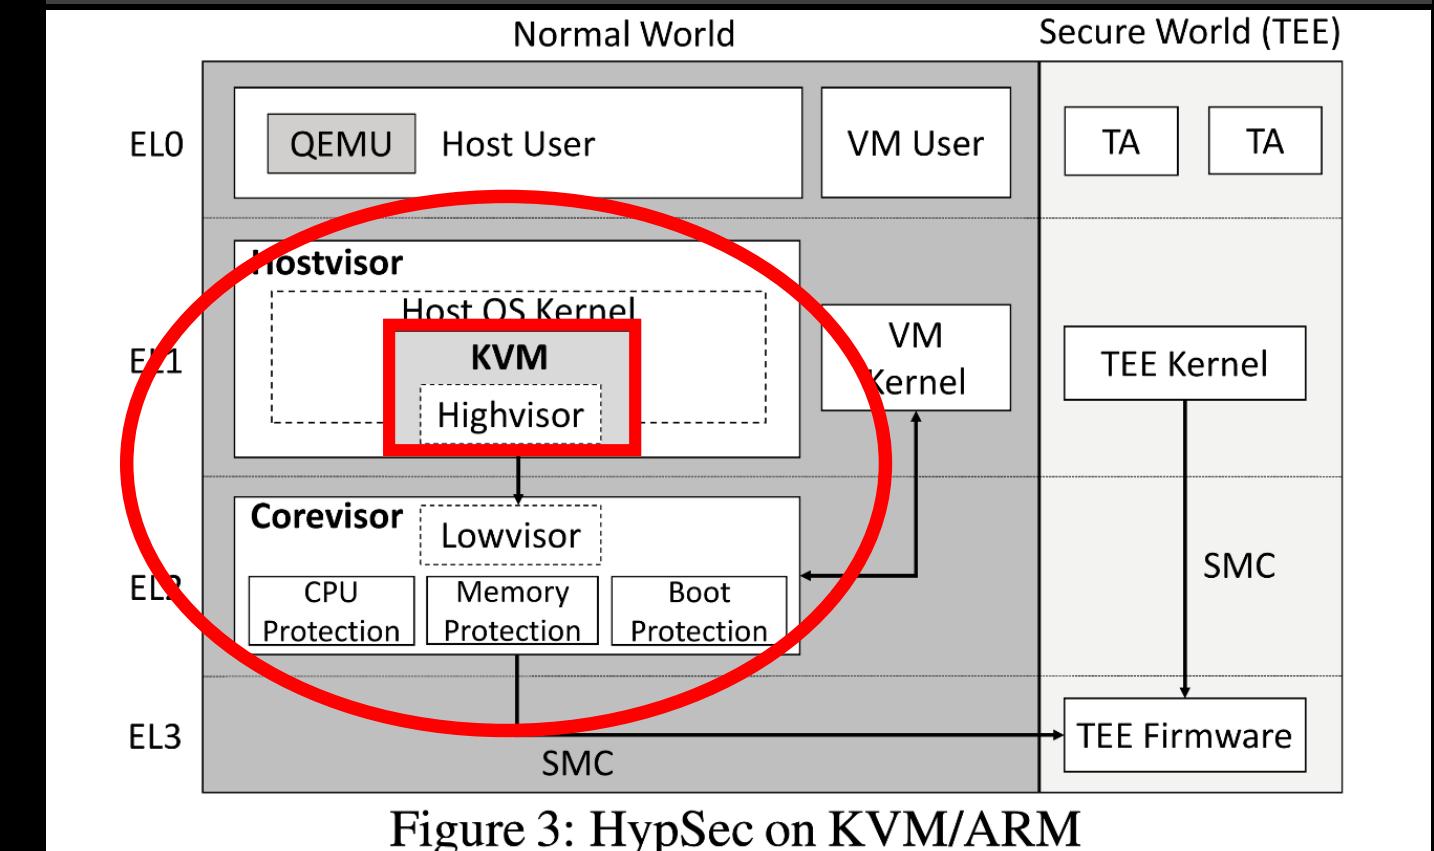
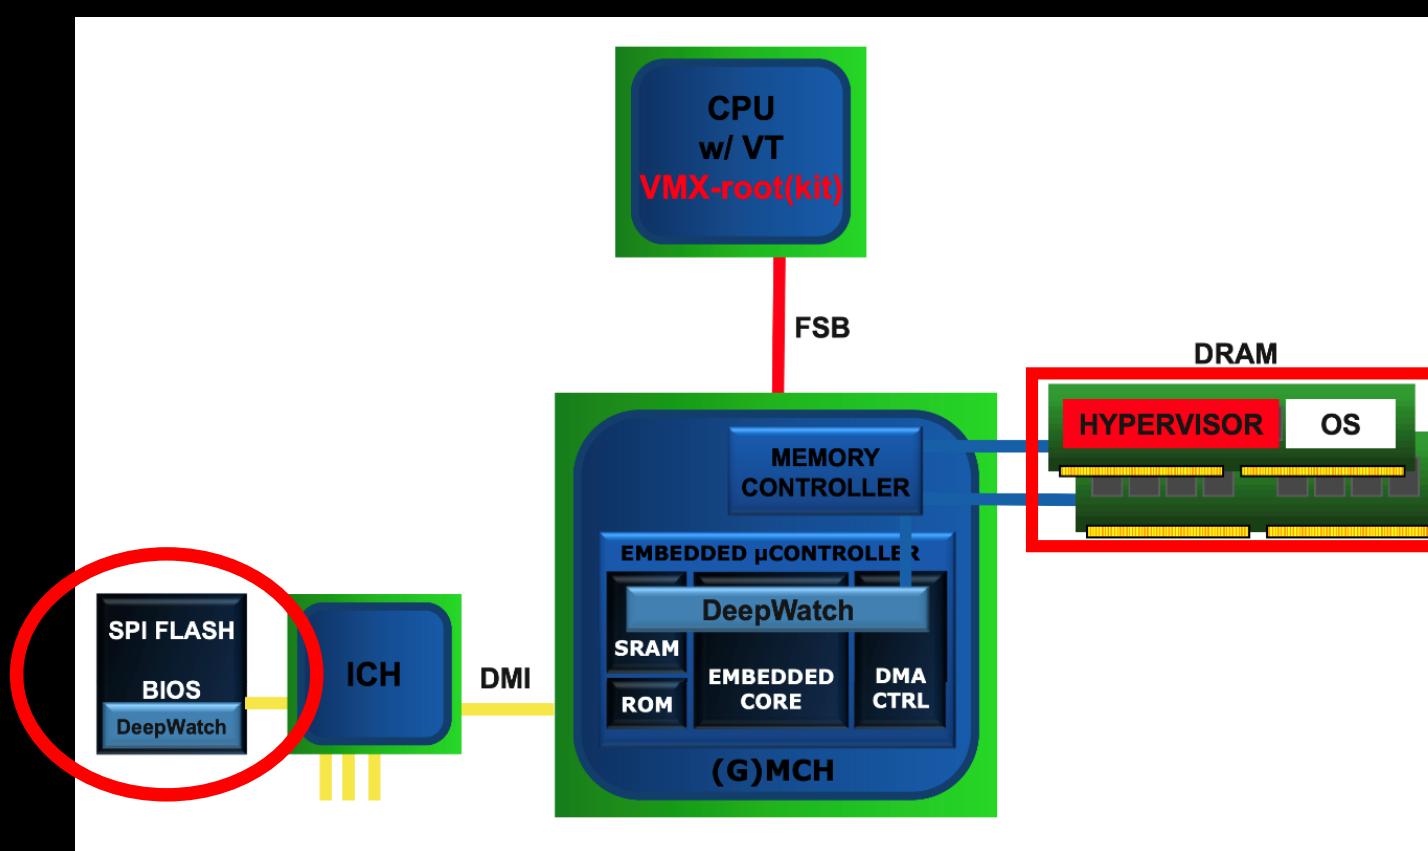
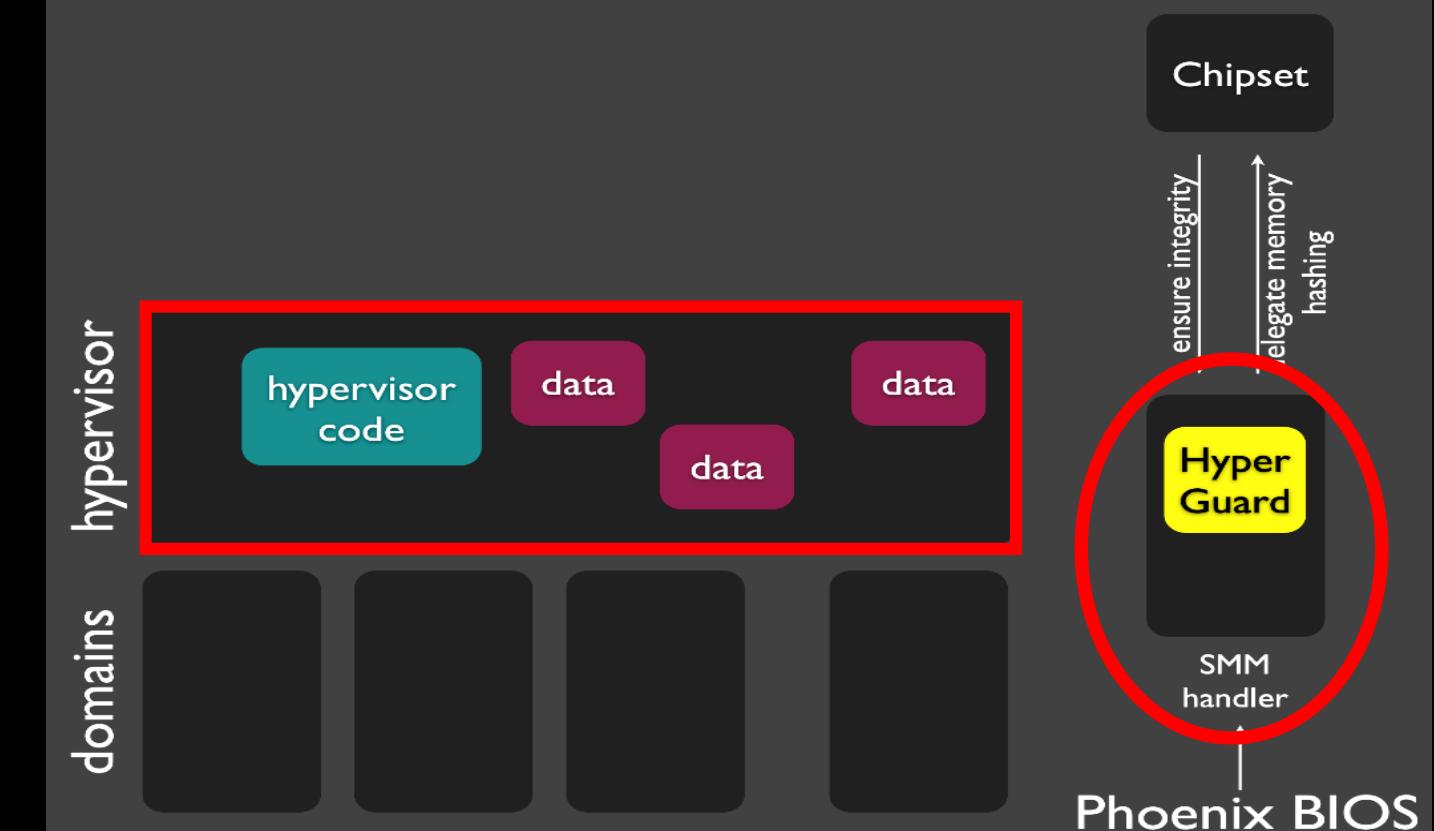


Figure 3: HypSec on KVM/ARM

Firmware Modification

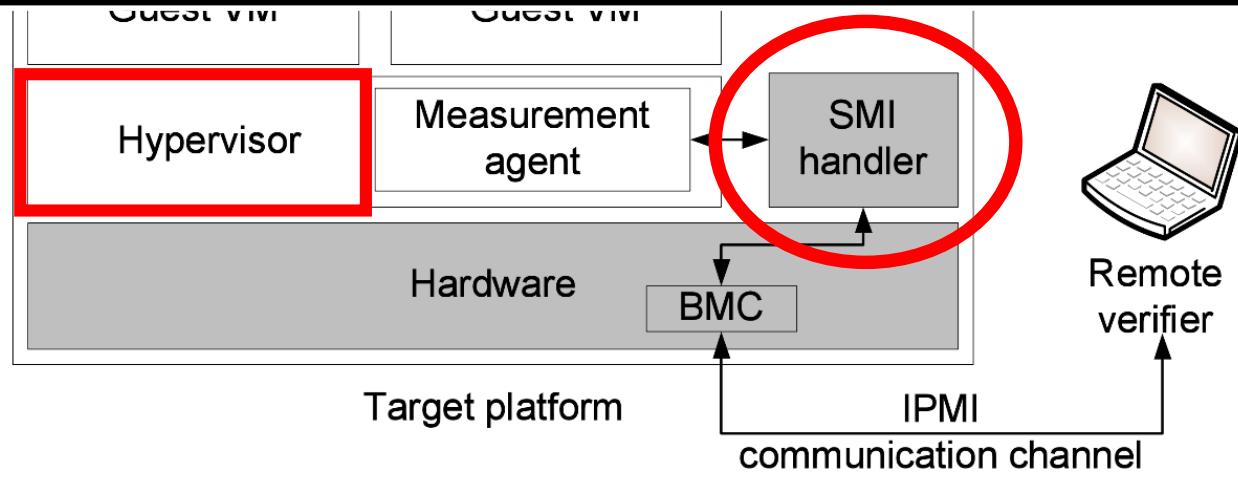
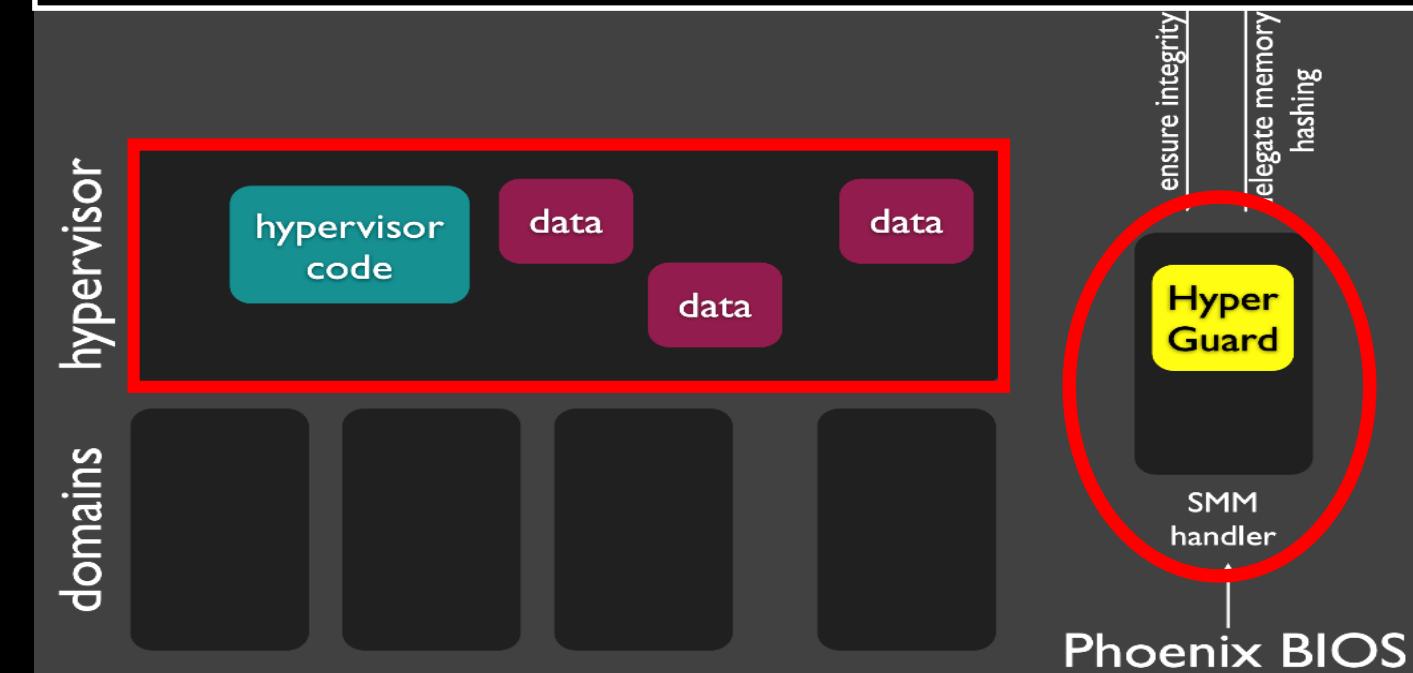
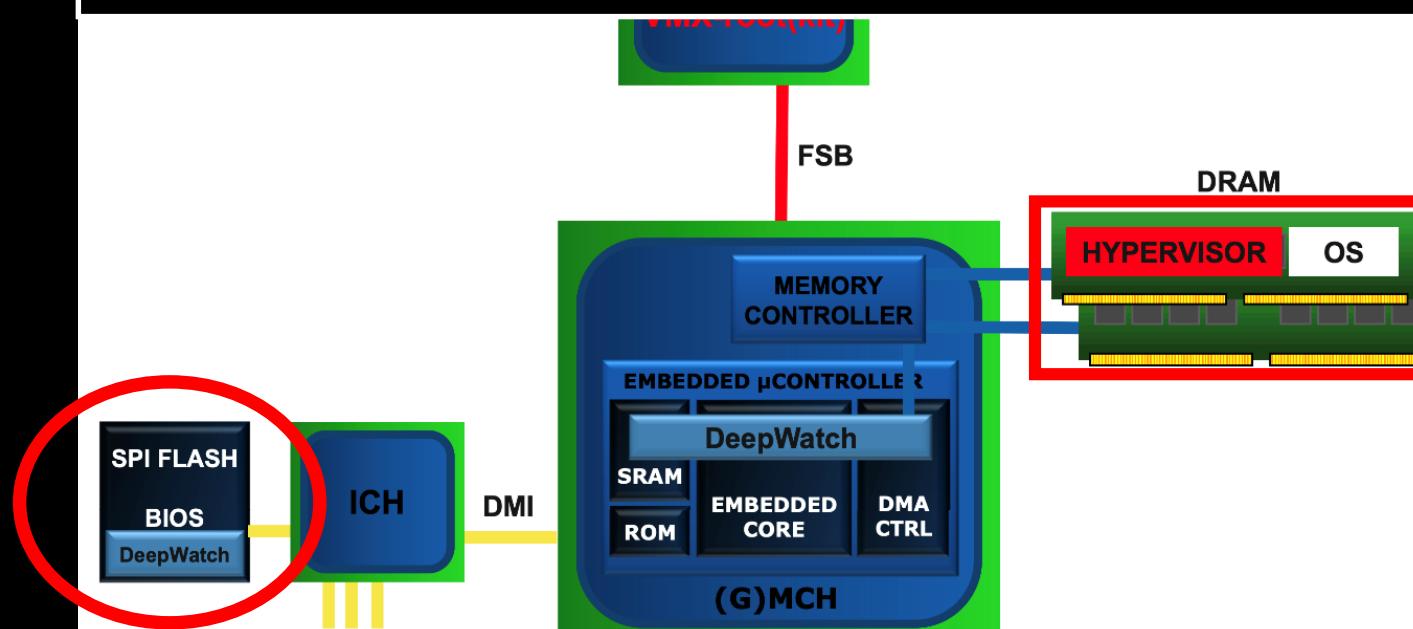


Figure 1: HyperSentry's architecture (Trusted components are shown in gray.)

Firmware Modification



Firmware Modification



Hypervisor Retrofit

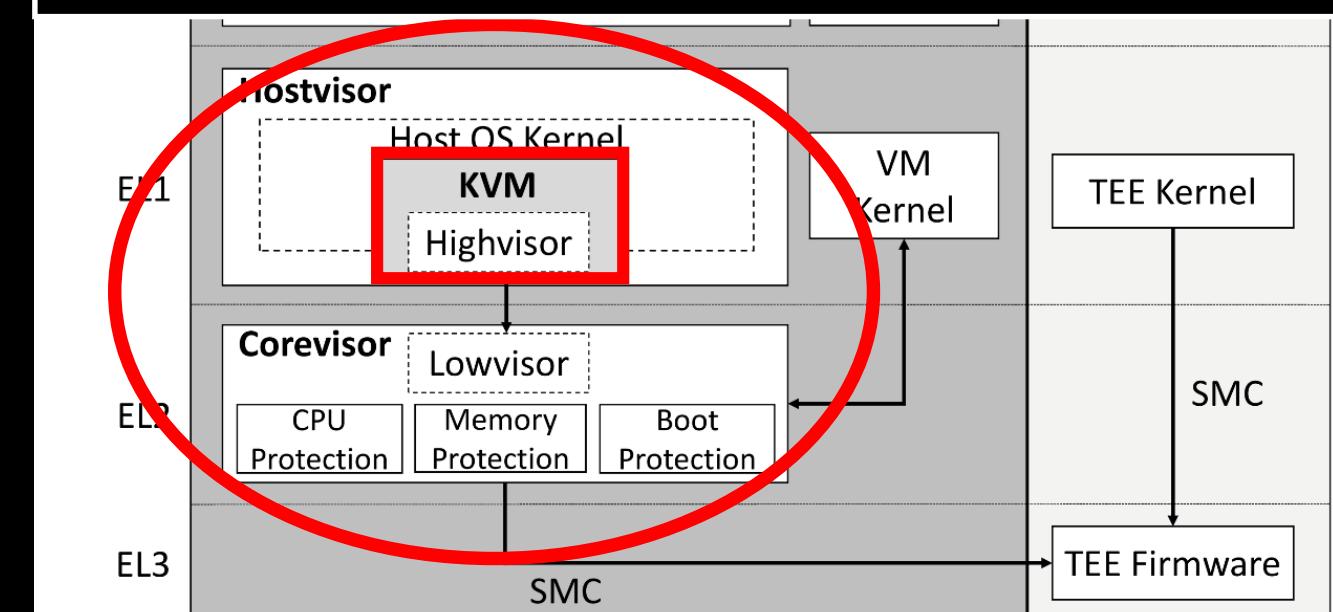


Figure 3: HypSec on KVM/ARM

Firmware Modification

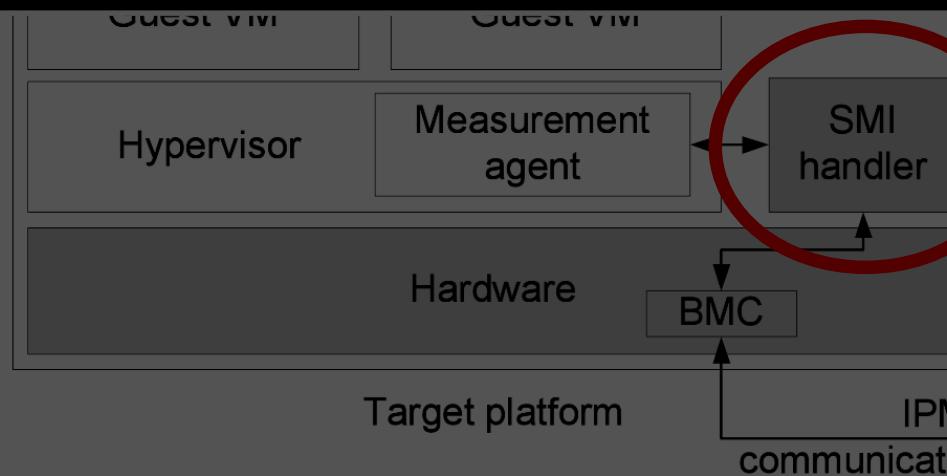
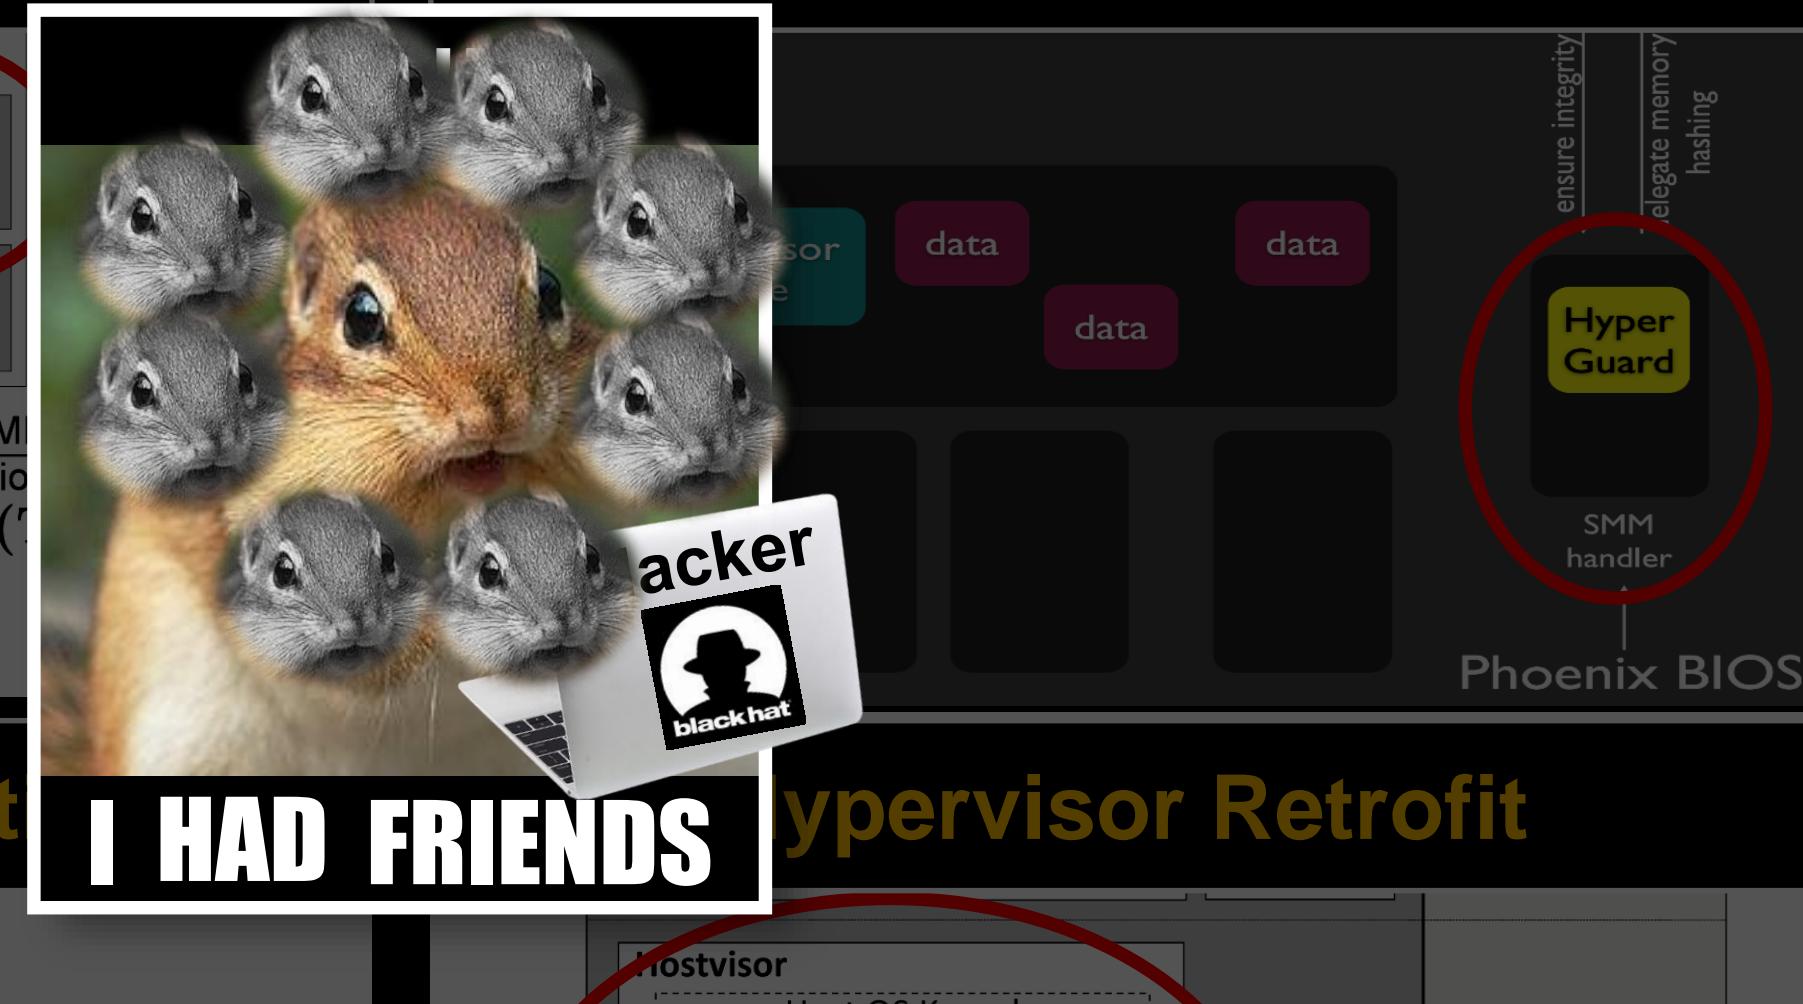


Figure 1: HyperSentry's architecture (The components shown in gray are the target platform components.)

Firmware Modification



I need a more **PRACTICAL** solution!

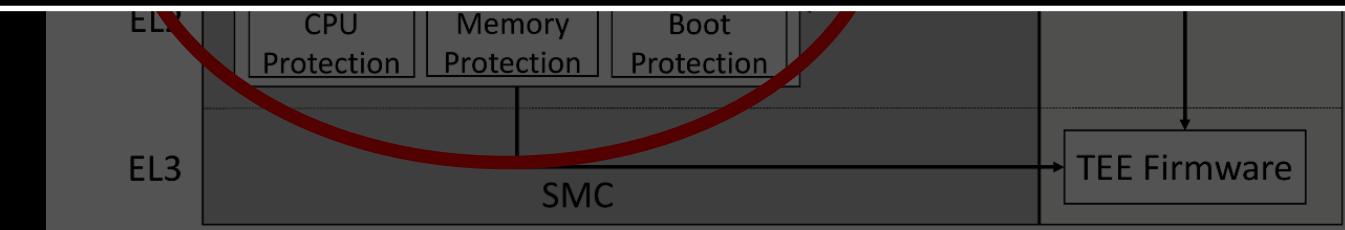
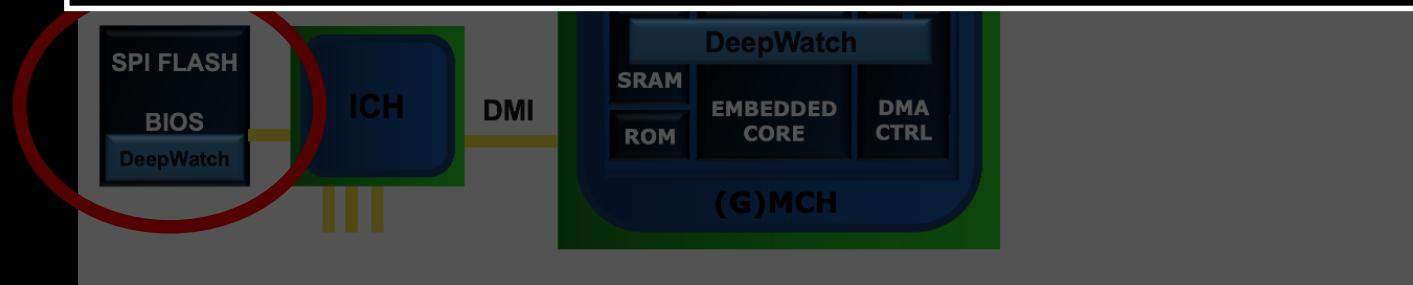
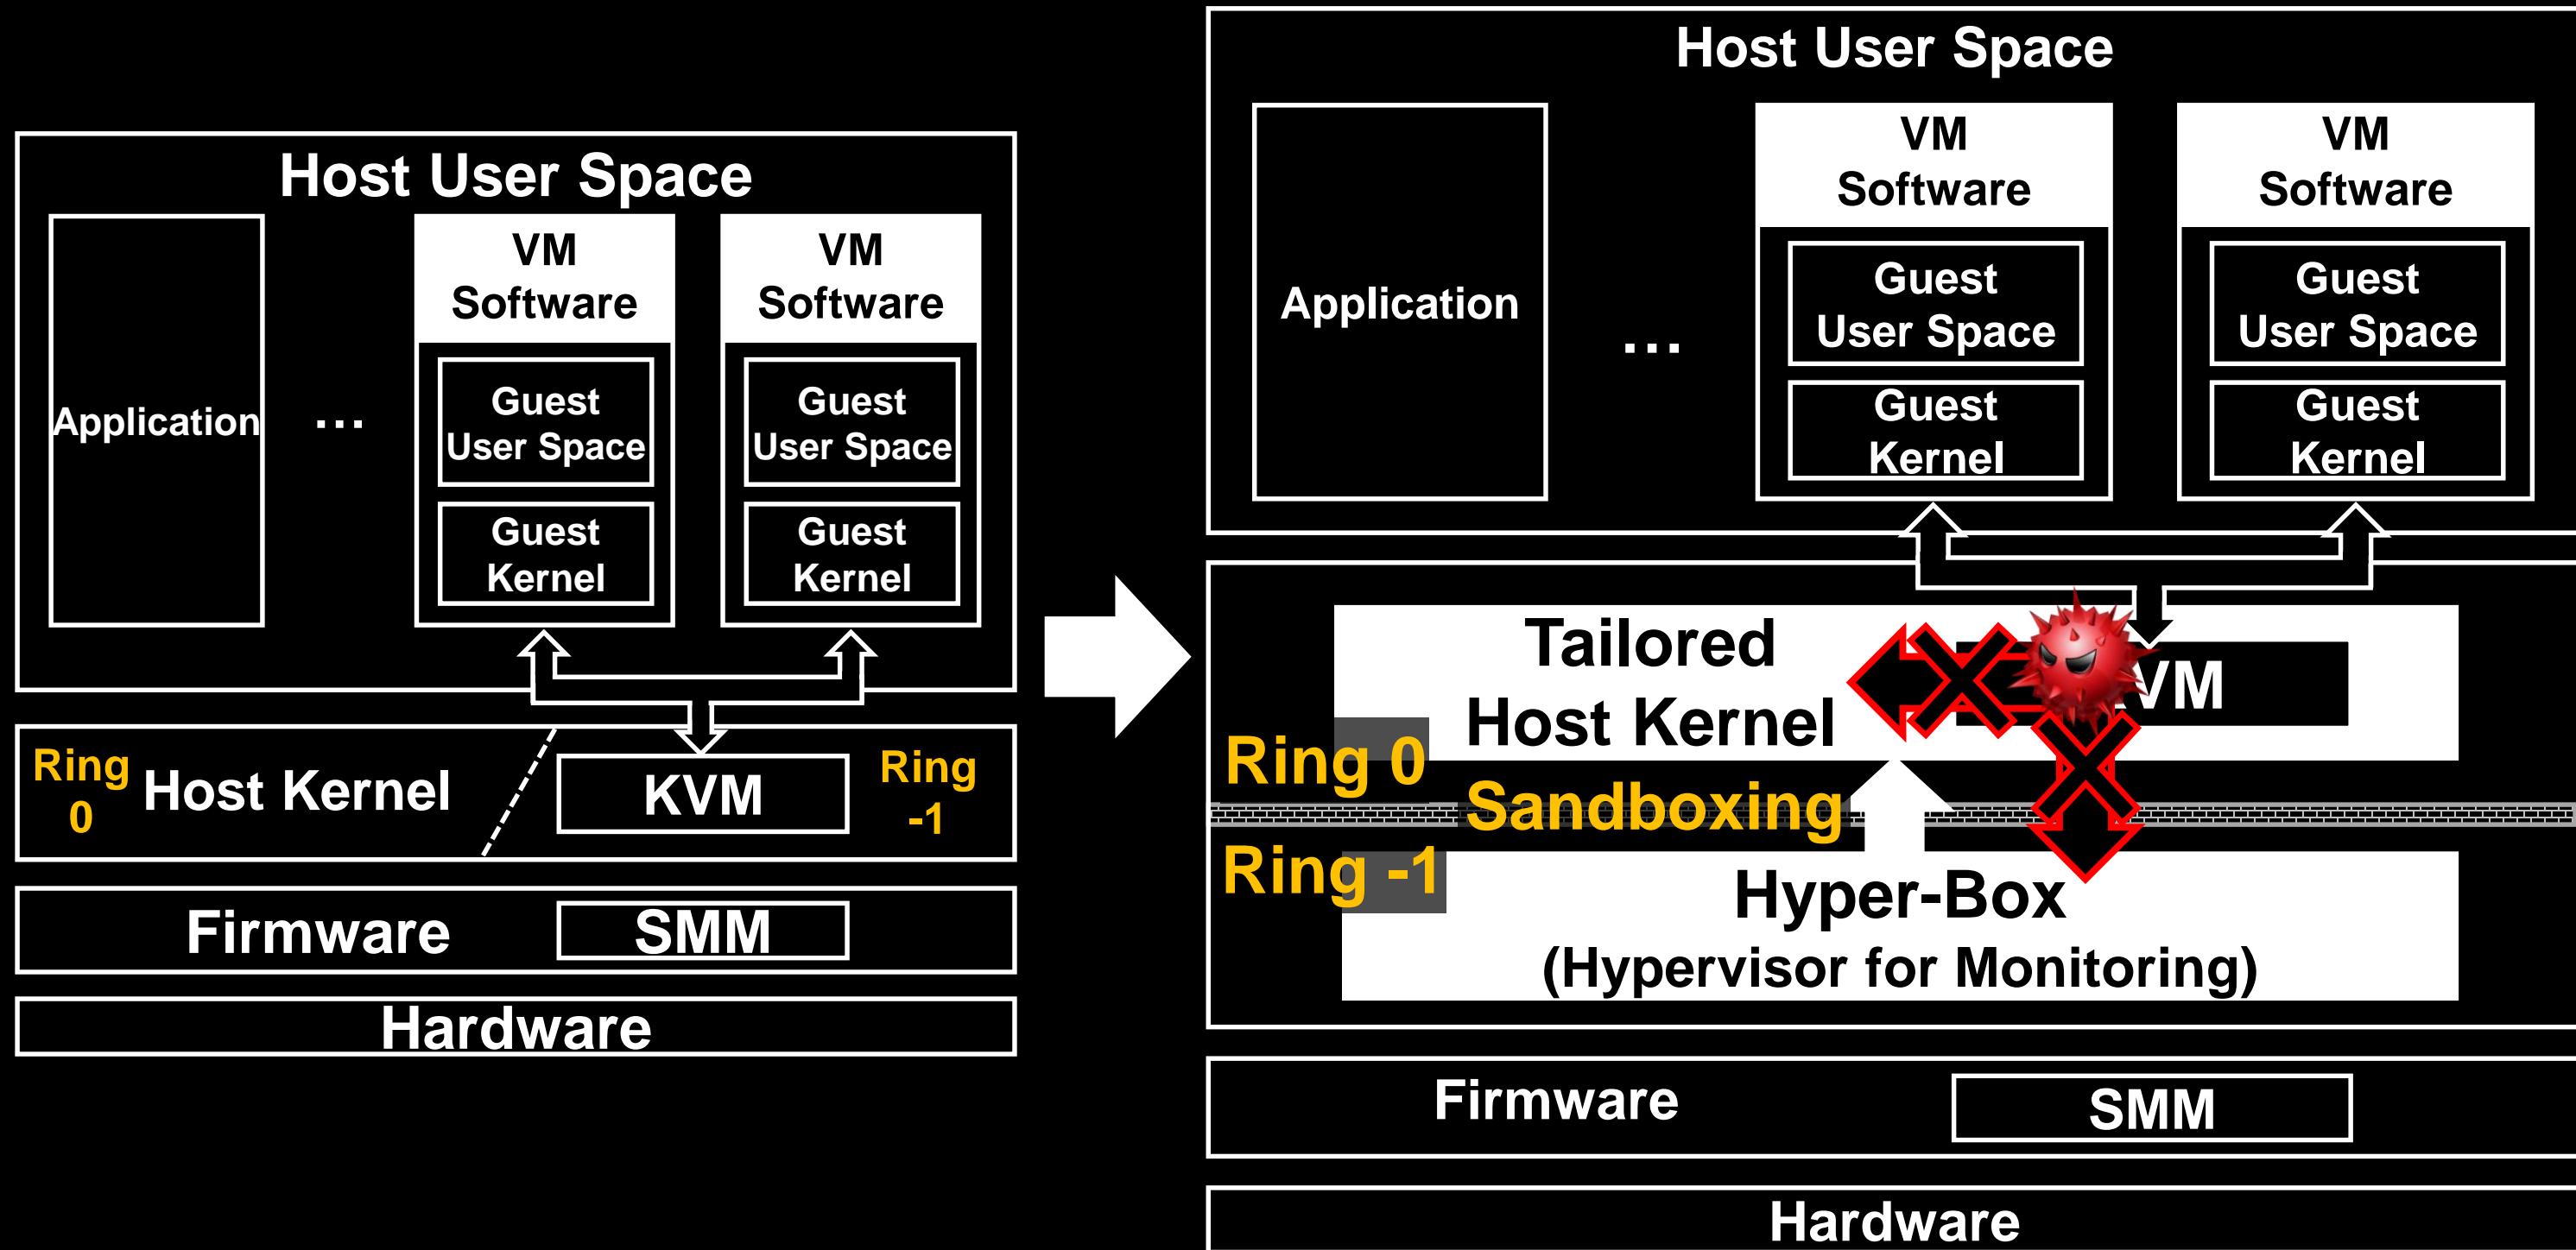
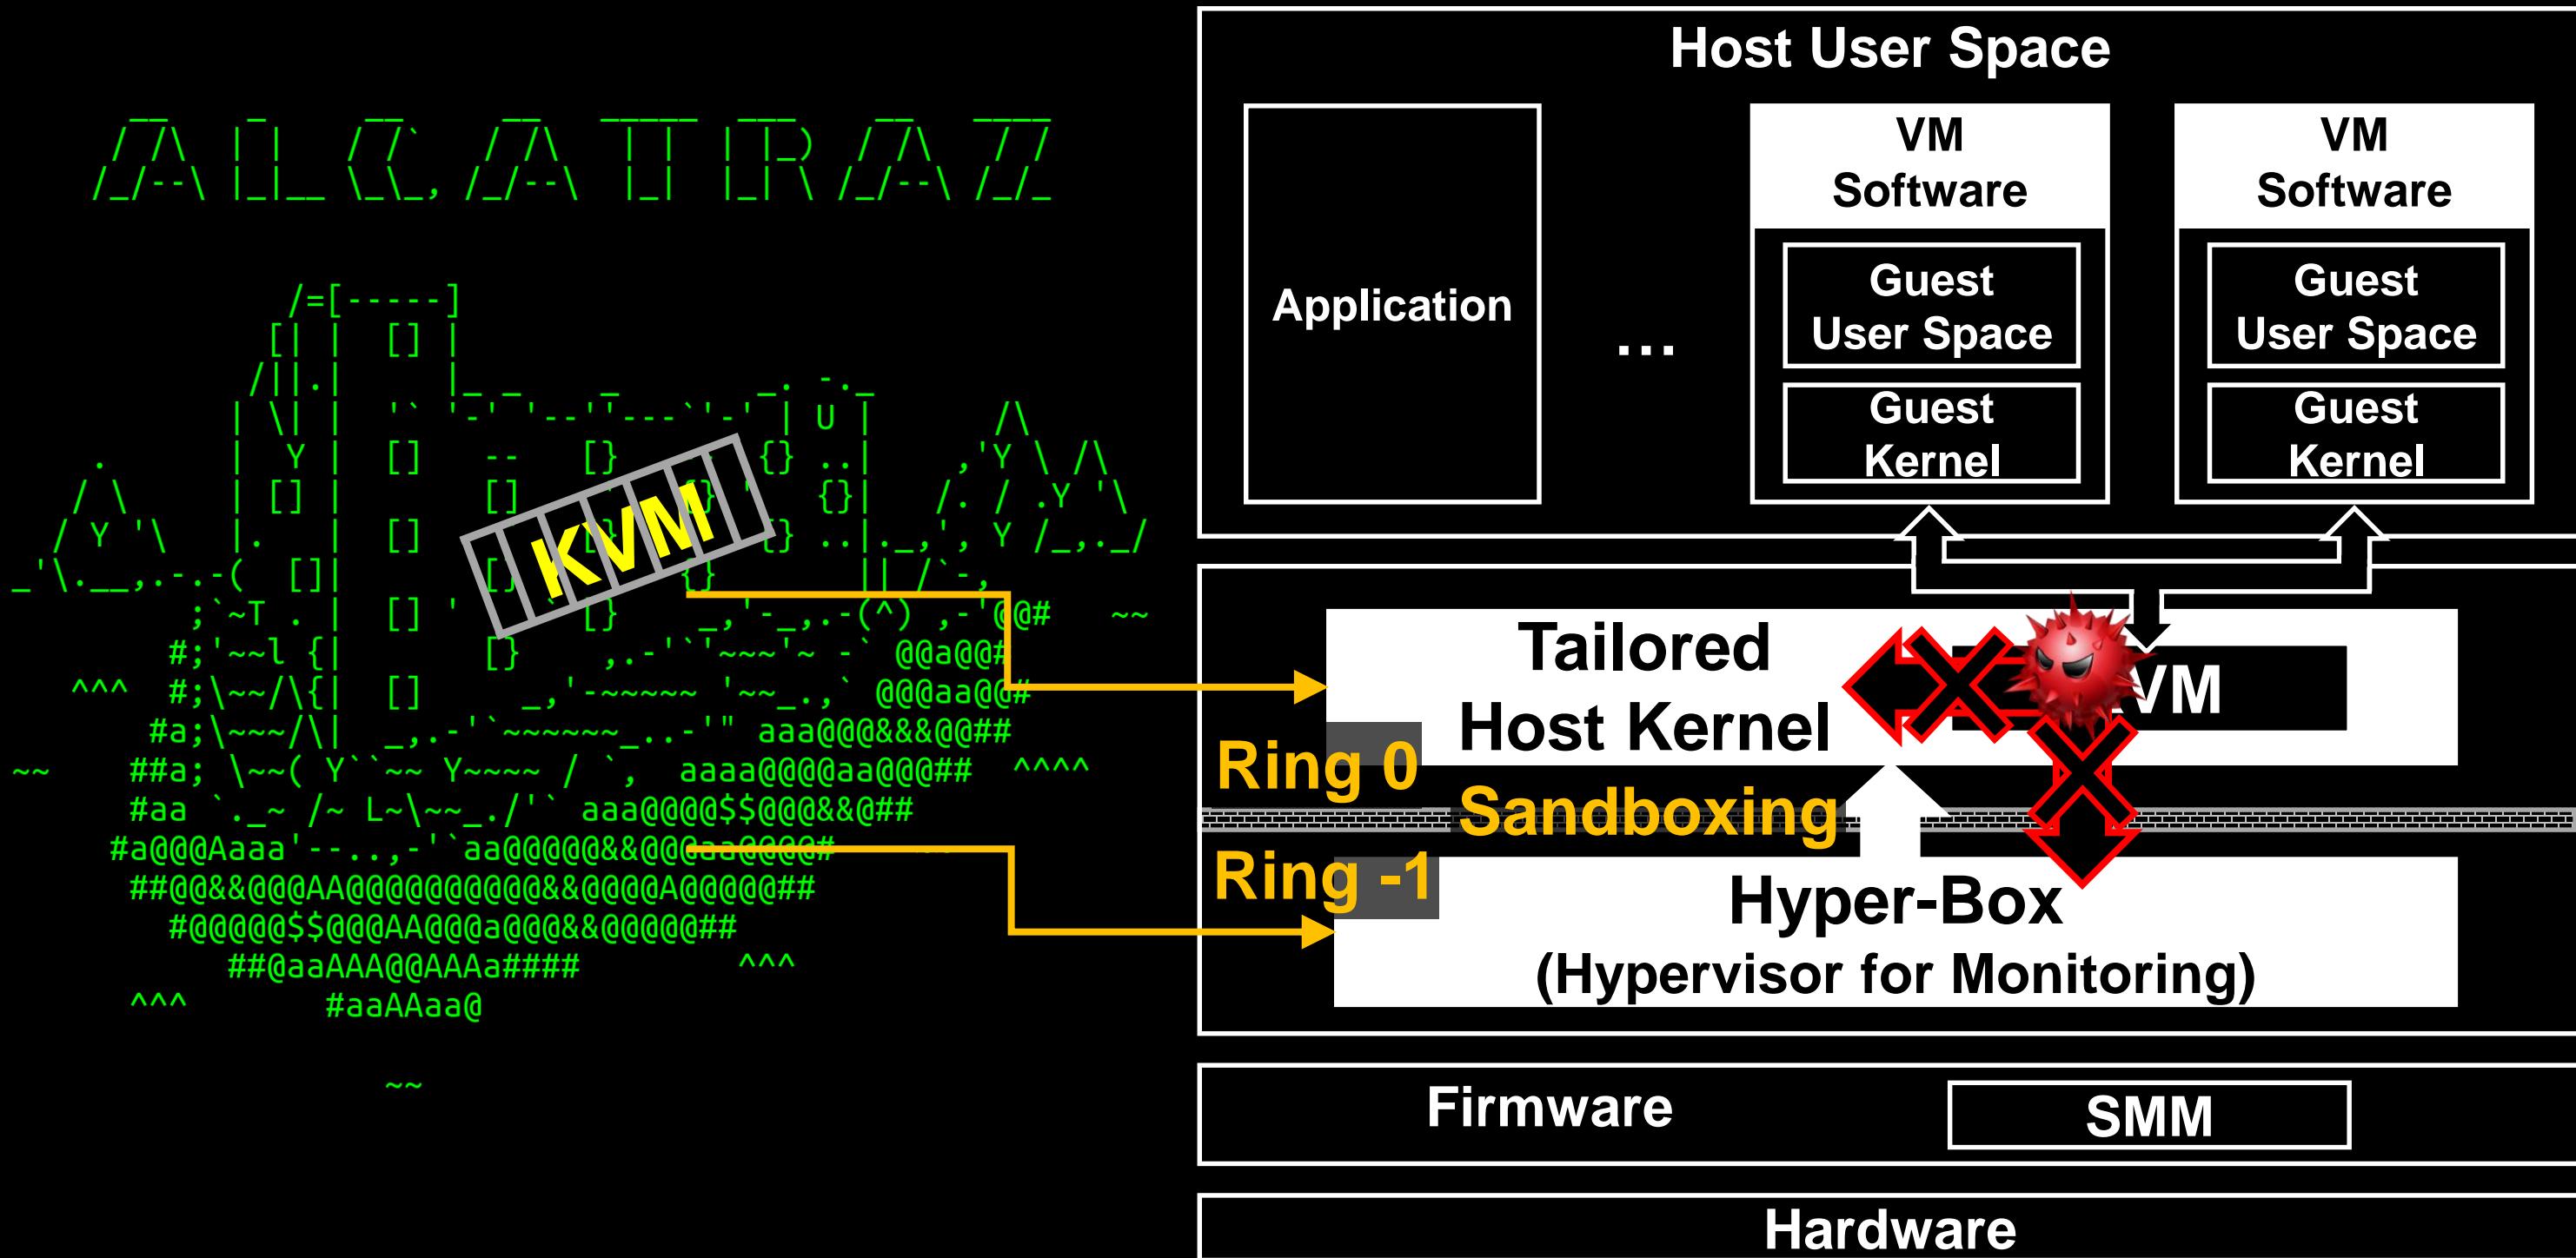


Figure 3: HypSec on KVM/ARM

The Solution: Sandboxing Hypervisor



The Solution: Sandboxing Hypervisor

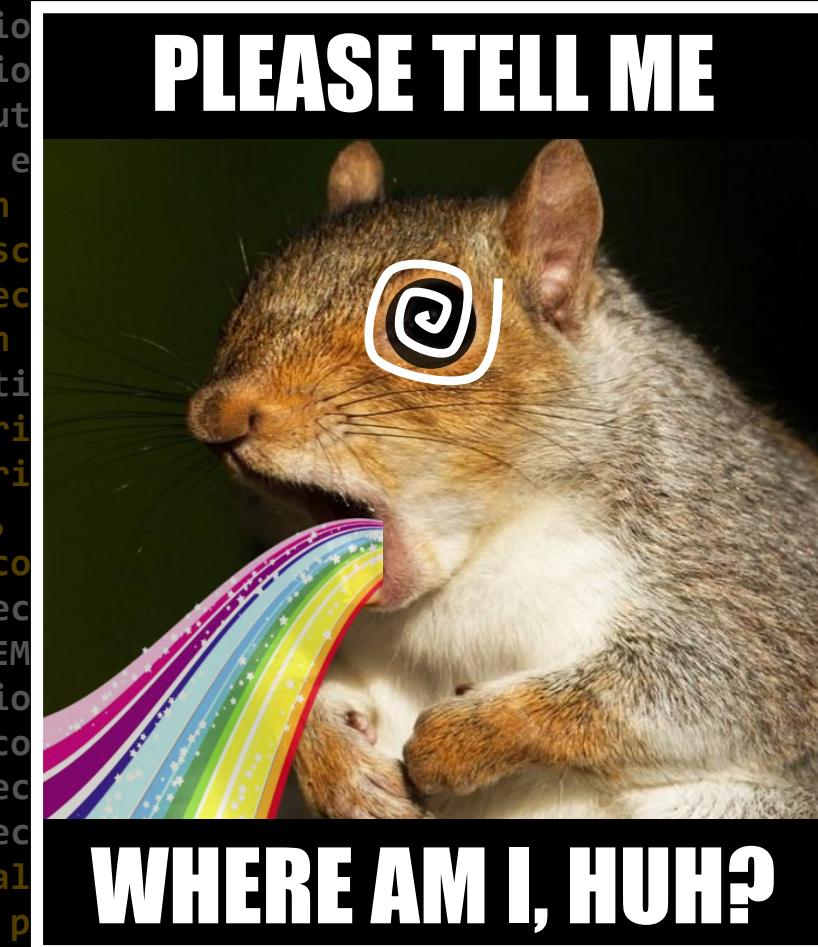


- Background
- Analysis of Critical Paths of Escapes 
- Design and Implementation of Alcatraz
- Evaluation and Demo
- Conclusion and Black Hat Sound Bytes

CVE-2021-22543: Improper page handling of KVM, privilege escalation
CVE-2021-20181: TOCTOU in 9pfs of QEMU, privilege escalation
CVE-2021- 3546: OOB in VIRTIO in QEMU, code execution
CVE-2021- 3409: OOB in SDC of QEMU, code execution
CVE-2020-35506: UAF in SCSI of QEMU, code execution
CVE-2020-17380: OOB in SDC of QEMU, code execution
CVE-2020-15863: OOB in NET of QEMU, code execution
CVE-2020-14364: OOB in USB of QEMU, code execution
CVE-2020- 1711: OOB in iSCSI of QEMU, code execution
CVE-2019-18389: OOB in VIRTIO-GPU in QEMU, code execution
CVE-2019-14835: OOB in KVM, privilege escalation
CVE-2019-14821: OOB in MMIO of KVM, privilege escalation
CVE-2019- 7221: UAF in HR timer of KVM, code execution
CVE-2018-16882: UAF in KVM, privilege escalation
CVE-2018-16847: OOB in NVME of QEMU, code execution
CVE-2018-12904: Improper VMX handling of KVM, privilege escalation
CVE-2018-10901: Improper GDT handling of KVM, privilege escalation
CVE-2018- 1087: Improper stack switching of KVM, privilege escalation
CVE-2017-12188: Improper page handling of KVM, code execution
CVE-2017- 7980: OOB in display of QEMU, code execution
CVE-2017- 5931: Integer overflow in VIRTIO of QEMU, code execution
CVE-2017- 5667: OOB in SDC of QEMU, code execution
CVE-2017- 2630: Stack overflow in NBD of QEMU, code execution
CVE-2017- 2620: OOB in display of QEMU, code execution
CVE-2017- 2615: OOB in display of QEMU, code execution
CVE-2016-10150: UAF in KVM, local privilege escalation
CVE-2016- 9777: Improper IOAPIC handing in KVM, privilege escalation
CVE-2016- 9603: OOB in display of QEMU, code execution
CVE-2016- 7161: OOB in NET of QEMU, code execution
CVE-2016- 6351: OOB in SCSI of QEMU, code execution
CVE-2016- 5338: OOB in SCSI of QEMU, code execution
CVE-2016- 5126: OOB in iSCSI of QEMU, code execution
CVD-2016- 4440: Improper APICv handling of KVM, code execution
CVE-2016- 4439: OOB in SCSI of QEMU, code execution
CVE-2016- 4002: OOB in NET of QEMU, code execution
CVE-2016- 3710: OOB in VESA VGA of QEMU, code execution
CVE-2016- 1714: OOB in NVRAM of QEMU, code execution
CVE-2016- 1586: OOB in IDE of QEMU, code execution

CVE-2016- 0749: OOB in SPICE of QEMU, code execution
CVE-2015- 7512: OOB in NET of QEMU, code execution
CVE-2015- 7504: OOB in NET of QEMU, code execution
CVE-2015- 5279: OOB in NET of QEMU, code execution
CVE-2015- 5278: OOB in NET of QEMU, code execution
CVE-2015- 5260: OOB in SPICE of QEMU, code execution
CVE-2015- 5225: OOB in VNC of QEMU, code execution
CVE-2015- 5154: OOB in IDE of QEMU, code execution
CVE-2015- 3456: OOB in FDC of QEMU (VENOM), code execution
CVE-2015- 3214: OOB in PIT of QEMU, code execution
CVE-2015- 3209: OOB in NET of QEMU, code execution
CVE-2014- 8106: OOB in display of QEMU, code execution
CVE-2014- 7840: OOB in migration function of QEMU, code execution
CVE-2014- 3461: OOB in USB of QEMU, code execution
CVE-2014- 0182: OOB in VIRTIO of QEMU, code execution
CVE-2014- 0150: Integer overflow in VIRTIO-NET of QEMU, code execution
CVE-2014- 0144: OOB in CLOOP of QEMU, code execution
CVE-2014- 0049: OOB in KVM, code execution
CVE-2013- 6399: OOB in VIRTIO of QEMU, code execution
CVE-2013- 4544: OOB in NET of QEMU, code execution
CVE-2013- 4587: OOB in KVM, local privilege escalation
CVE-2013- 4542: OOB in SCSI of QEMU, code execution
CVE-2013- 4541: OOB in USB of QEMU, code execution
CVE-2013- 4540: OOB in GPIO of QEMU, code execution
CVE-2013- 4539: OOB in Audio of QEMU, code execution
CVE-2013- 4538: OOB in Display of QEMU, code execution
CVE-2013- 4537: OOB in SDC of QEMU, code execution
CVE-2013- 4536: OOB in PIC of QEMU, code execution
CVE-2013- 4532: OOB in NET of QEMU, code execution
CVE-2013- 4529: OOB in PCI of QEMU, code execution
CVE-2013- 4527: OOB in HPET of QEMU, code execution
CVE-2013- 4526: OOB in AHCI of QEMU, code execution
CVE-2013- 4151: OOB in VIRTIO of QEMU, code execution
CVE-2013- 4150: OOB in VIRTIO-NET of QEMU, code execution
CVE-2013- 4149: OOB in VIRTIO-NET of QEMU, code execution
CVE-2013- 4148: OOB in VIRTIO-NET of QEMU, code execution
CVE-2013- 1943: Improper page handling of KVM, local privilege escalation
CVE-2013- 0311: Improper descriptor handling of KVM, privilege escalation

CVE-2021-22543: Improper page handling of KVM, privilege escalation
CVE-2021-20181: TOCTOU in 9pfs of QEMU, privilege escalation
CVE-2021- 3546: OOB in VIRTIO in QEMU, code execution
CVE-2021- 3409: OOB in SDC of QEMU, code execution
CVE-2020-35506: UAF in SCSI of QEMU, code execution
CVE-2020-17380: OOB in SDC of QEMU, code execution
CVE-2020-15863: OOB in NET of QEMU, code execution
CVE-2020-14364: OOB in USB of QEMU, code execution
CVE-2020- 1711: OOB in iSCSI of QEMU, code execution
CVE-2019-18389: OOB in VIRTIO-GPU in QEMU, code execution
CVE-2019-14835: OOB in KVM, privilege escalation
CVE-2019-14821: OOB in MMIO of KVM, privilege escalation
CVE-2019- 7221: UAF in HR timer of KVM, code execution
CVE-2018-16882: UAF in KVM, privilege escalation
CVE-2018-16847: OOB in NVME of QEMU, code execution
CVE-2018-12904: Improper VMX handling of KVM, privilege escalation
CVE-2018-10901: Improper GDT handling of KVM, privilege escalation
CVE-2018- 1087: Improper stack switching of KVM, privilege escalation
CVE-2017-12188: Improper page handling of KVM, code execution
CVE-2017- 7980: OOB in display of QEMU, code execution
CVE-2017- 5931: Integer overflow in VIRTIO of QEMU, code execution
CVE-2017- 5667: OOB in SDC of QEMU, code execution
CVE-2017- 2630: Stack overflow in NBD of QEMU, code execution
CVE-2017- 2620: OOB in display of QEMU, code execution
CVE-2017- 2615: OOB in display of QEMU, code execution
CVE-2016-10150: UAF in KVM, local privilege escalation
CVE-2016- 9777: Improper IOAPIC handing in KVM, privilege escalation
CVE-2016- 9603: OOB in display of QEMU, code execution
CVE-2016- 7161: OOB in NET of QEMU, code execution
CVE-2016- 6351: OOB in SCSI of QEMU, code execution
CVE-2016- 5338: OOB in SCSI of QEMU, code execution
CVE-2016- 5126: OOB in iSCSI of QEMU, code execution
CVD-2016- 4440: Improper APICv handling of KVM, code execution
CVE-2016- 4439: OOB in SCSI of QEMU, code execution
CVE-2016- 4002: OOB in NET of QEMU, code execution
CVE-2016- 3710: OOB in VESA VGA of QEMU, code execution
CVE-2016- 1714: OOB in NVRAM of QEMU, code execution
CVE-2016- 1586: OOB in IDE of QEMU, code execution



PLEASE TELL ME

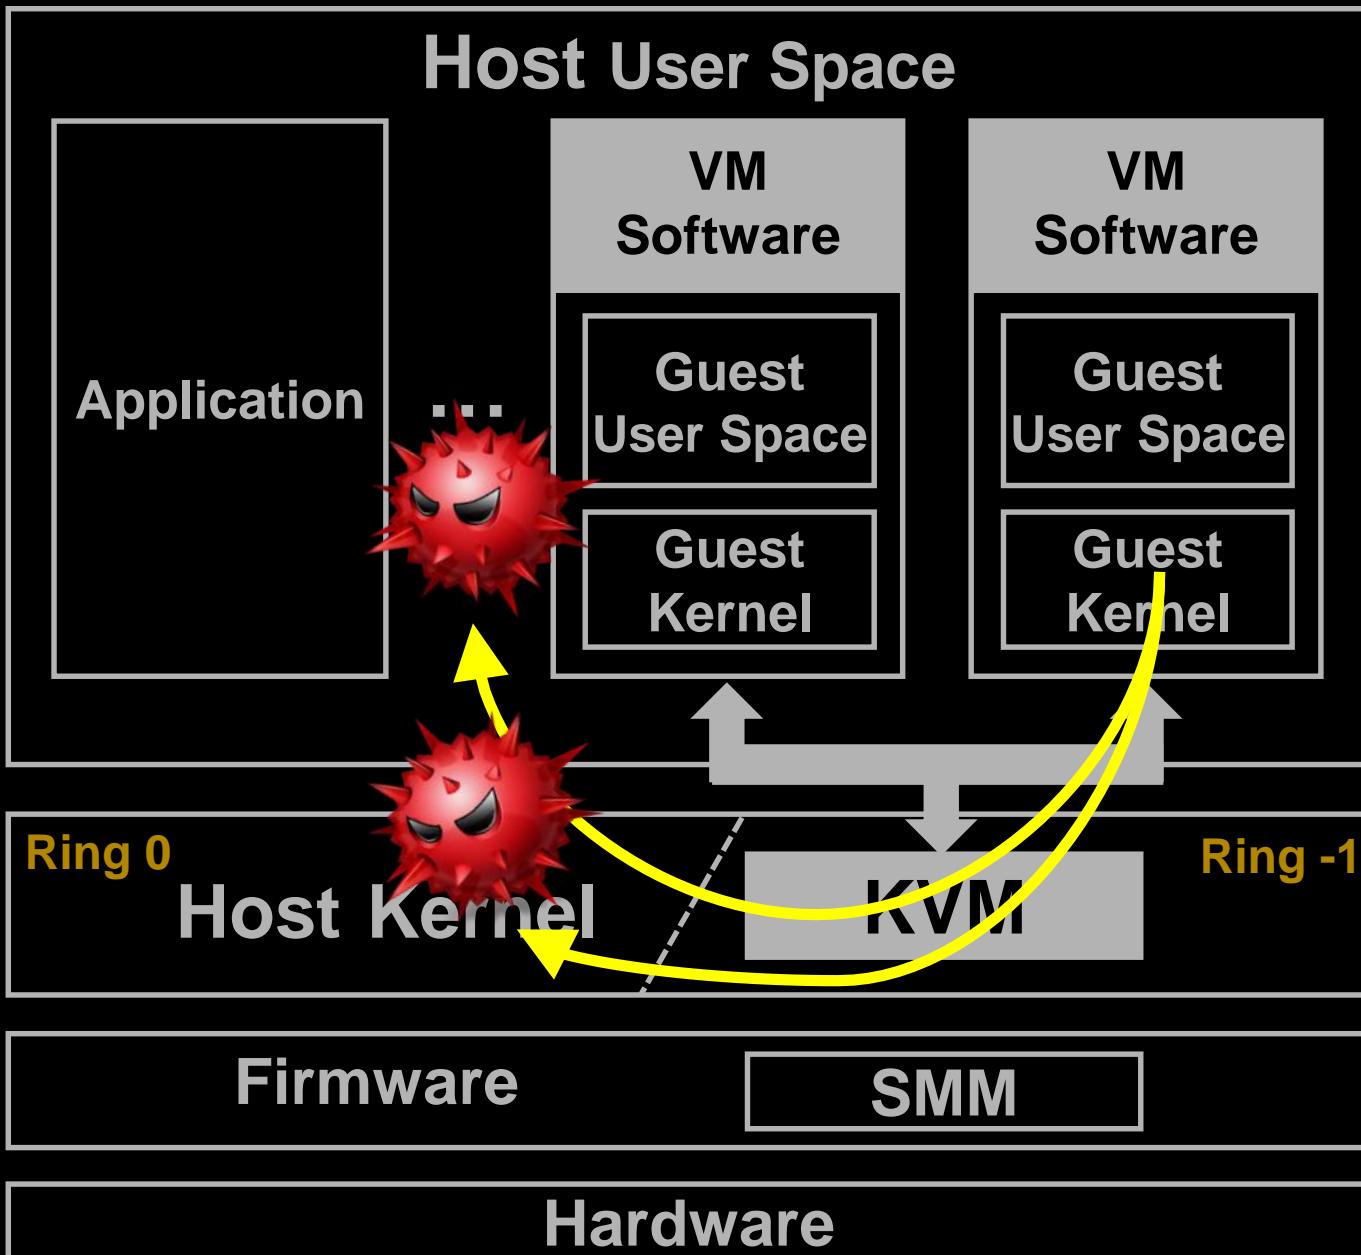
WHERE AM I, HUH?

CVE-2016- 0749: OOB in SPICE of QEMU, code execution
CVE-2015- 7512: OOB in NET of QEMU, code execution
CVE-2015- 7504: OOB in NET of QEMU, code execution
CVE-2015- 5279: OOB in NET of QEMU, code execution
CVE-2015- 5278: OOB in NET of QEMU, code execution
CVE-2015- 5260: OOB in SPICE of QEMU, code execution
OOB in VNC of QEMU, code execution
OOB in IDE of QEMU, code execution
OOB in FDC of QEMU (VENOM), code execution
OOB in PIT of QEMU, code execution
OOB in NET of QEMU, code execution
OOB in display of QEMU, code execution
OOB in migration function of QEMU, code execution
OOB in USB of QEMU, code execution
OOB in VIRTIO of QEMU, code execution
Integer overflow in VIRTIO-NET of QEMU, code execution
OOB in CLOOP of QEMU, code execution
OOB in KVM, code execution
OOB in VIRTIO of QEMU, code execution
OOB in NET of QEMU, code execution
OOB in KVM, local privilege escalation
OOB in SCSI of QEMU, code execution
OOB in USB of QEMU, code execution
OOB in GPIO of QEMU, code execution
OOB in Audio of QEMU, code execution
OOB in Display of QEMU, code execution
OOB in SDC of QEMU, code execution
CVE-2013- 4536: OOB in PIC of QEMU, code execution
CVE-2013- 4532: OOB in NET of QEMU, code execution
CVE-2013- 4529: OOB in PCI of QEMU, code execution
CVE-2013- 4527: OOB in HPET of QEMU, code execution
CVE-2013- 4526: OOB in AHCI of QEMU, code execution
CVE-2013- 4151: OOB in VIRTIO of QEMU, code execution
CVE-2013- 4150: OOB in VIRTIO-NET of QEMU, code execution
CVE-2013- 4149: OOB in VIRTIO-NET of QEMU, code execution
CVE-2013- 4148: OOB in VIRTIO-NET of QEMU, code execution
CVE-2013- 1943: Improper page handling of KVM, local privilege escalation
CVE-2013- 0311: Improper descriptor handling of KVM, privilege escalation

Extracting Escape Paths from CVEs (1)

Code Execution with a KVM Vulnerability

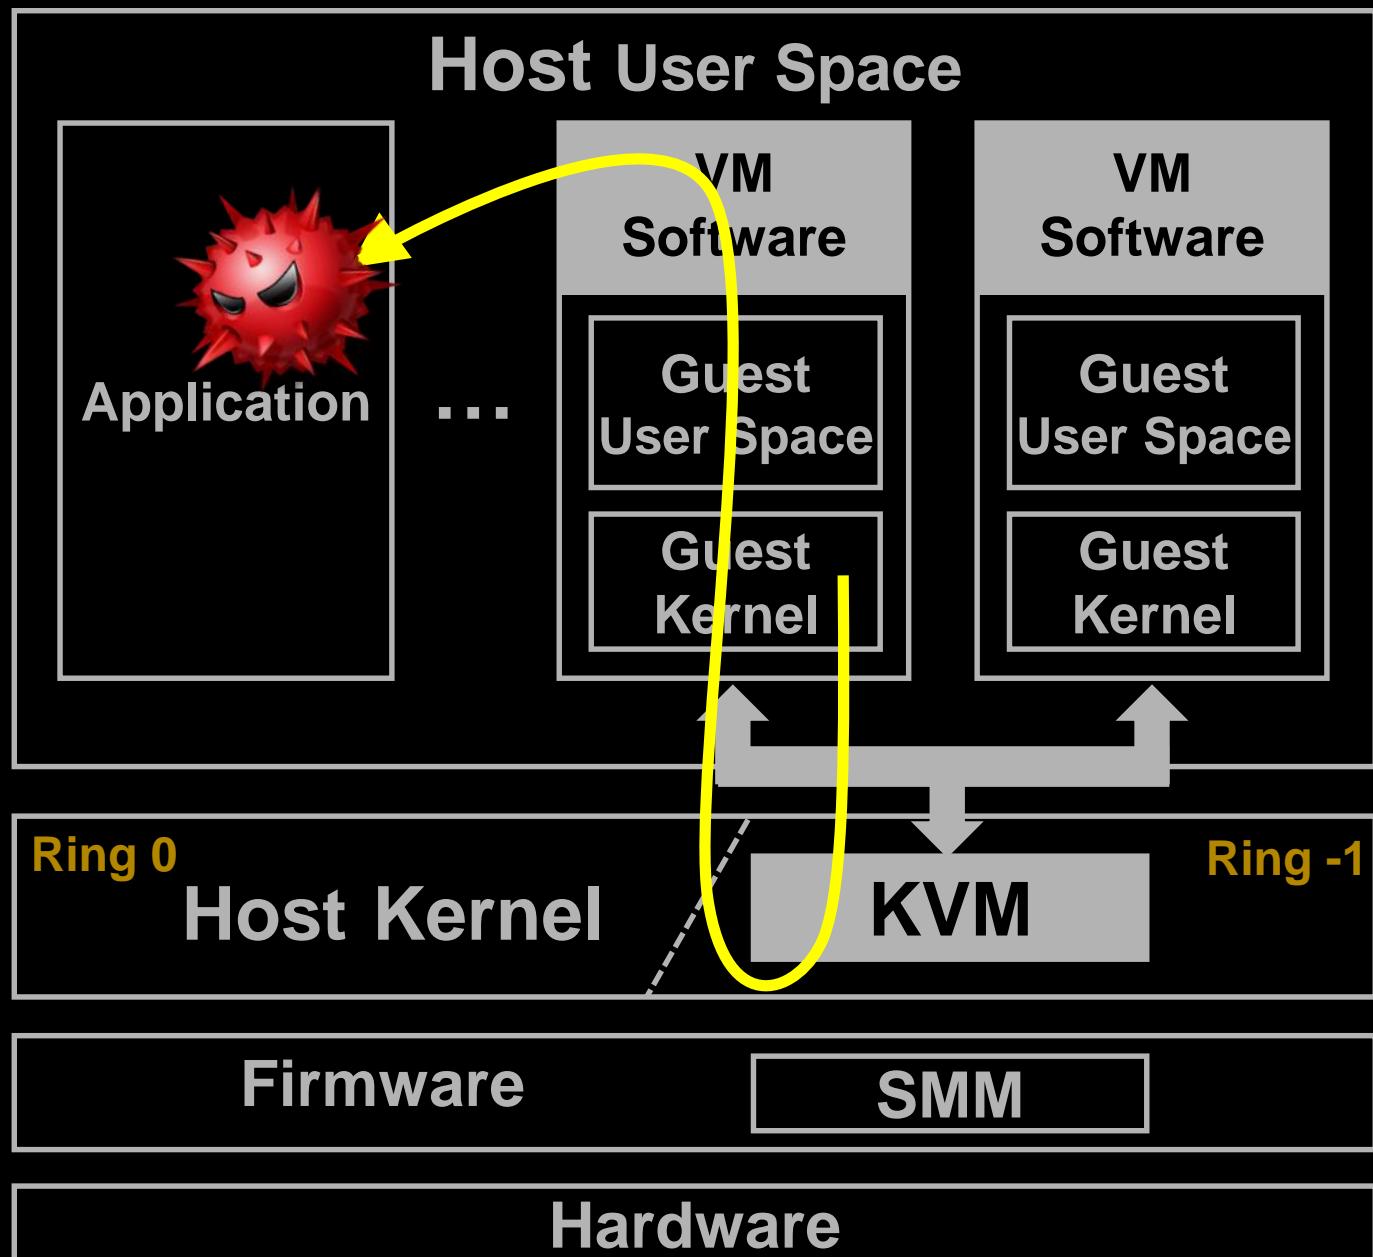
- Executes code outside of a VM
 - Small shellcode is used for **creating a ROOT privilege process of a host**
 - Process creation: Guest → Host
 - Privilege escalation: Host user → root
 - It is also possible to modify kernel or KVM code
 - But, disabling security features are needed! 😞



Extracting Escape Paths from CVEs (2)

Code Execution with a VM Vulnerability

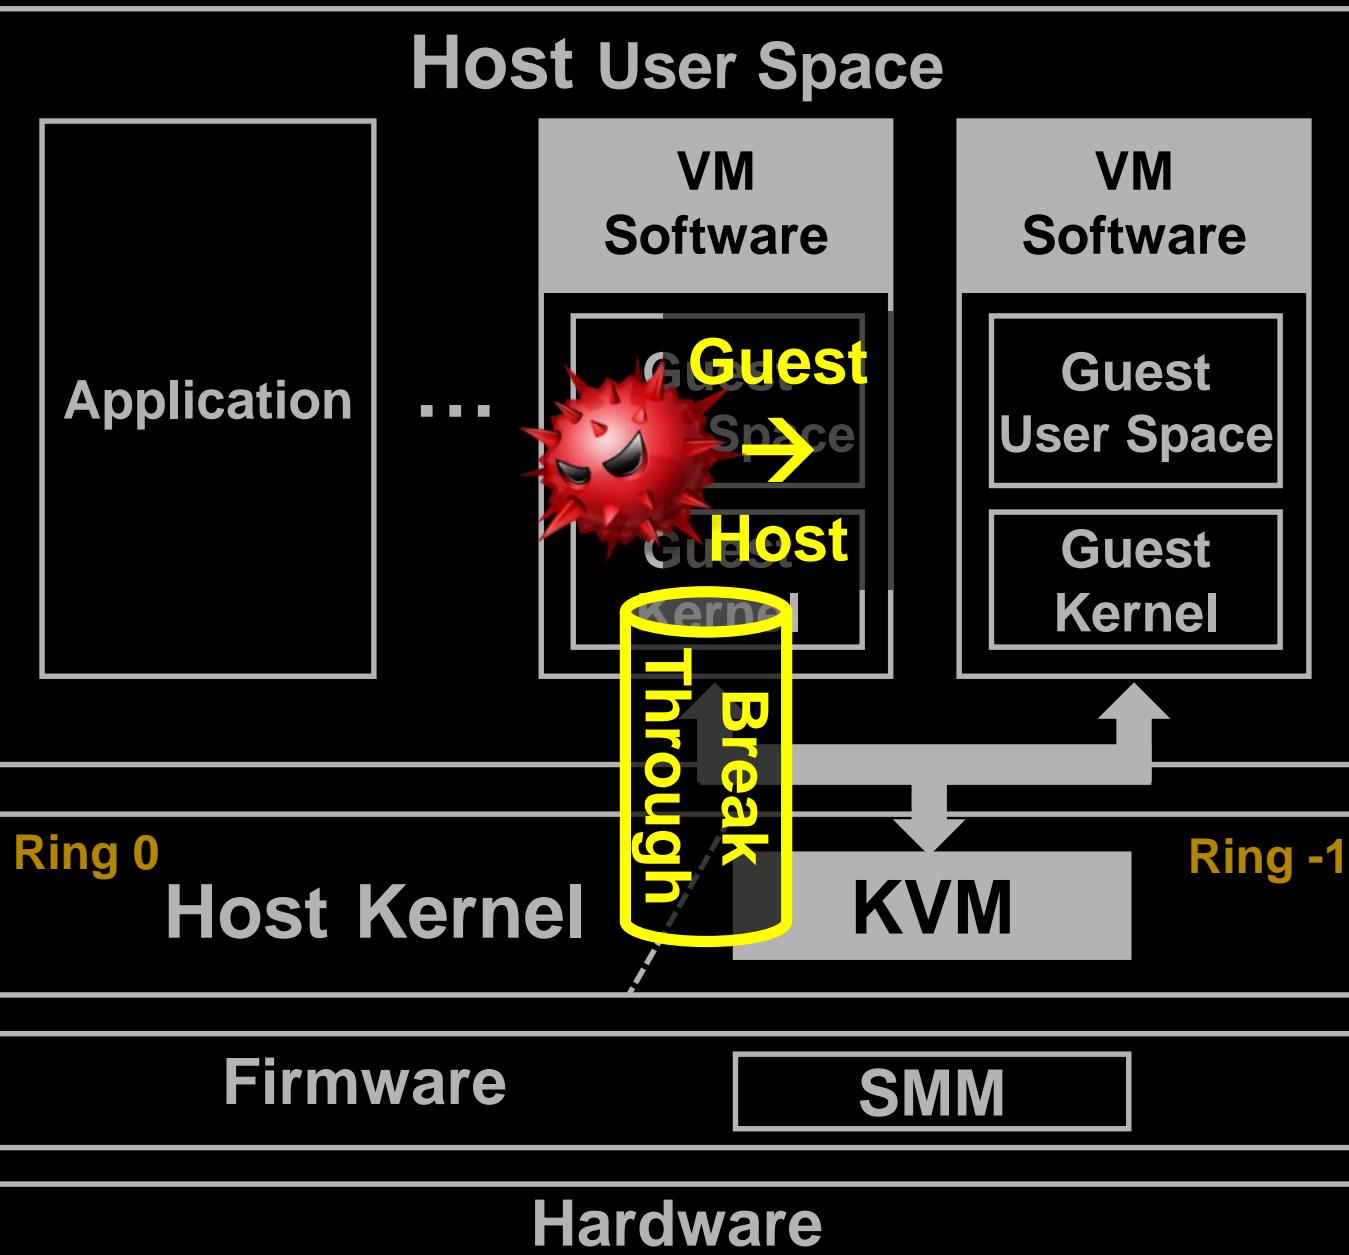
- Also executes code outside of a VM
 - A VM software has lots of code for services and hardware emulations
 - Full-featured VMs (QEMU) has a large code base
 - Others have smaller ones, but ... 😞
 - Small shellcode is used for creating a USER privilege process of a host
 - Process creation: Guest → Host
 - Privilege escalation is needed!



Extracting Escape Paths from CVEs (3)

Privilege Escalation with a KVM Vulnerability

- Elevates the privilege of a VM
 - A VM is a process and runs in the guest privilege of a host
 - But, improper handlings give HOST privilege to the VM
 - Page tables, interrupts, registers, etc.
 - It is possible to create a root privilege process of a host and modify kernel code

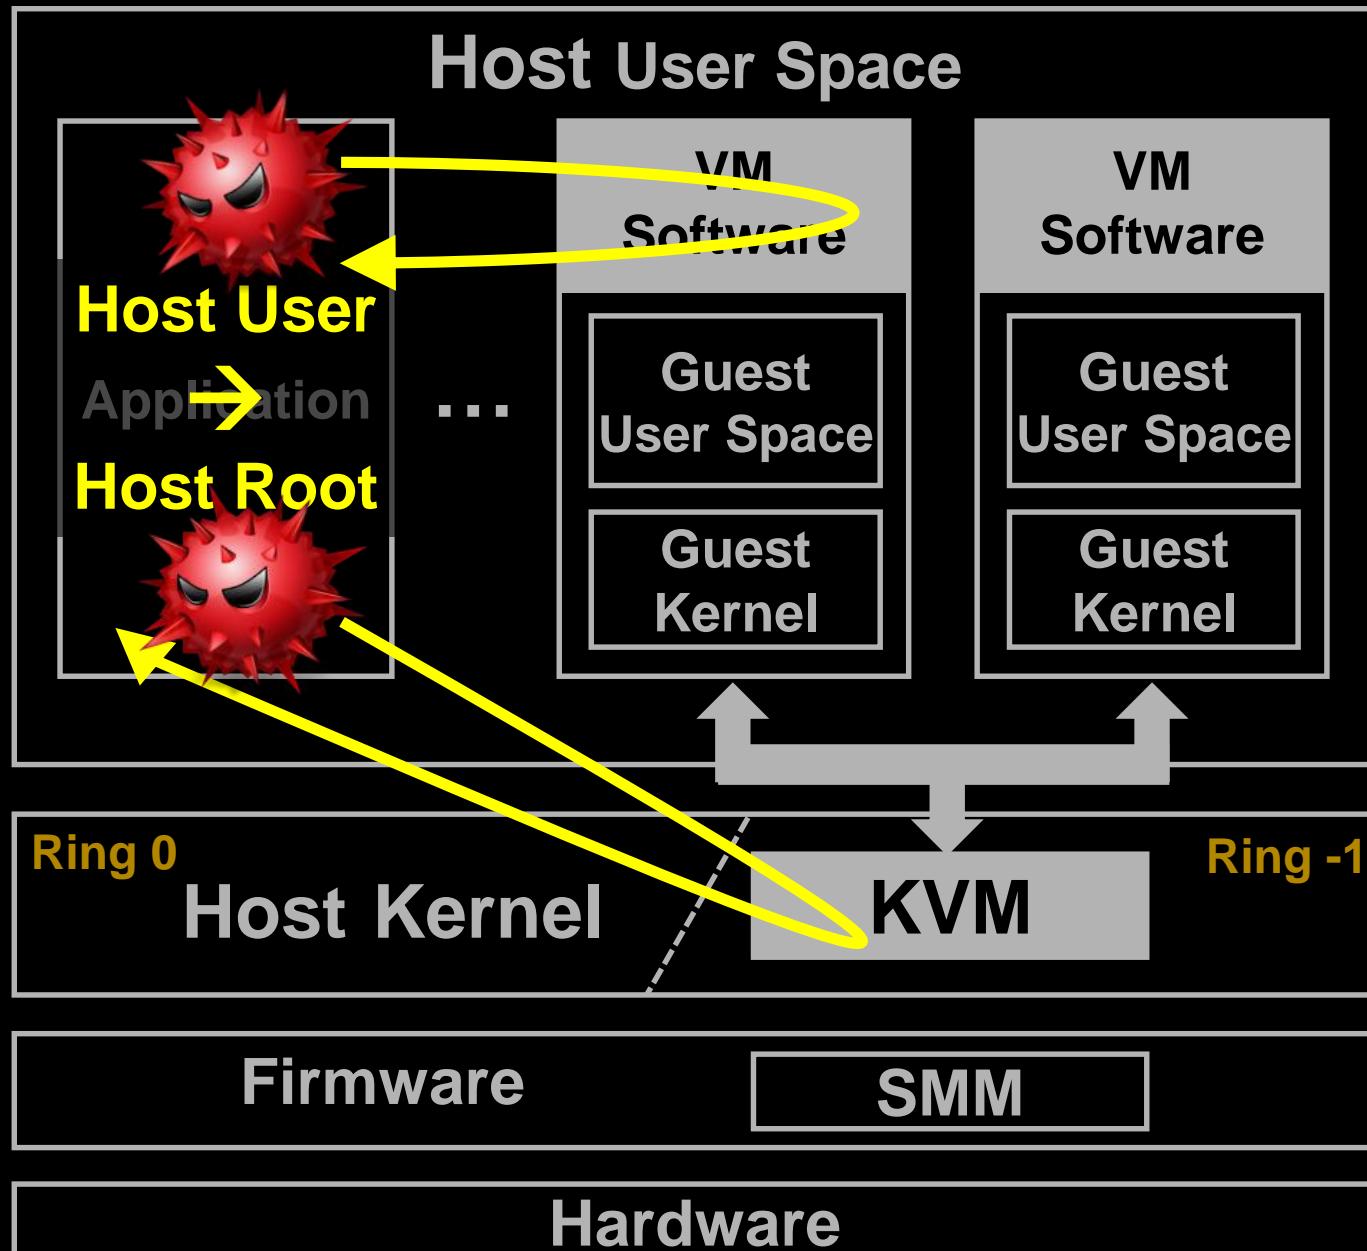


Extracting Escape Paths from CVEs (4)

Local Privilege Escalation with a KVM and VM Vuln.

- Elevates the privilege of a process

- Some KVM vulnerabilities can be exploited through /dev/kvm
 - It can be used to elevate privilege (and create a root privilege process)
- Some VM vulnerabilities can be used for privilege escalation if the VM or related daemon runs in root privilege
 - But, a VM usually runs in user privilege for security reasons 😞



3d Red Pill

A Guest-to-Host Escape on QEMU/KVM Virtio Device

College of Cyber Security Jinan University
Zhijian Shao, Jian Weng, Yue Zhang

OOB in VIRTIO-GPU of QEMU

Escape From The Docker-KVM-QEMU Machine

Shengping Wang, Xu Liu
Qihoo 360 Marvel Team

OOB in NET of QEMU

VENOM Vulnerability: Community Patching and Mitigation Update

May 13, 2015 Editorial Team Executive Viewpoint



OOB in FDC of QEMU

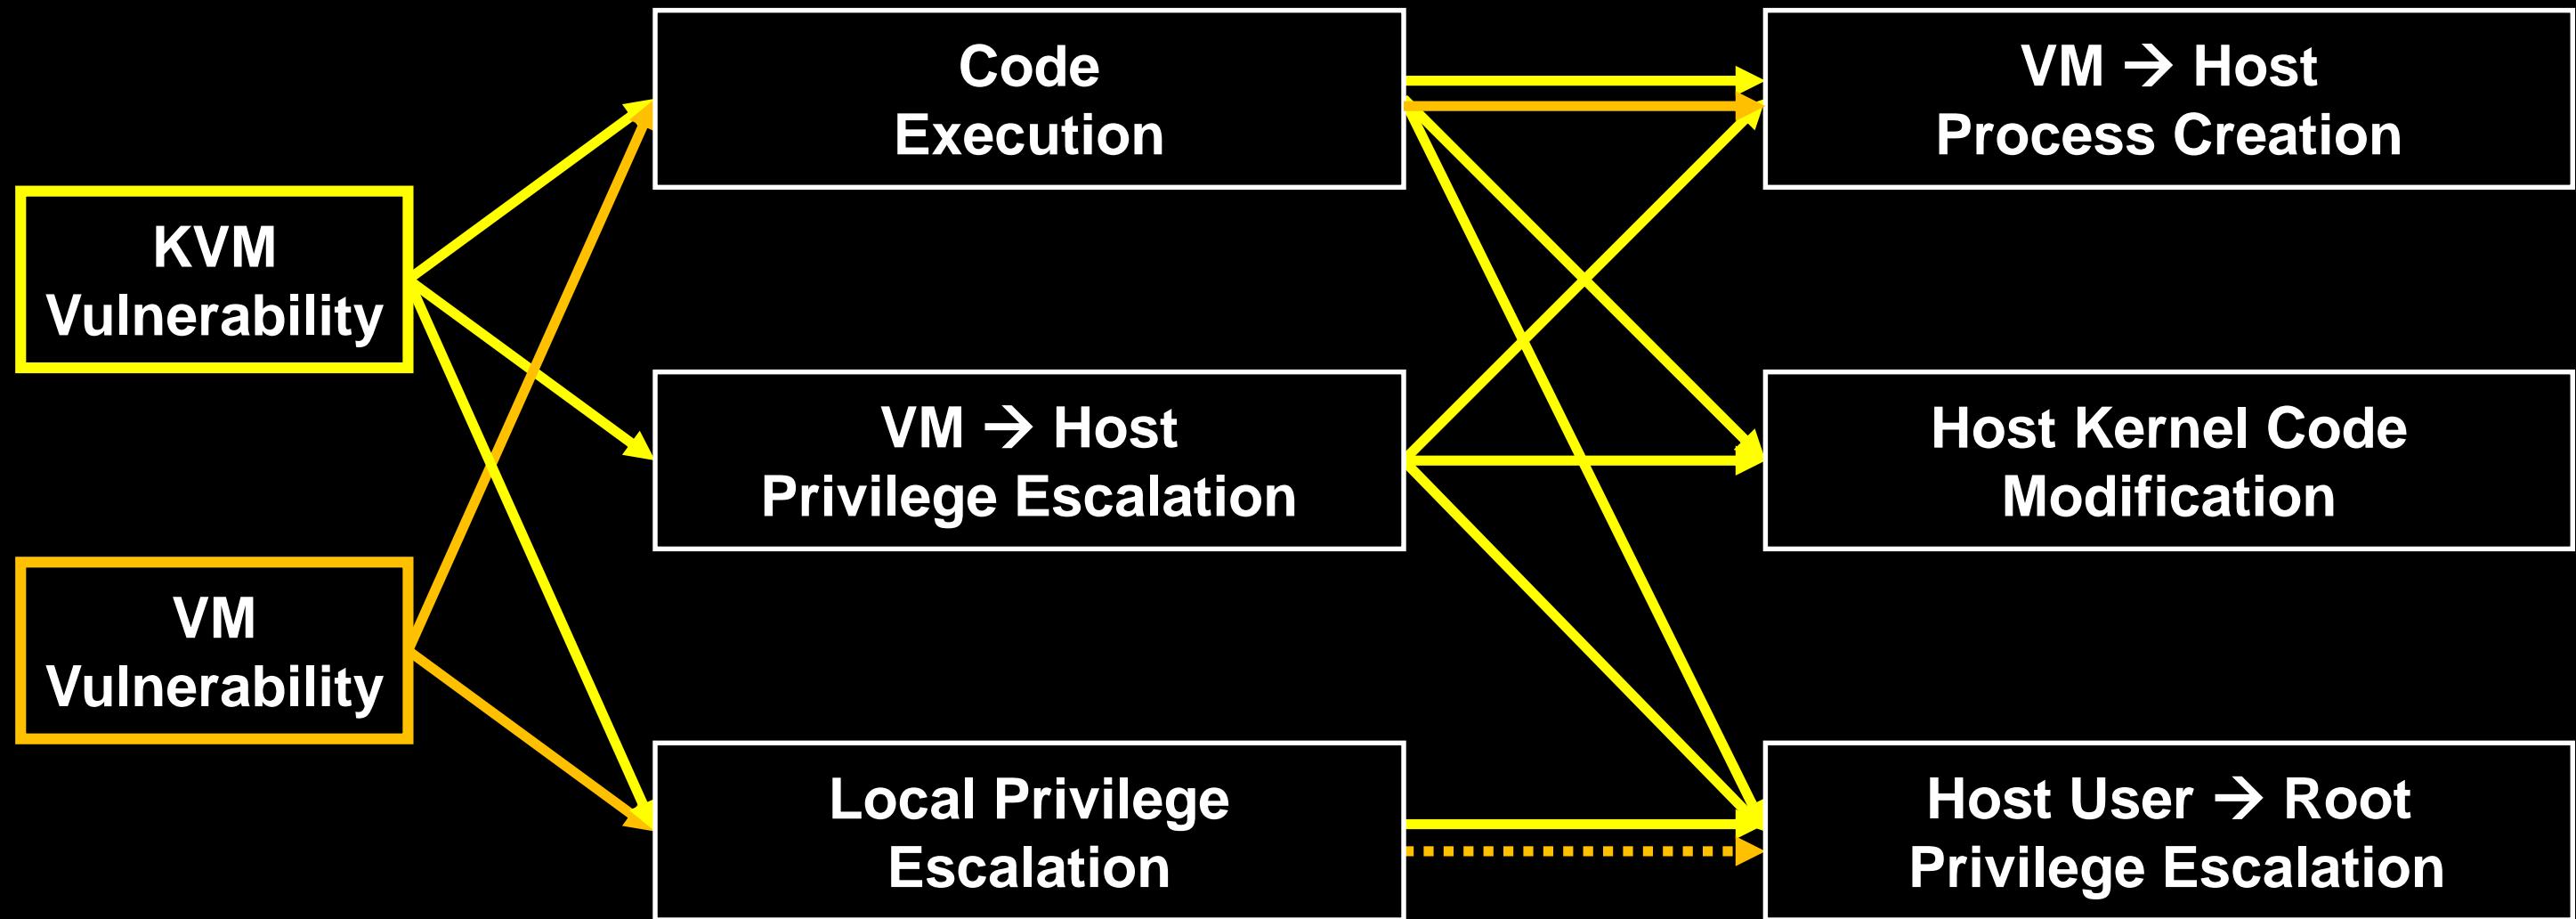
Virtunoid: Breaking out of KVM

Nelson Elhage

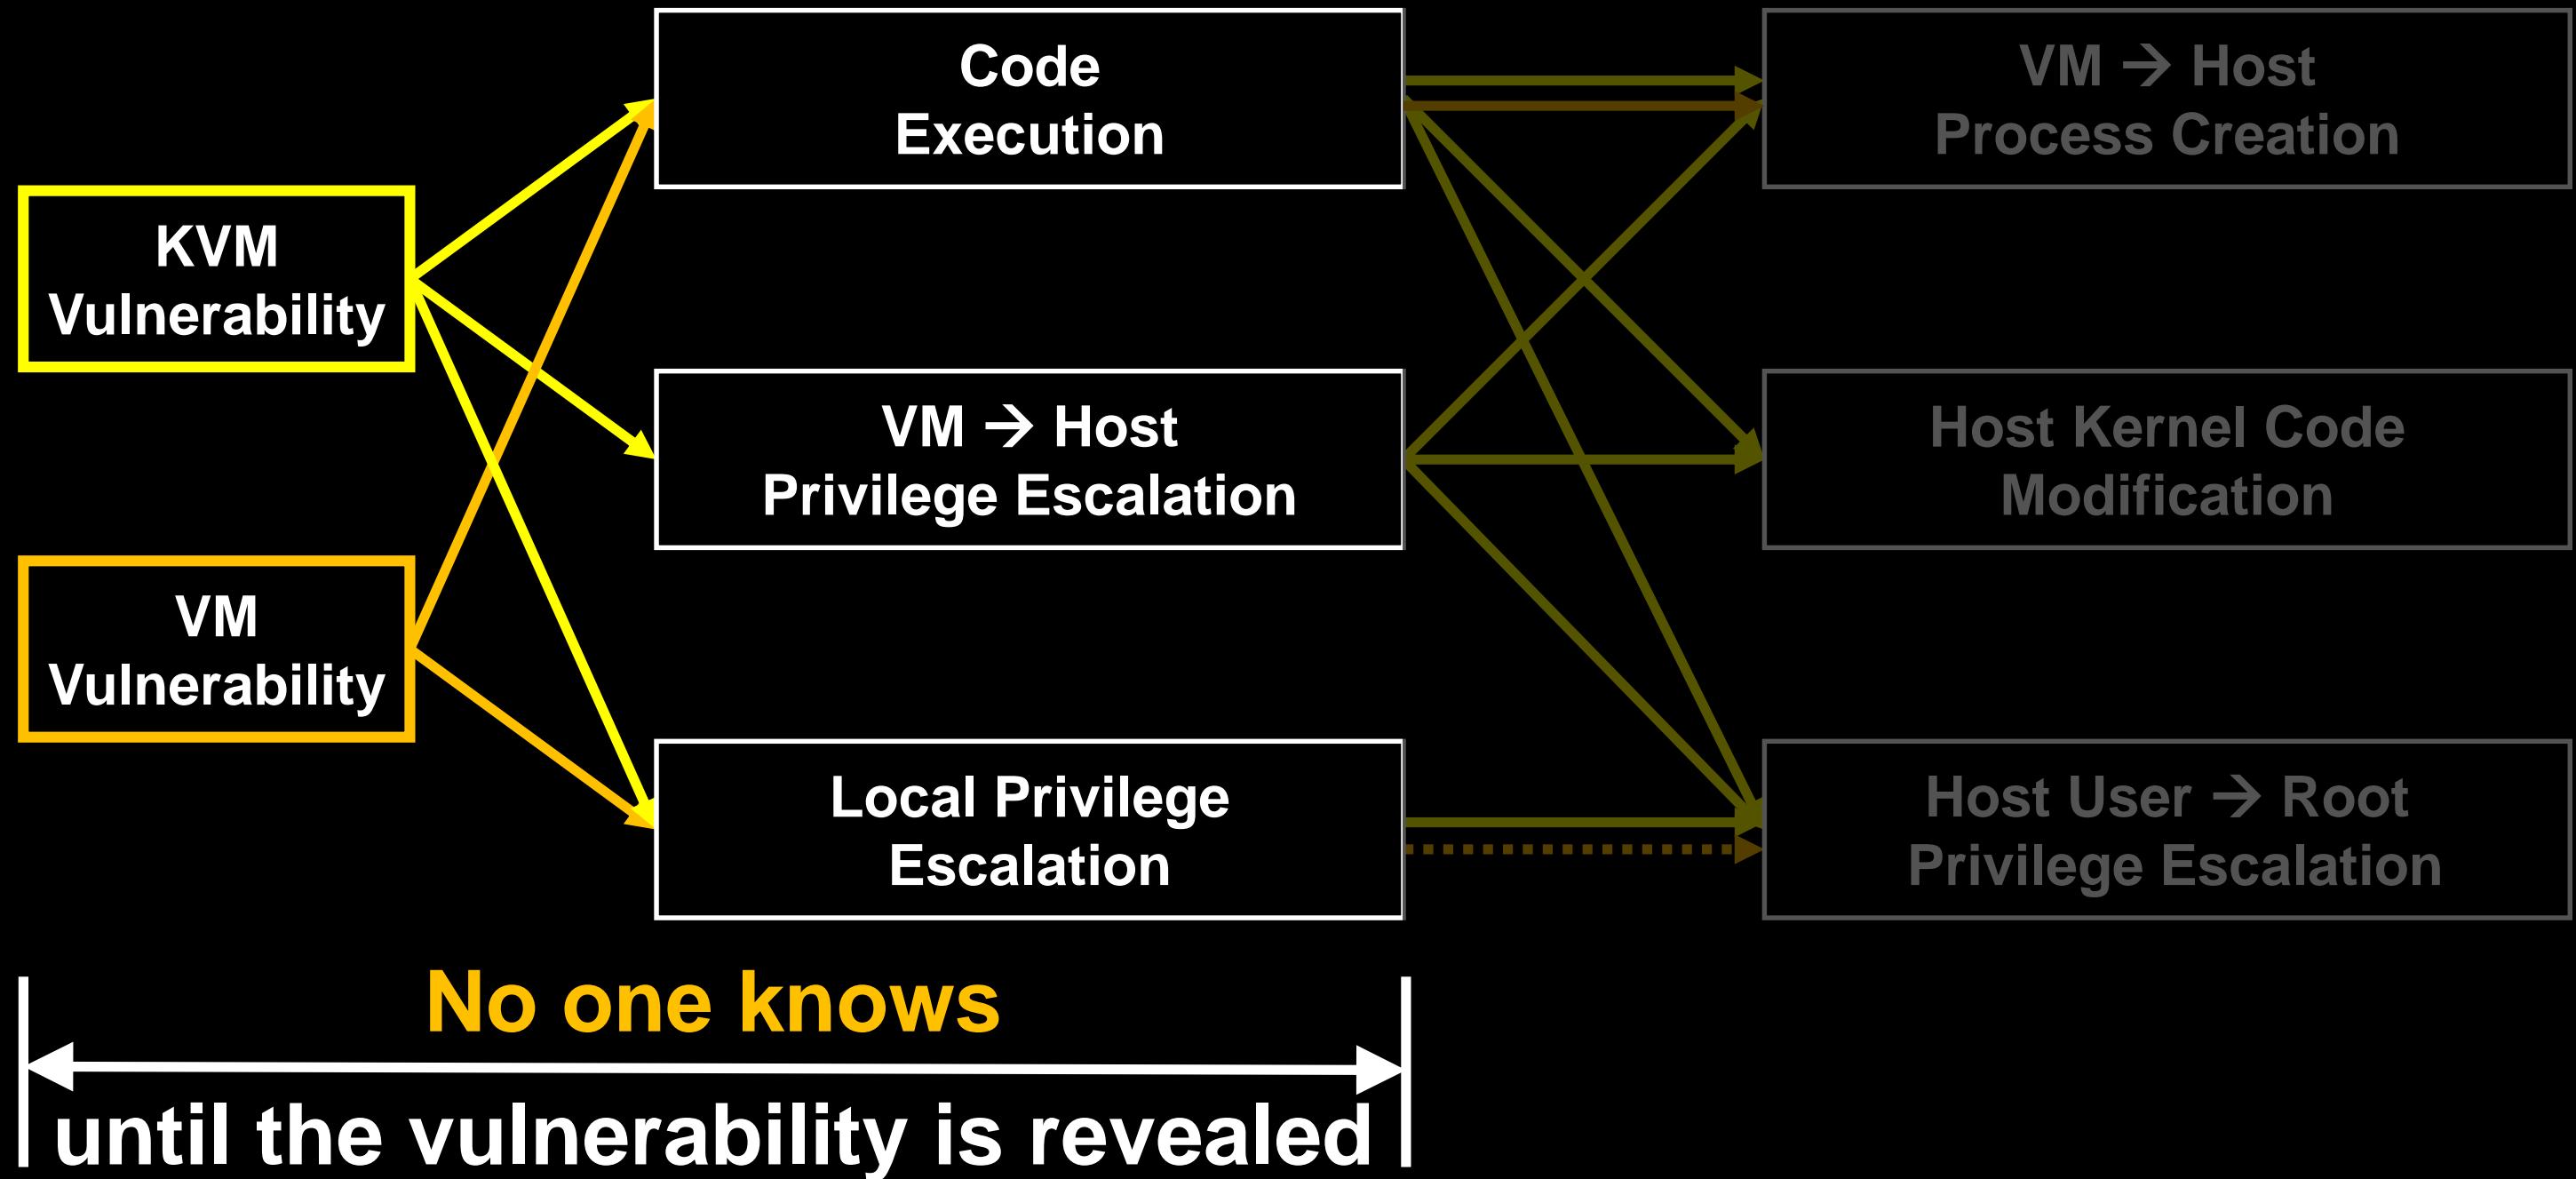
July 25, 2011

UAF in Timer of QEMU

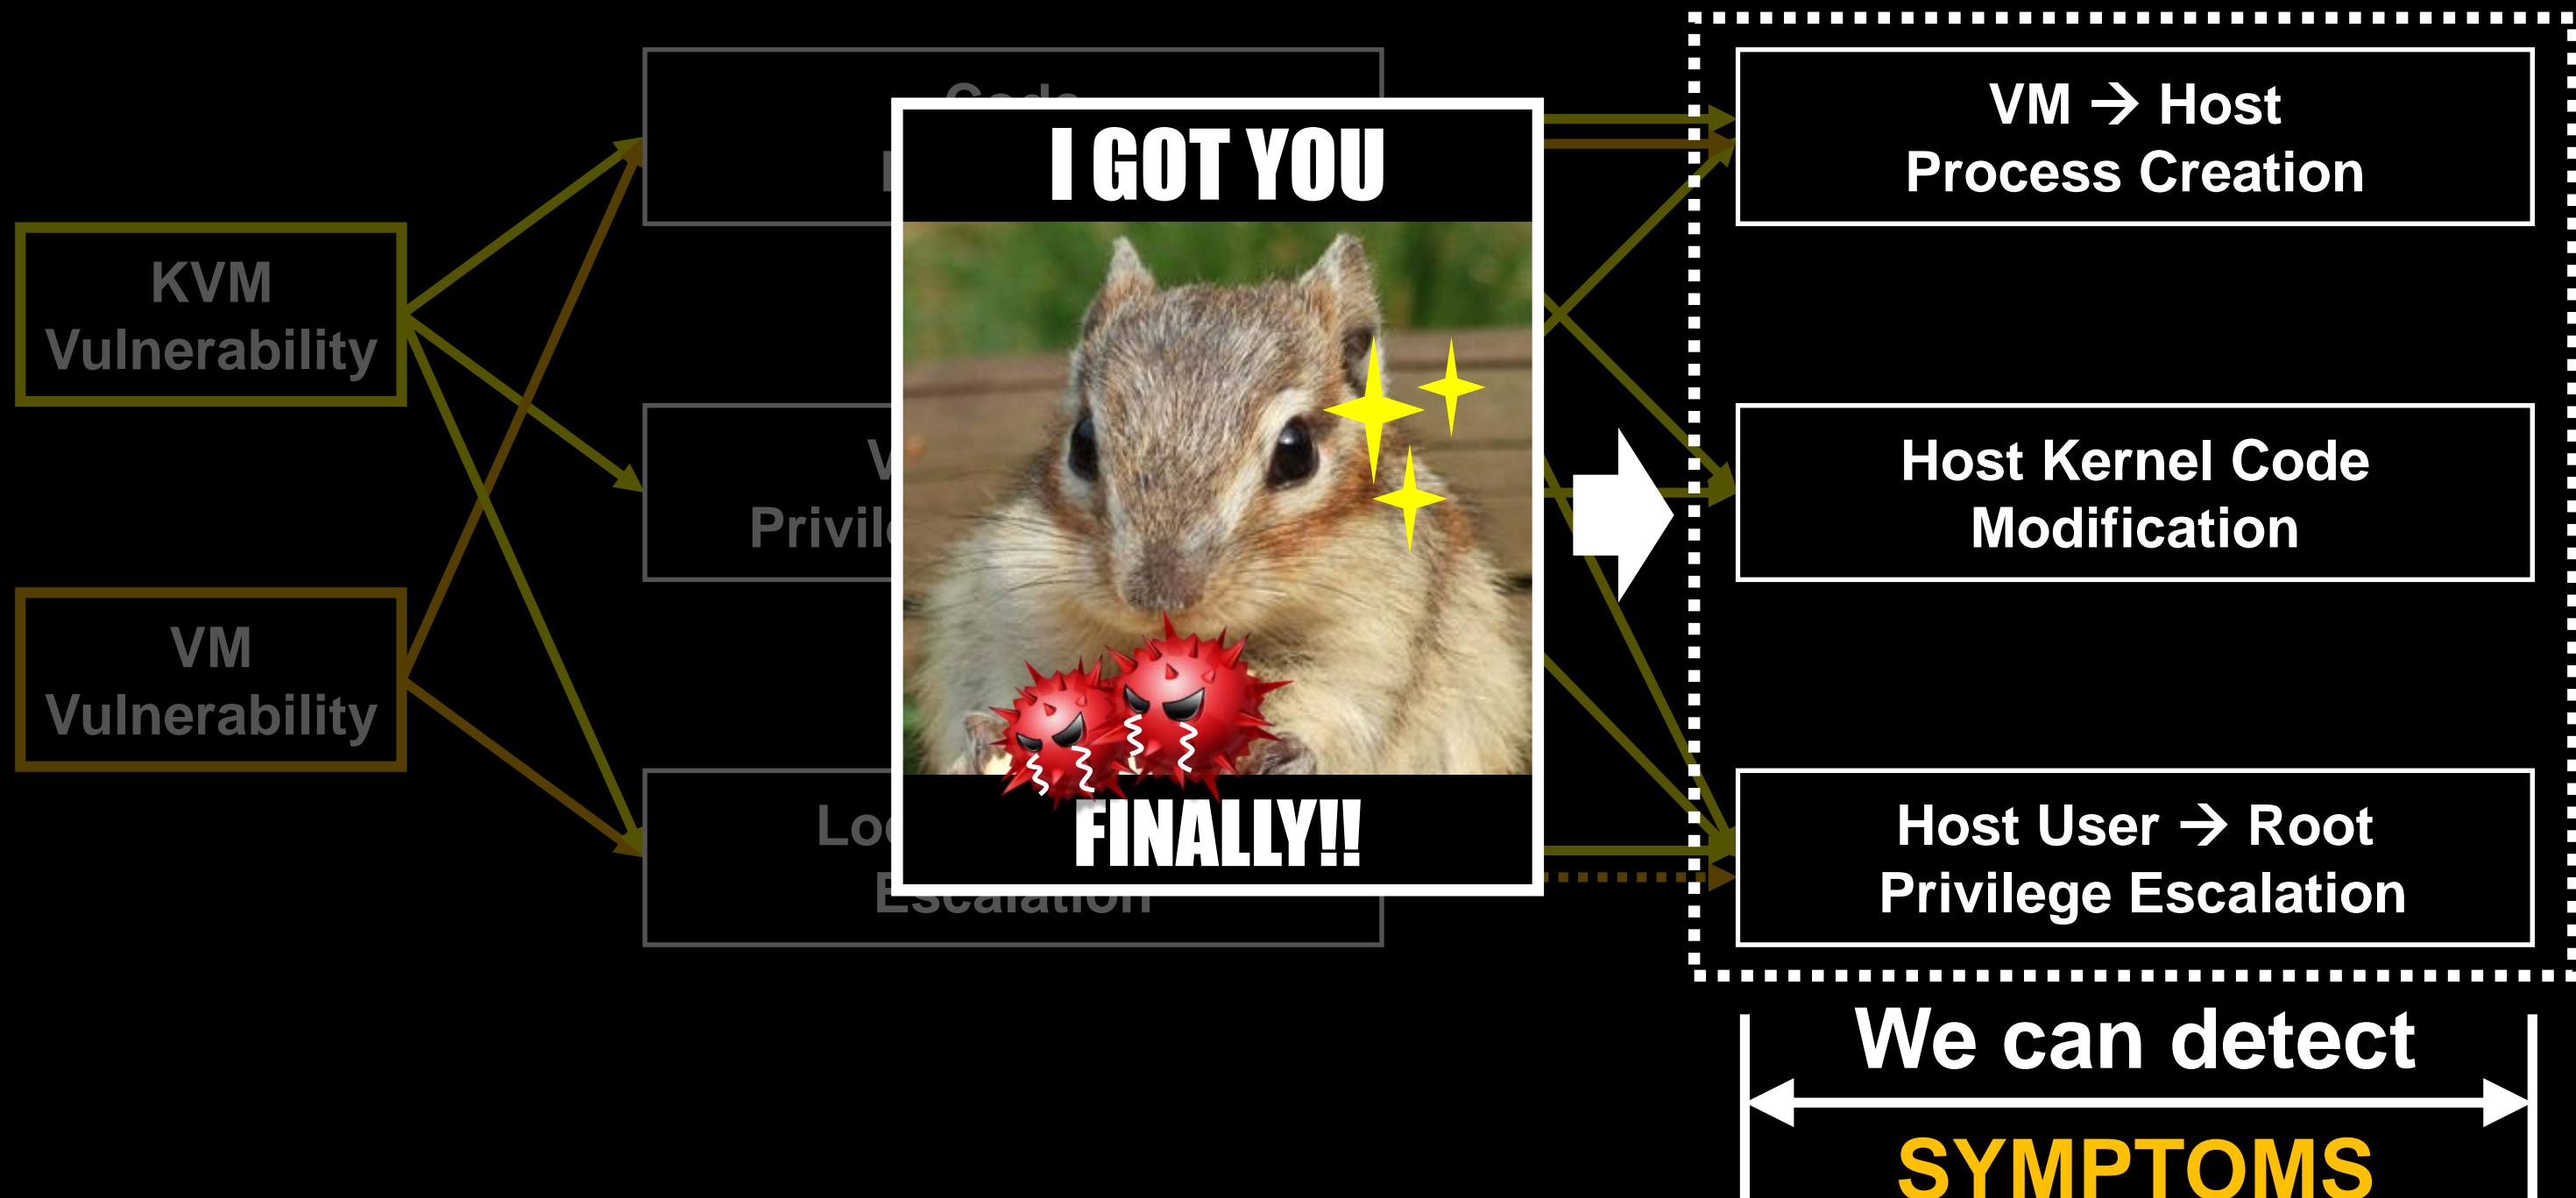
So, critical paths of escapes are ...



So, critical paths of escapes are ...



So, critical paths of escapes are ...

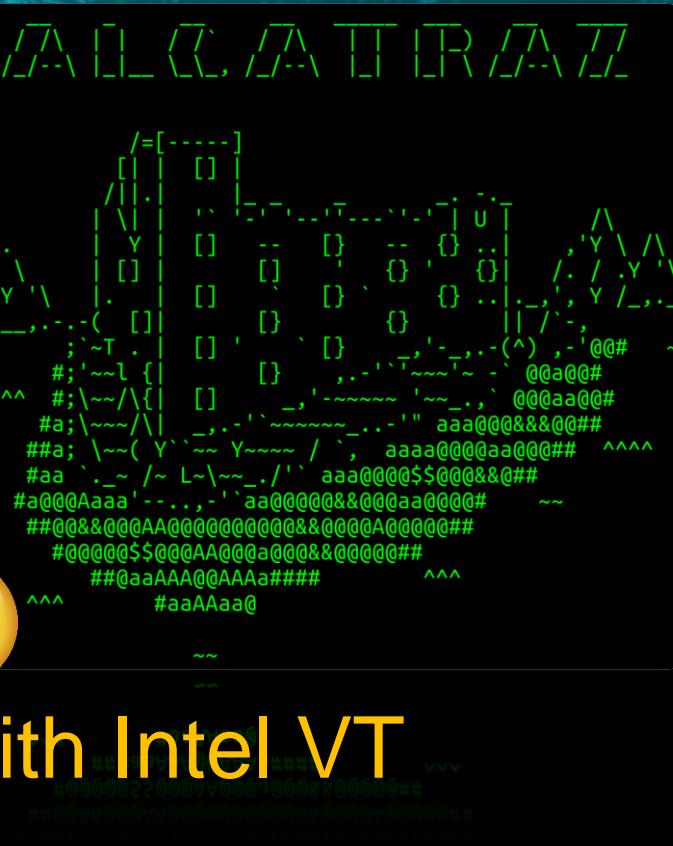


- Background
- Analysis of Critical Paths of Escapes
- Design and Implementation of Alcatraz
- Evaluation and Demo
- Conclusion and Black Hat Sound Bytes



What is Alcatraz?

- Alcatraz is a practical hypervisor sandbox to prevent escapes
 - It does not need SMM and retrofits of the hypervisor 😊
 - It downgrades KVM's privilege and makes a sandbox with Intel VT
- Alcatraz consists of **Hyper-box** and **tailored Linux kernel**
 - Hyper-box is a pico hypervisor made from scratch
 - It has core mechanisms for preventing escapes
 - Tailored Linux kernel has only one system call interface for x86_64
 - It narrows the attack surface and monitoring points



Hyper-Box (1)

- Makes a sandbox with the memory and register protection techniques
 - It leverages the Extended Page Table (EPT) of Intel VT
 - Guest Physical Address (PA) are translated to host PA with EPT
 - EPT can prevent code and RO data modifications!
 - It protects security-related bits of control registers (CRs)
 - Write Protect (WP) of CR0, Supervisor Mode Execution Protection (SMEP) of CR4
- Monitors system calls and prevents unauthorized behaviors
 - It sets H/W breakpoints to the system call entry point, fork, exit, and cred functions
 - `wake_up_new_task()`, `cgroup_release()`, `entry_SYSCALL_64()`, and `commit_creds()`
 - It detects and prevents unauthorized behaviors
 - code modification, process creation, and privilege escalation

Hyper-Box (2)

- Emulates VMX instructions for nested virtualization
 - It downgrades KVM's privilege (Ring -1) to a guest hypervisor (Ring 0)
 - So, It has to emulate VMX instructions of KVM because Ring 0 cannot execute them
 - However, performance overhead is high
- Uses VT features to reduce the performance overhead
 - Virtual Machine Control Structure (VMCS) shadowing
 - It allows a guest hypervisor to execute VMREAD and VMWRITE directly
 - So, emulations or interventions are reduced!
 - Virtual Process ID (VPID)
 - It gives unique ID to vCPU and Translation Lookaside Buffer (TLB)
 - It prevents TLB flushes when VMExit occurs!

Hyper-Box (2)

- Emulates VMX instructions
 - It downgrades KVM's privileged instructions
 - So, It has to emulate VMX instructions
 - However, performance is still better
- Uses VT features to reduce overhead
 - Virtual Machine Control Registers
 - It allows a guest hypervisor to access host hypervisor (Ring 0) directly
 - So, emulations or interventions are reduced
 - Virtual Process ID (VPID)
 - It gives unique ID to vCPU and Translation Lookaside Buffer (TLB)
 - It prevents TLB flushes when VMExit occurs!

AGAIN, WHERE AM I?

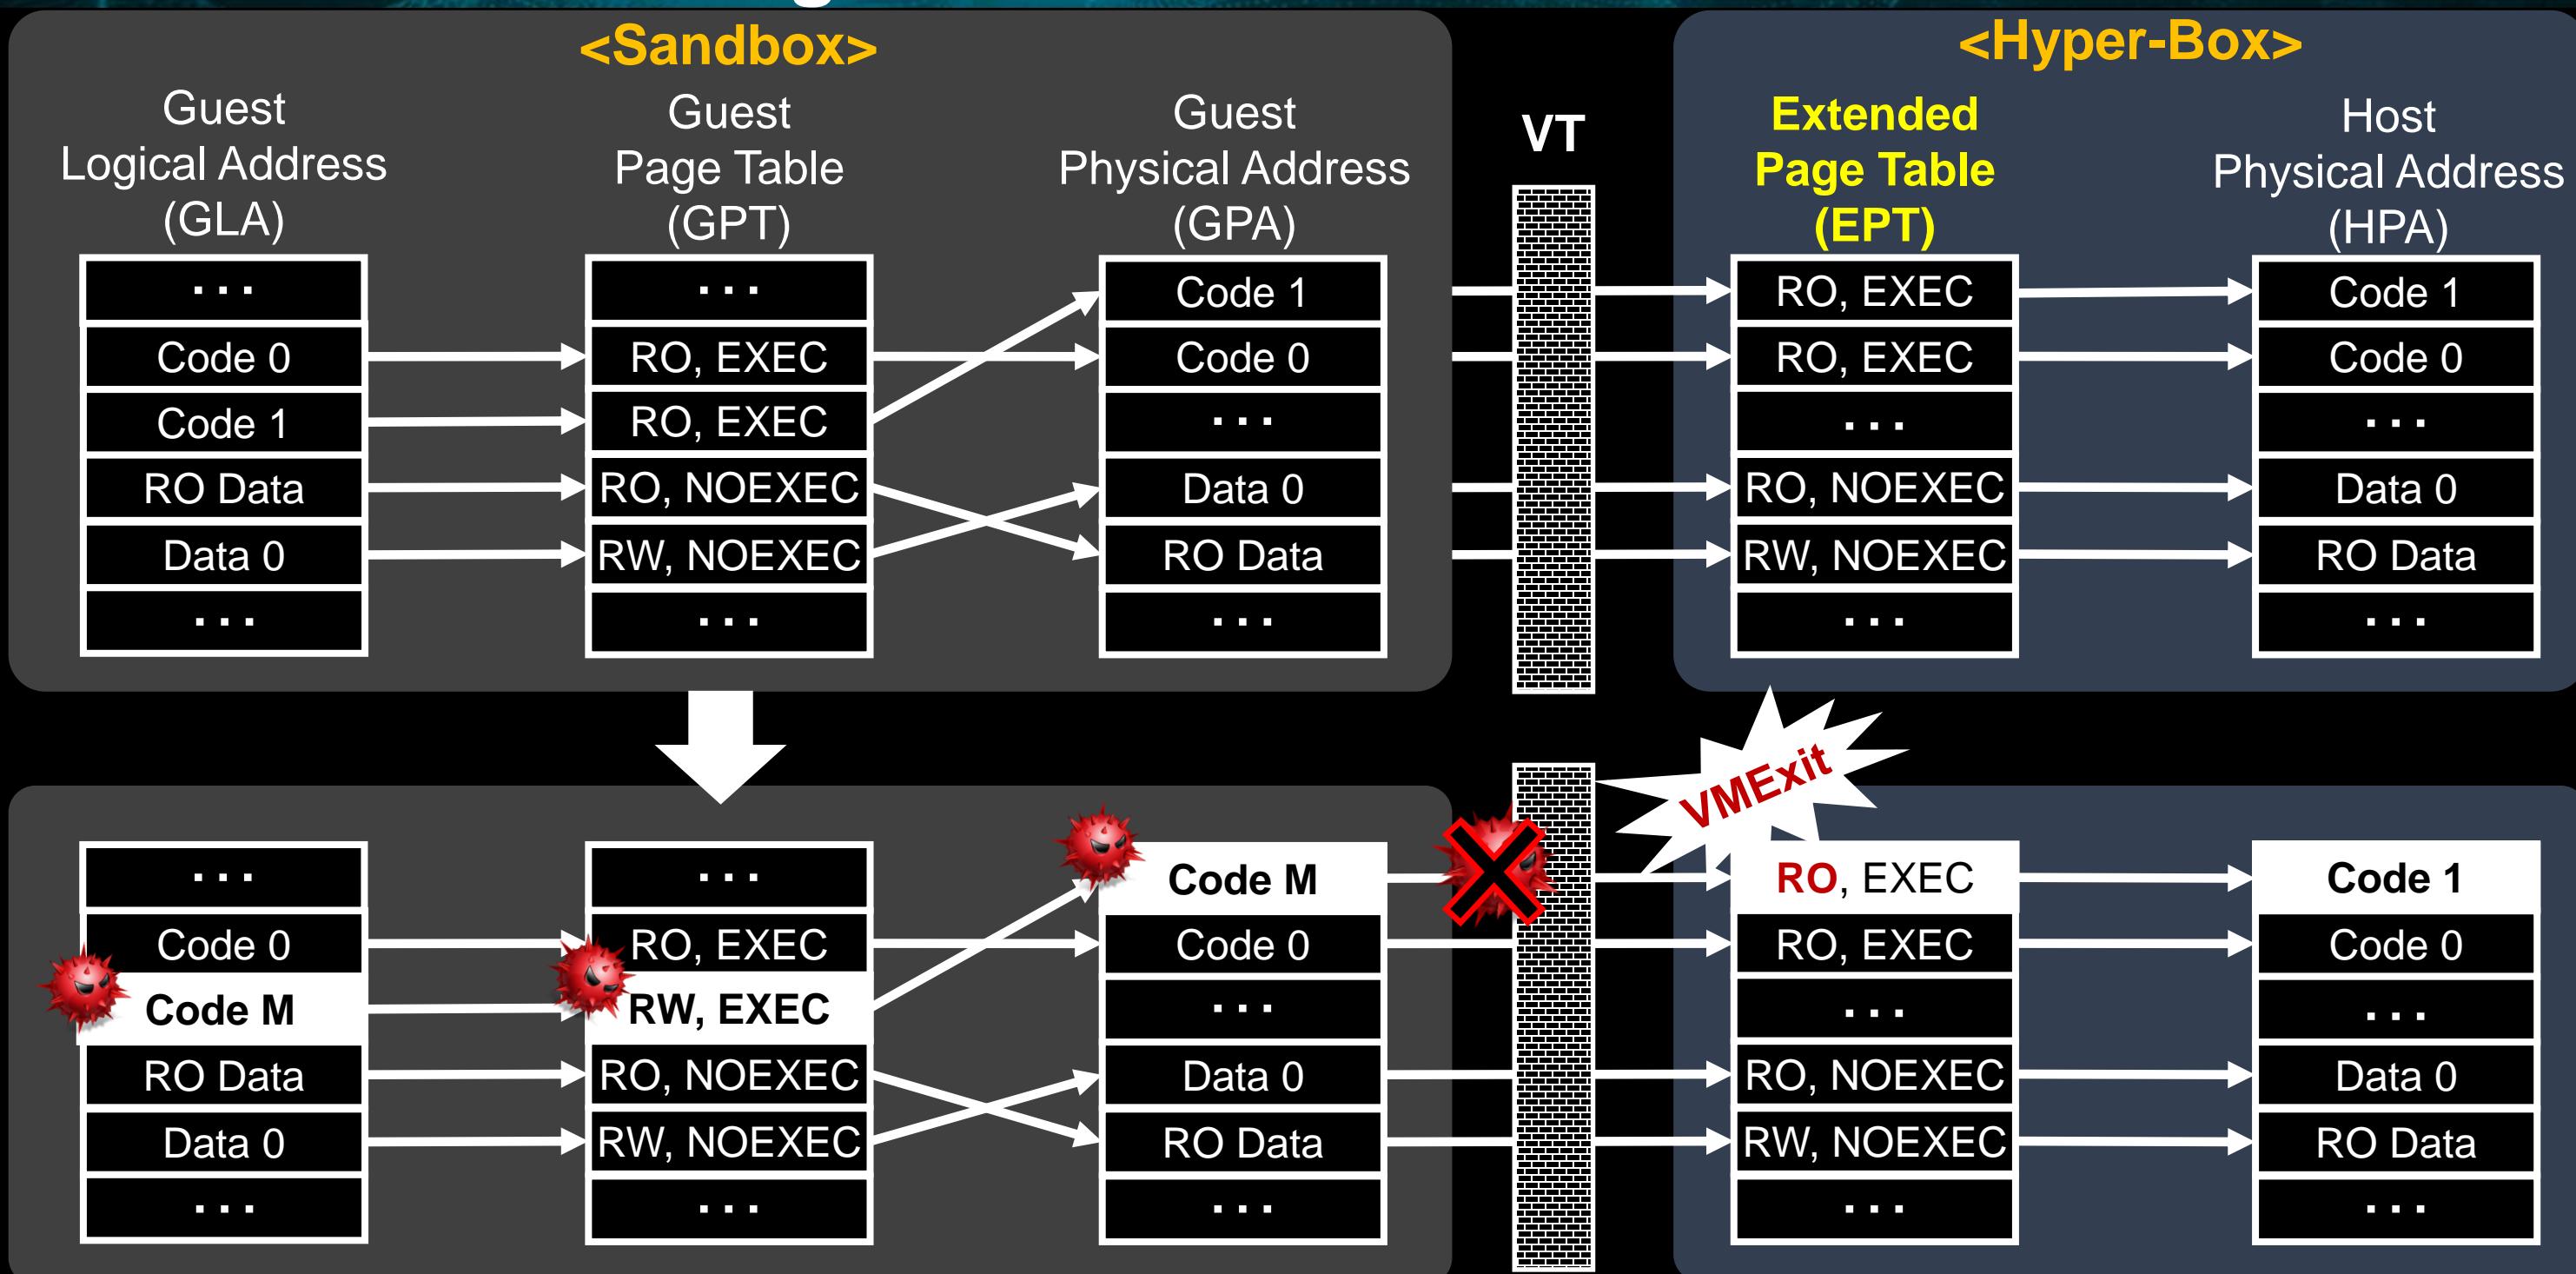


GIVE ME PICTURES!

- ualization
- st hypervisor (Ring 0) cause Ring 0 cannot execute them
- ce overhead
- ving
- D and VMWRITE directly

Hyper-Box Design

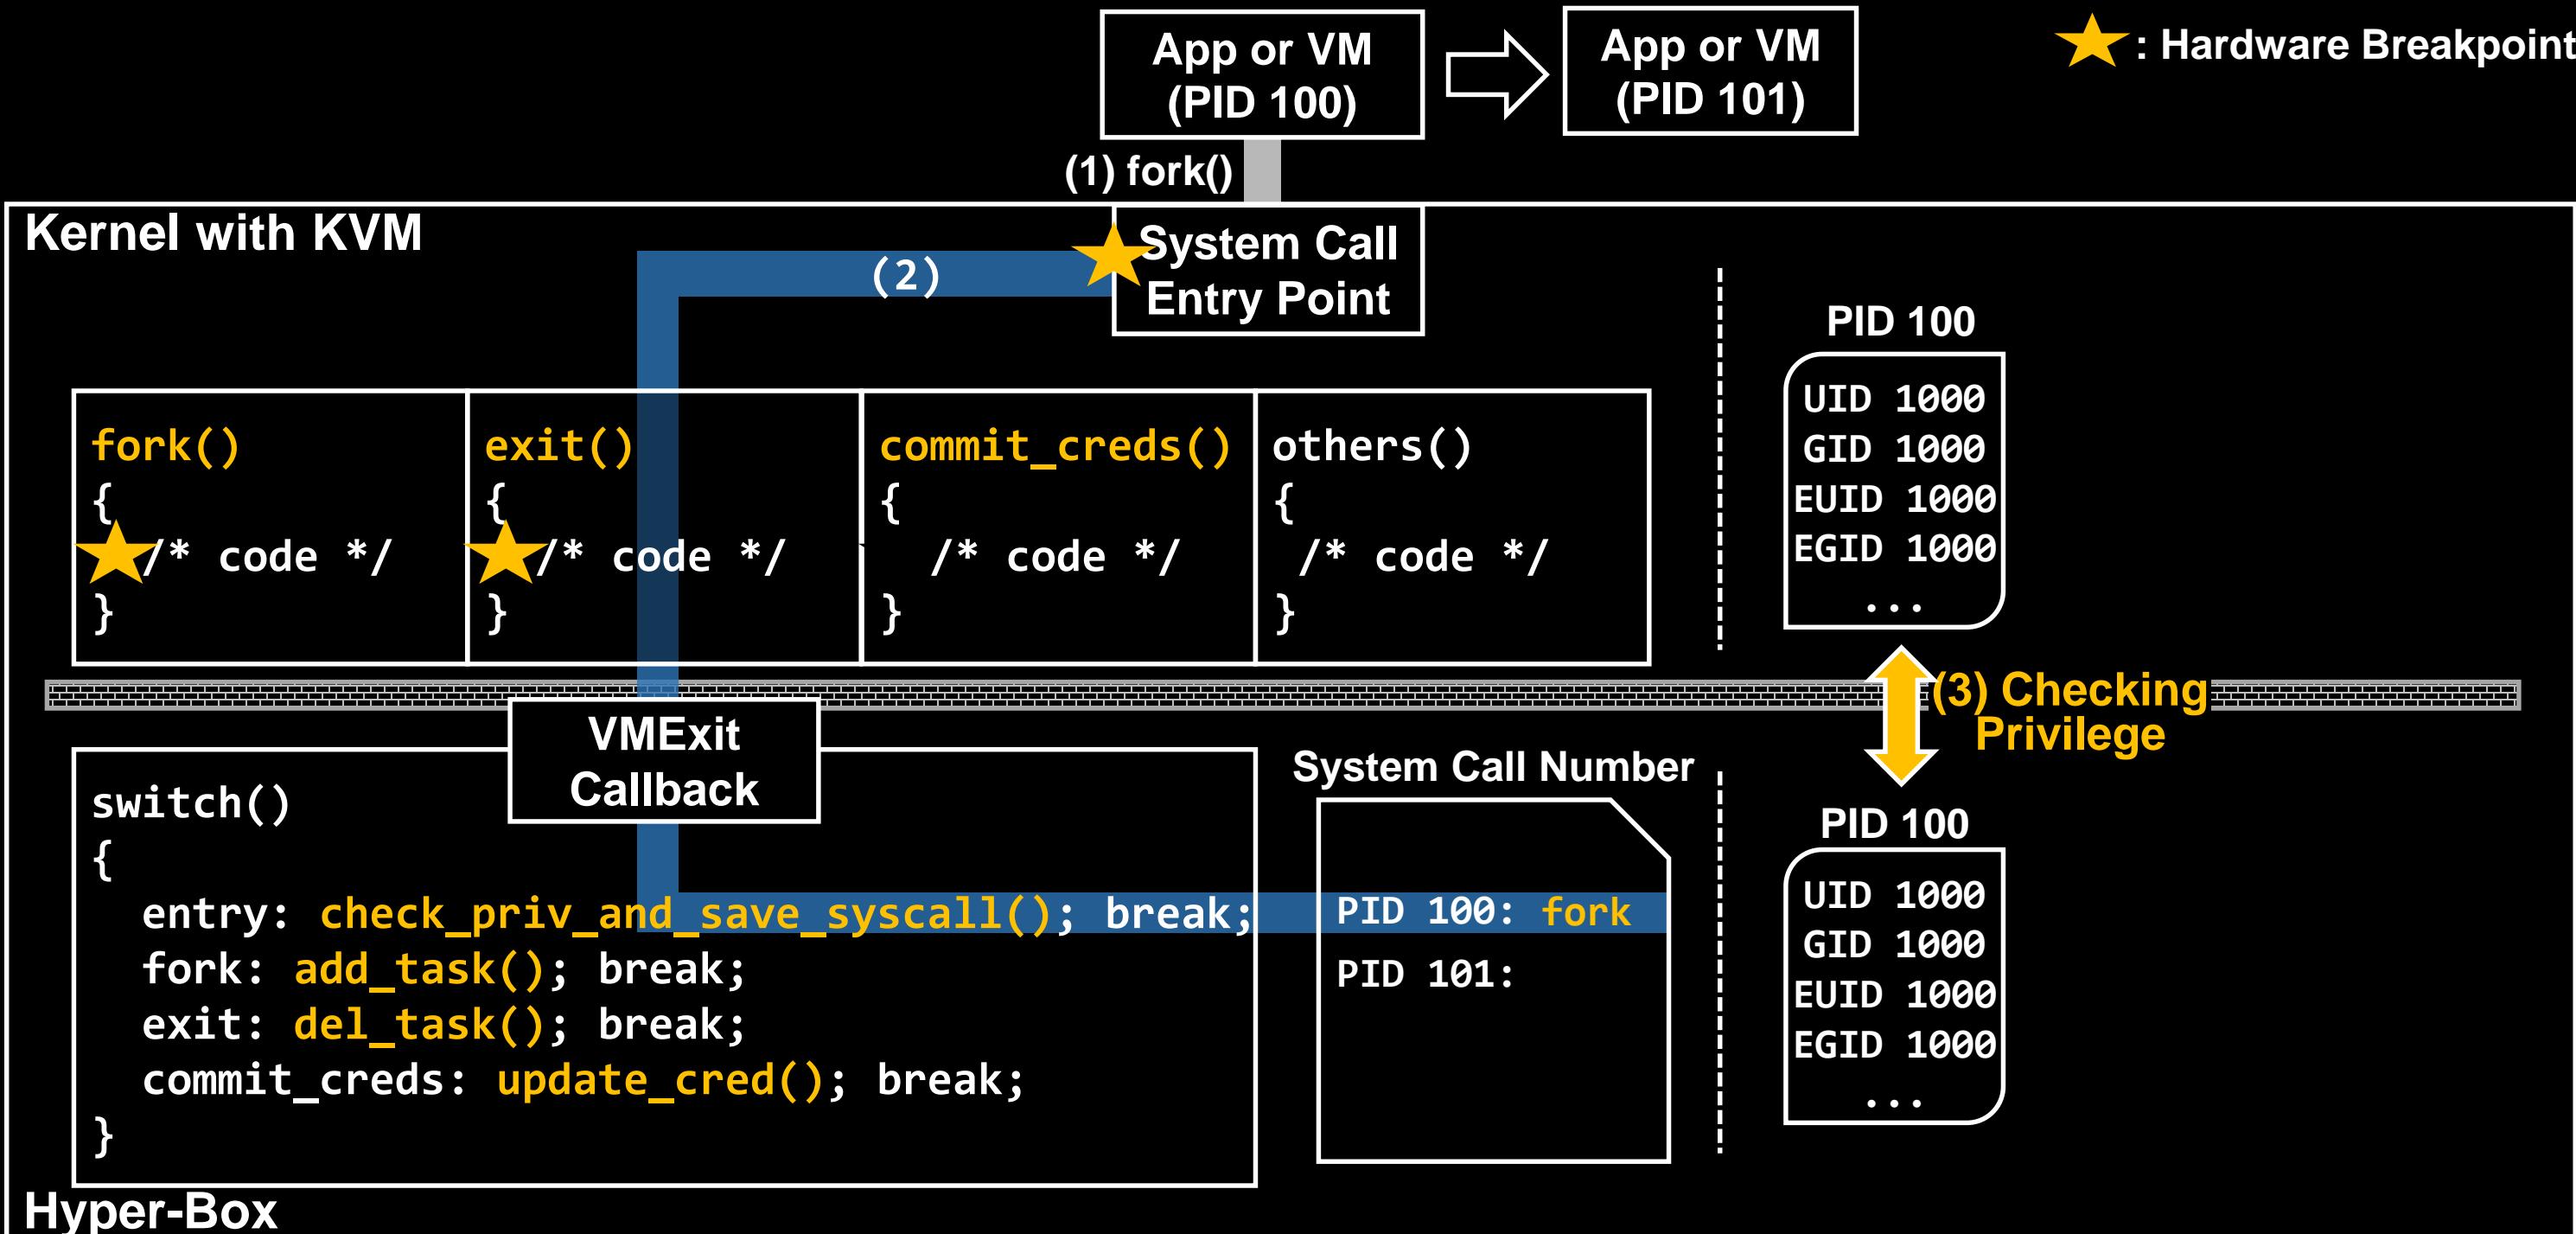
Preventing Unauthorized Code Modification



Hyper-Box Design

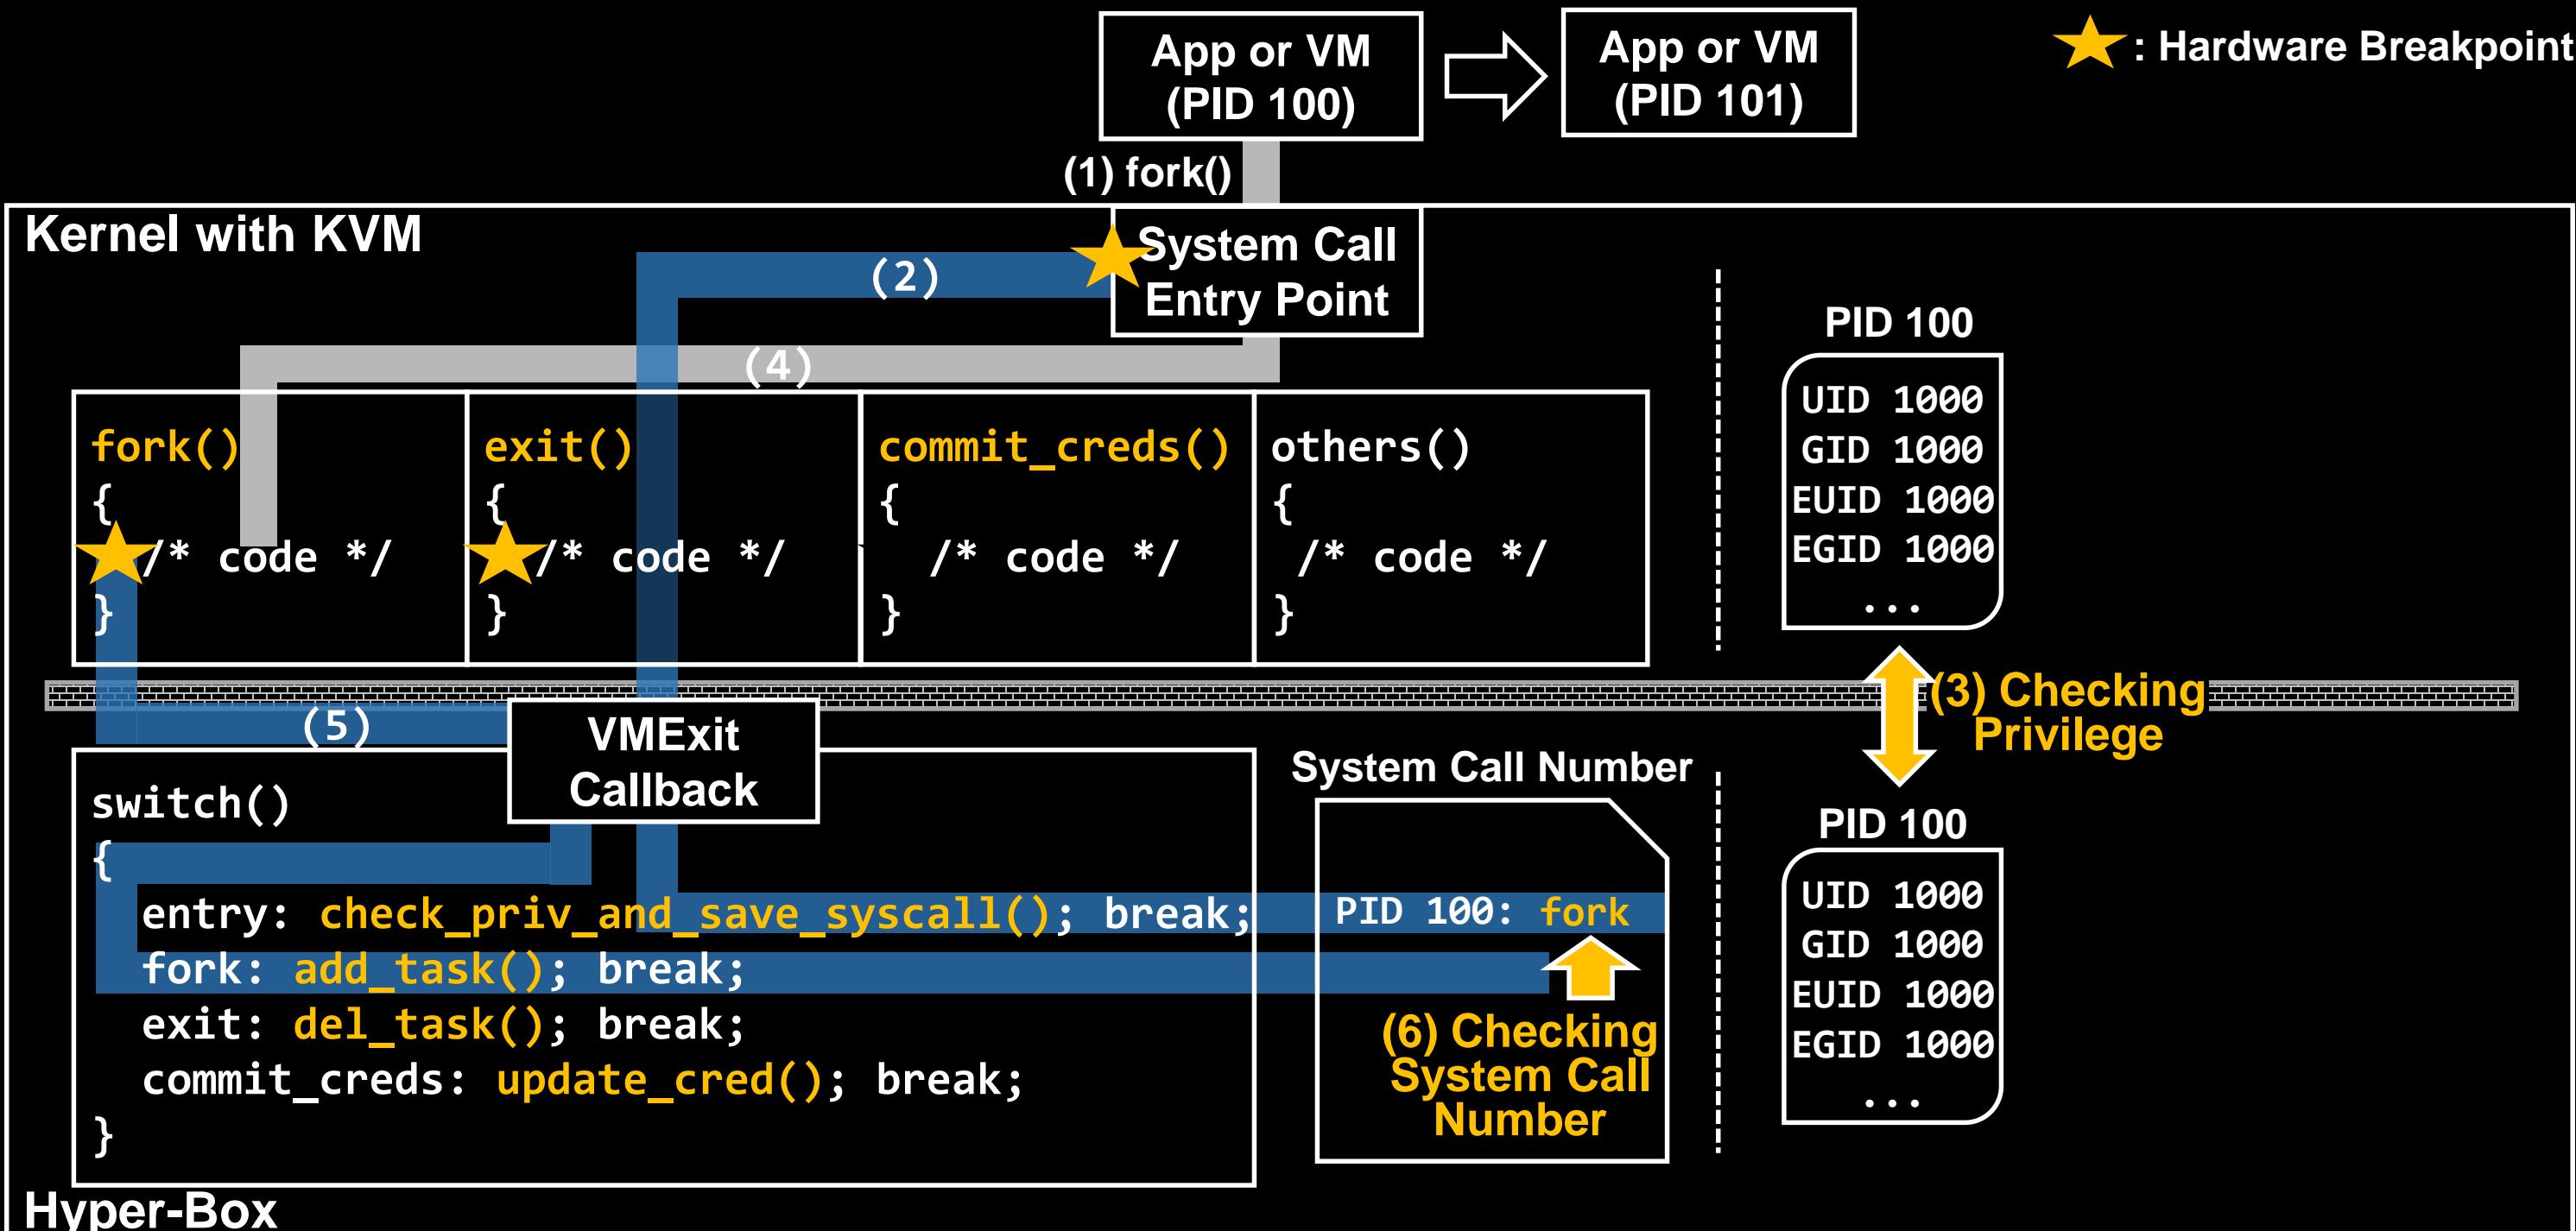
Preventing Unauthorized Process Creation (1)

★ : Hardware Breakpoint



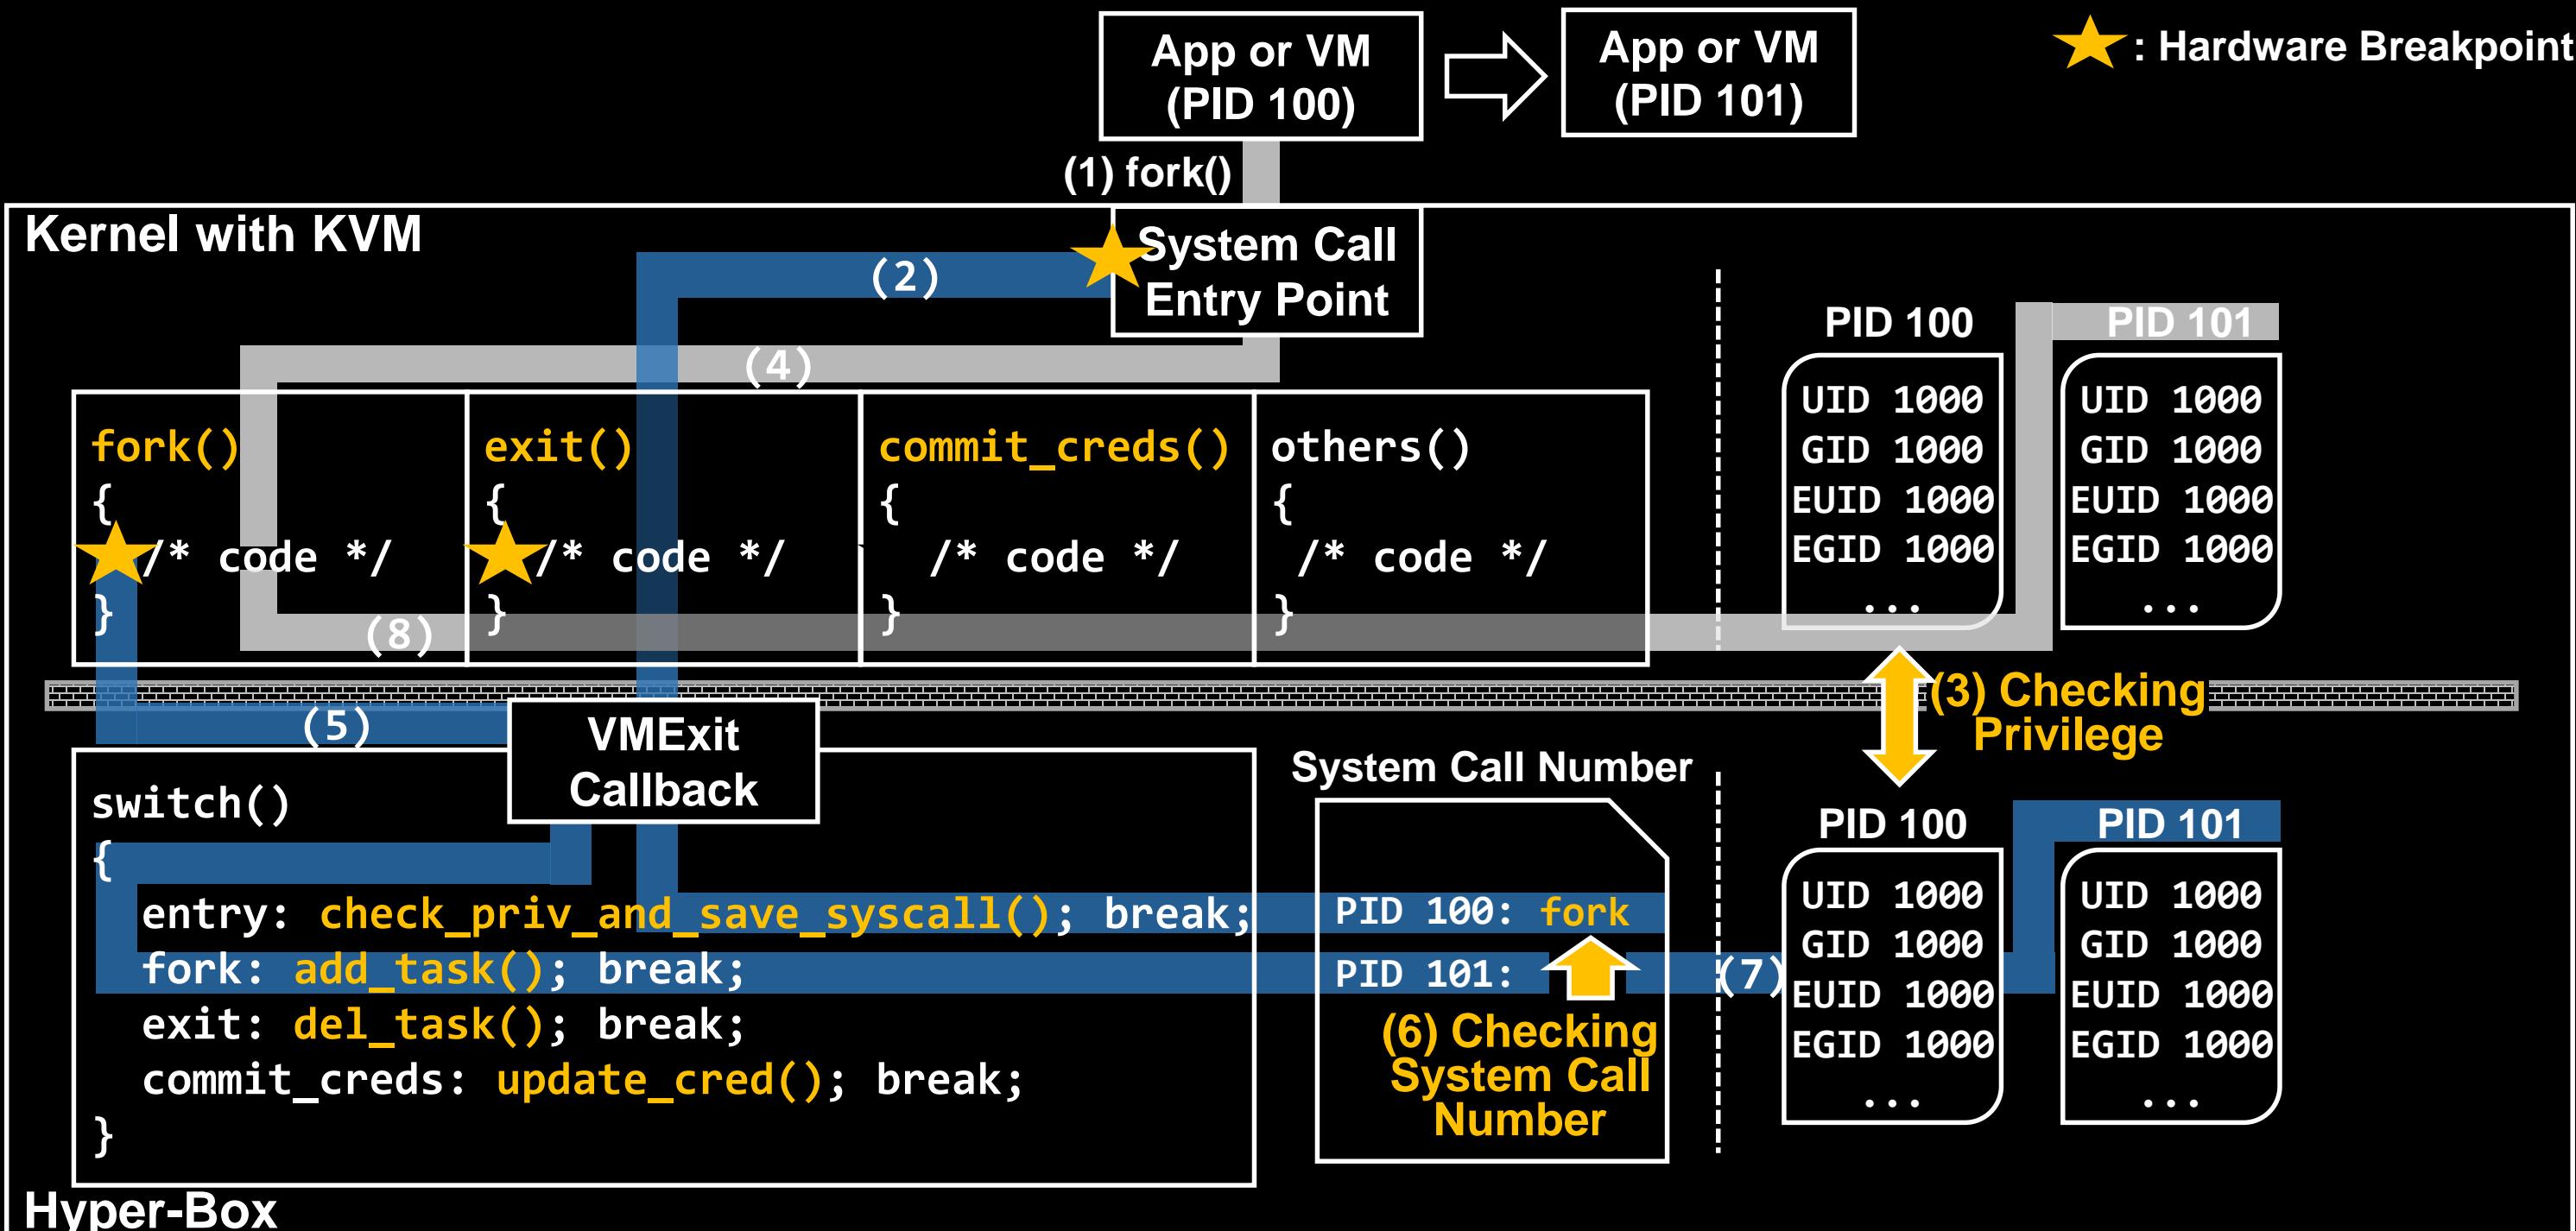
Hyper-Box Design

Preventing Unauthorized Process Creation (1)



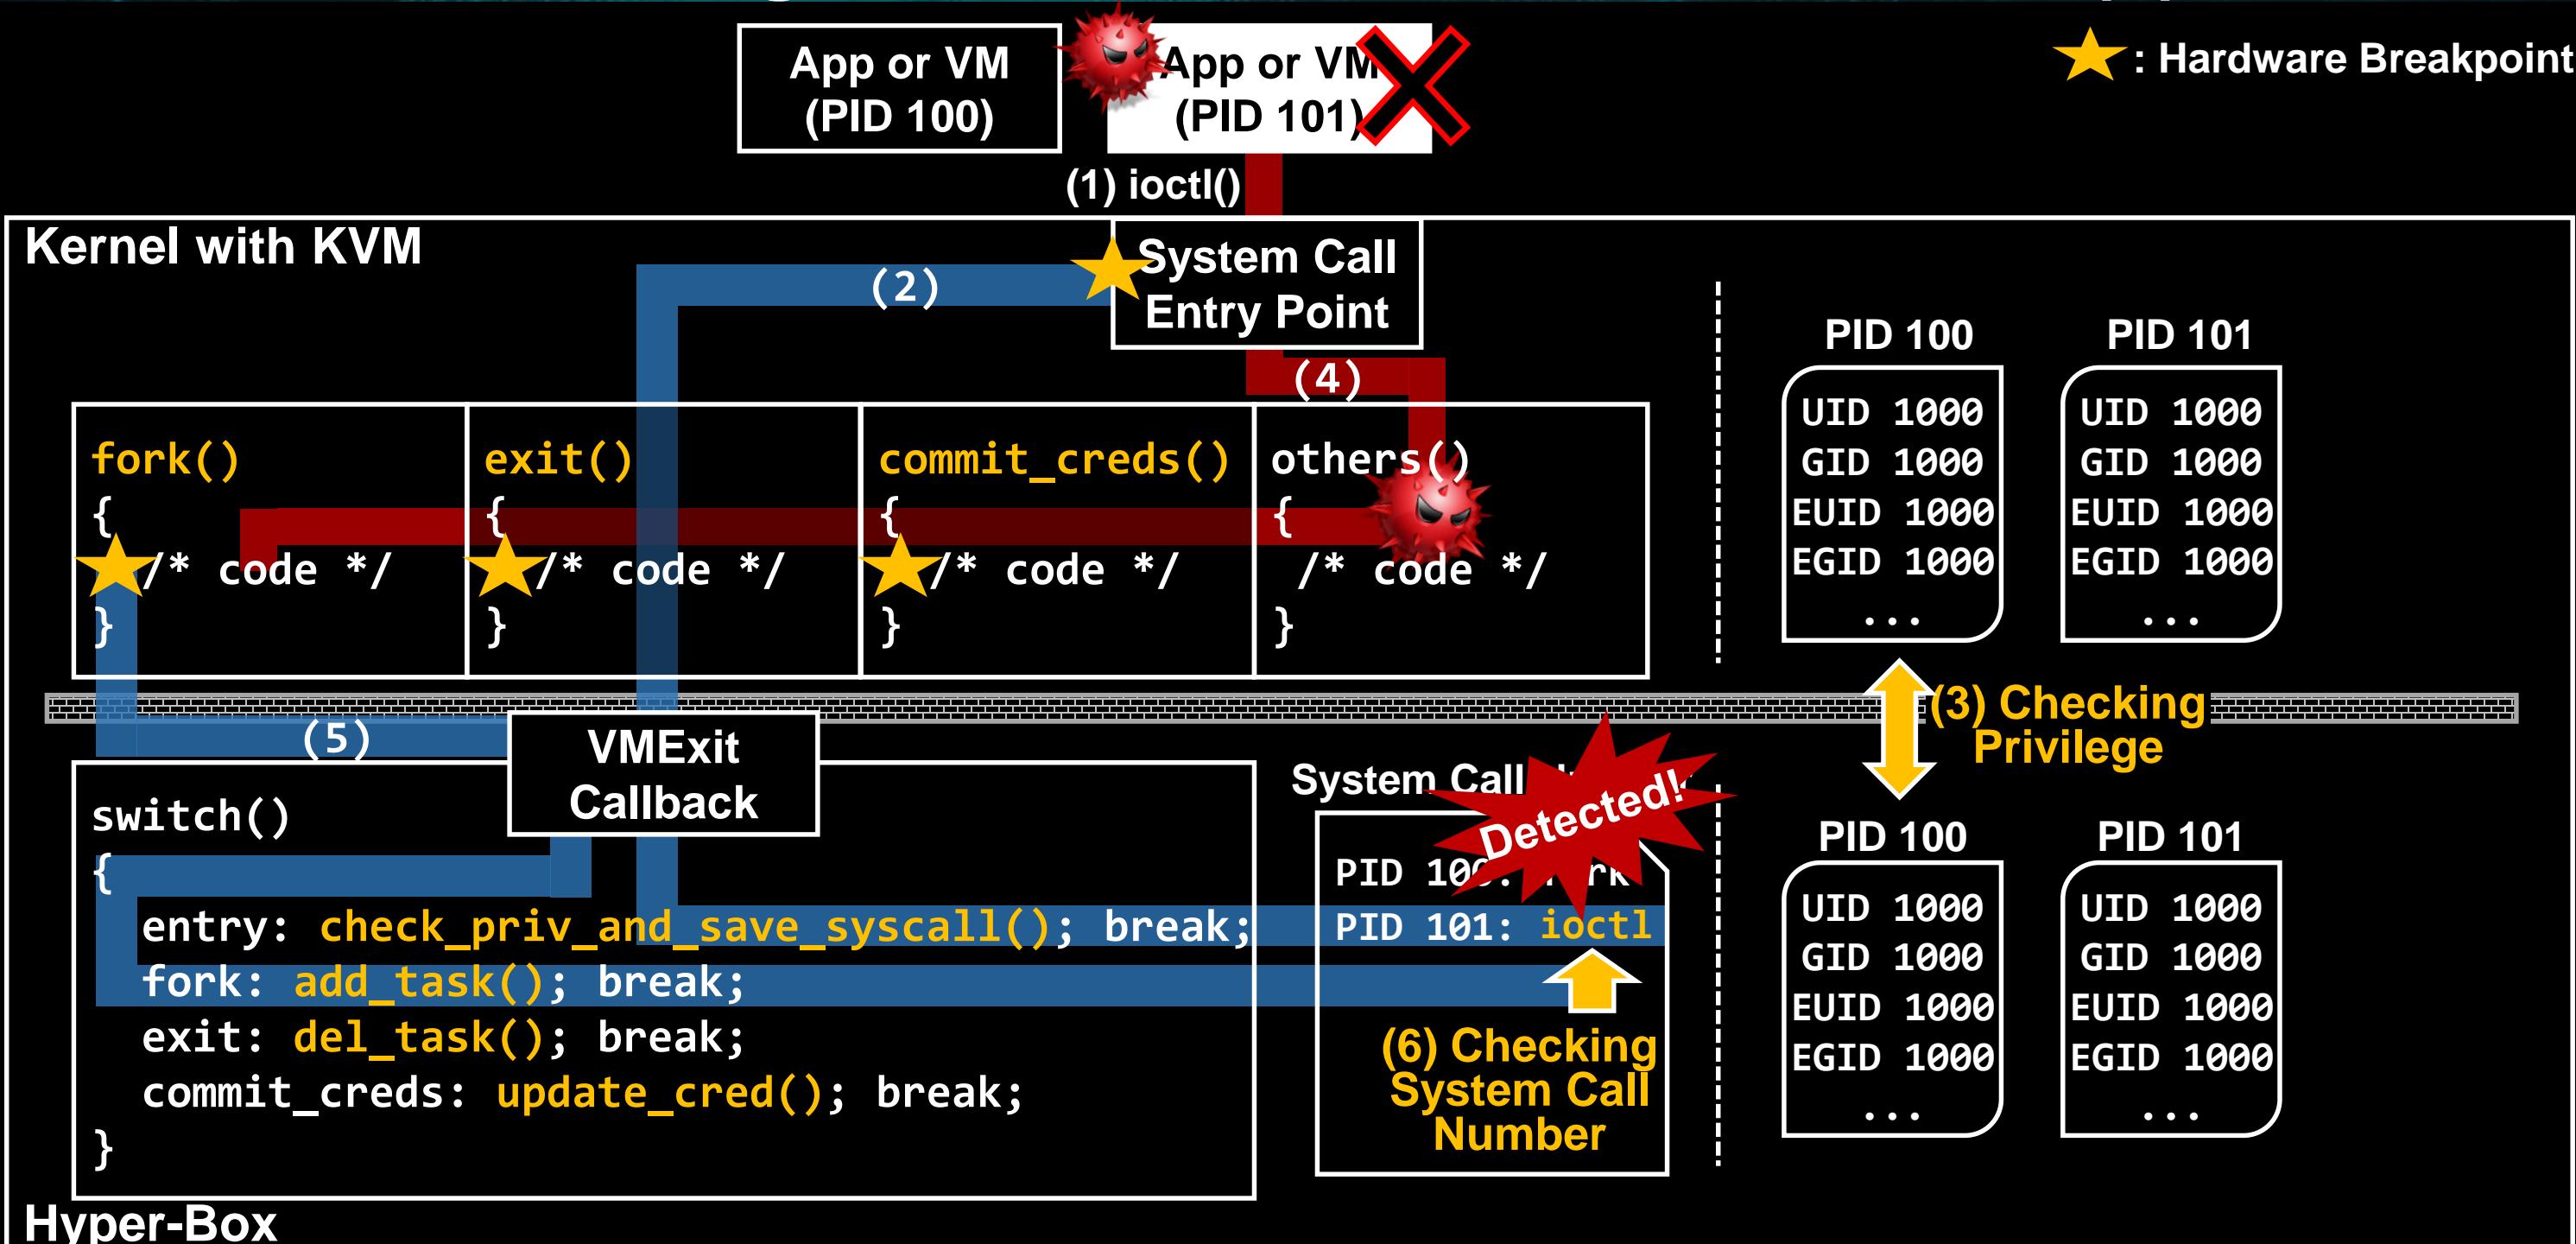
Hyper-Box Design

Preventing Unauthorized Process Creation (1)



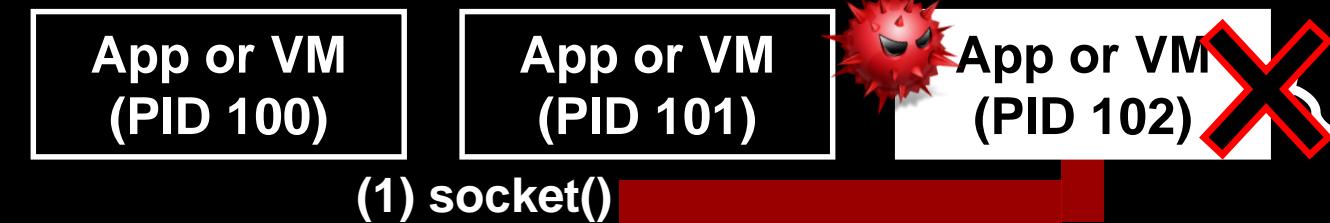
Hyper-Box Design

Preventing Unauthorized Process Creation (2)



Hyper-Box Design

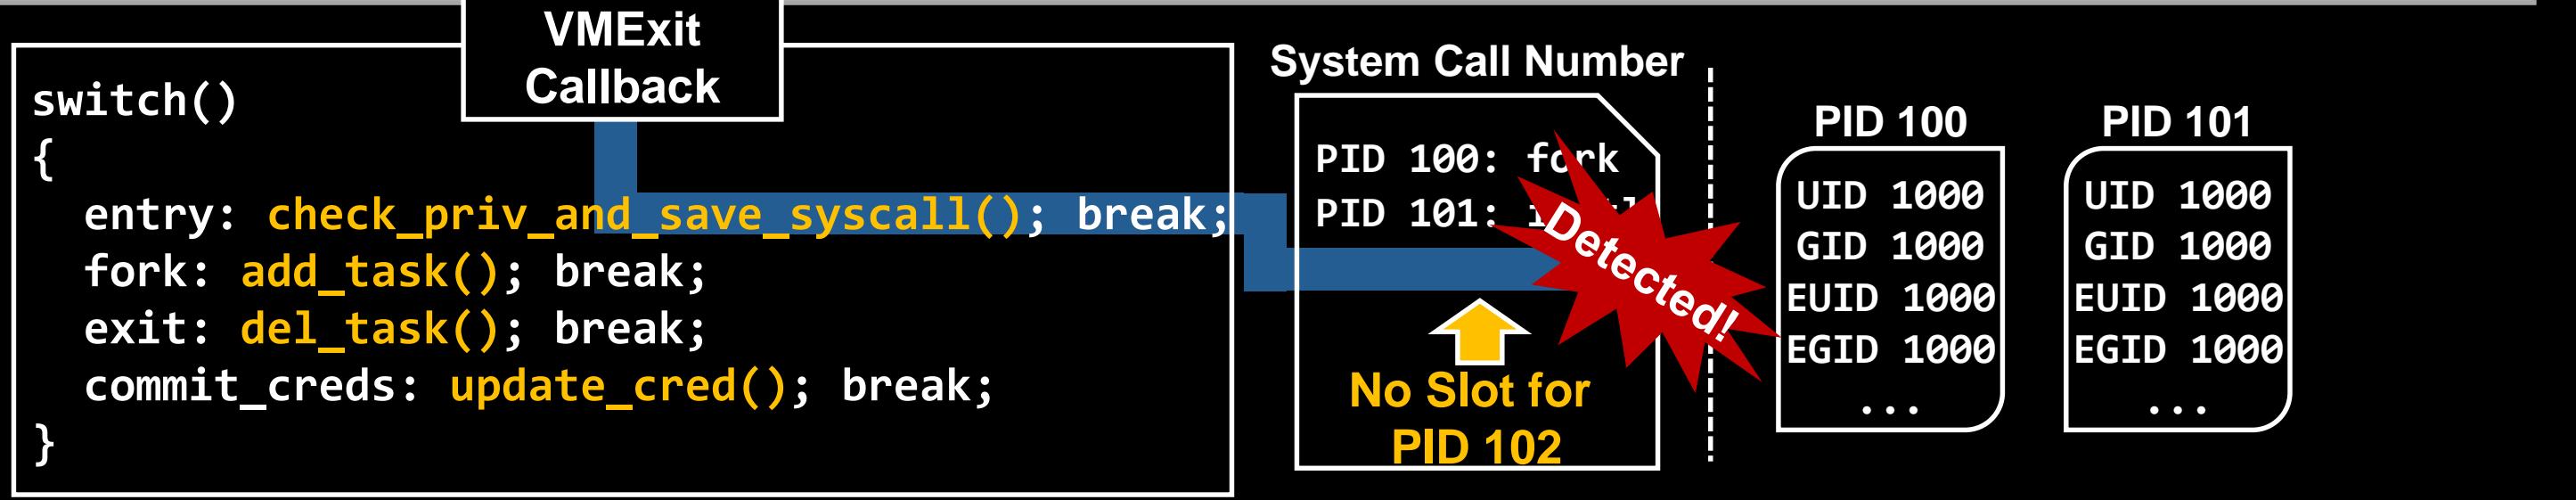
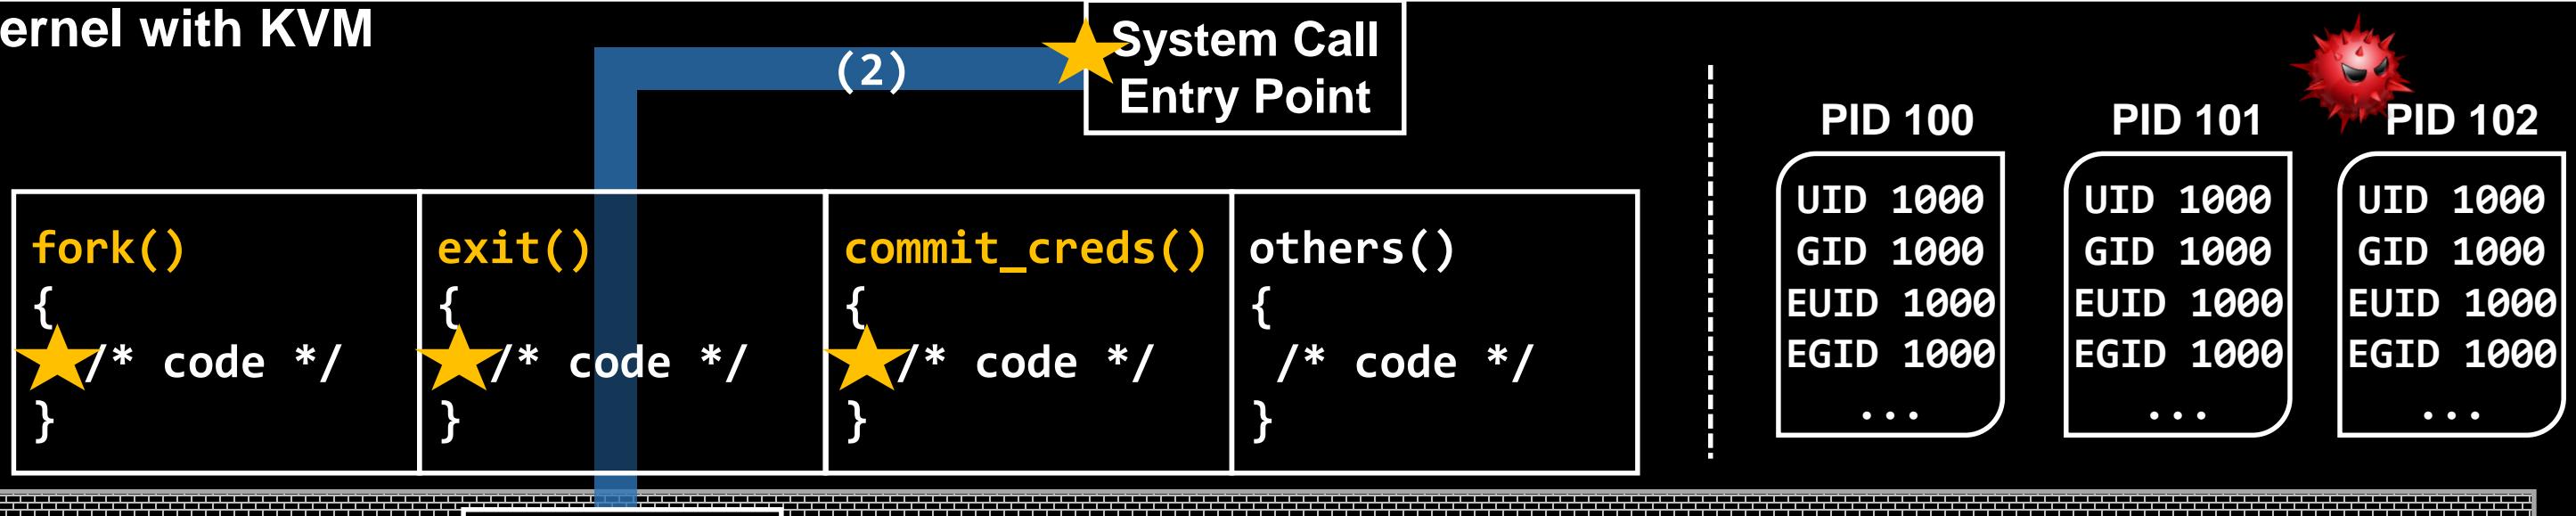
Preventing Unauthorized Process Creation (3)



 : Hardware Breakpoint

This process was created
by detouring H/W breakpoints

Kernel with KVM



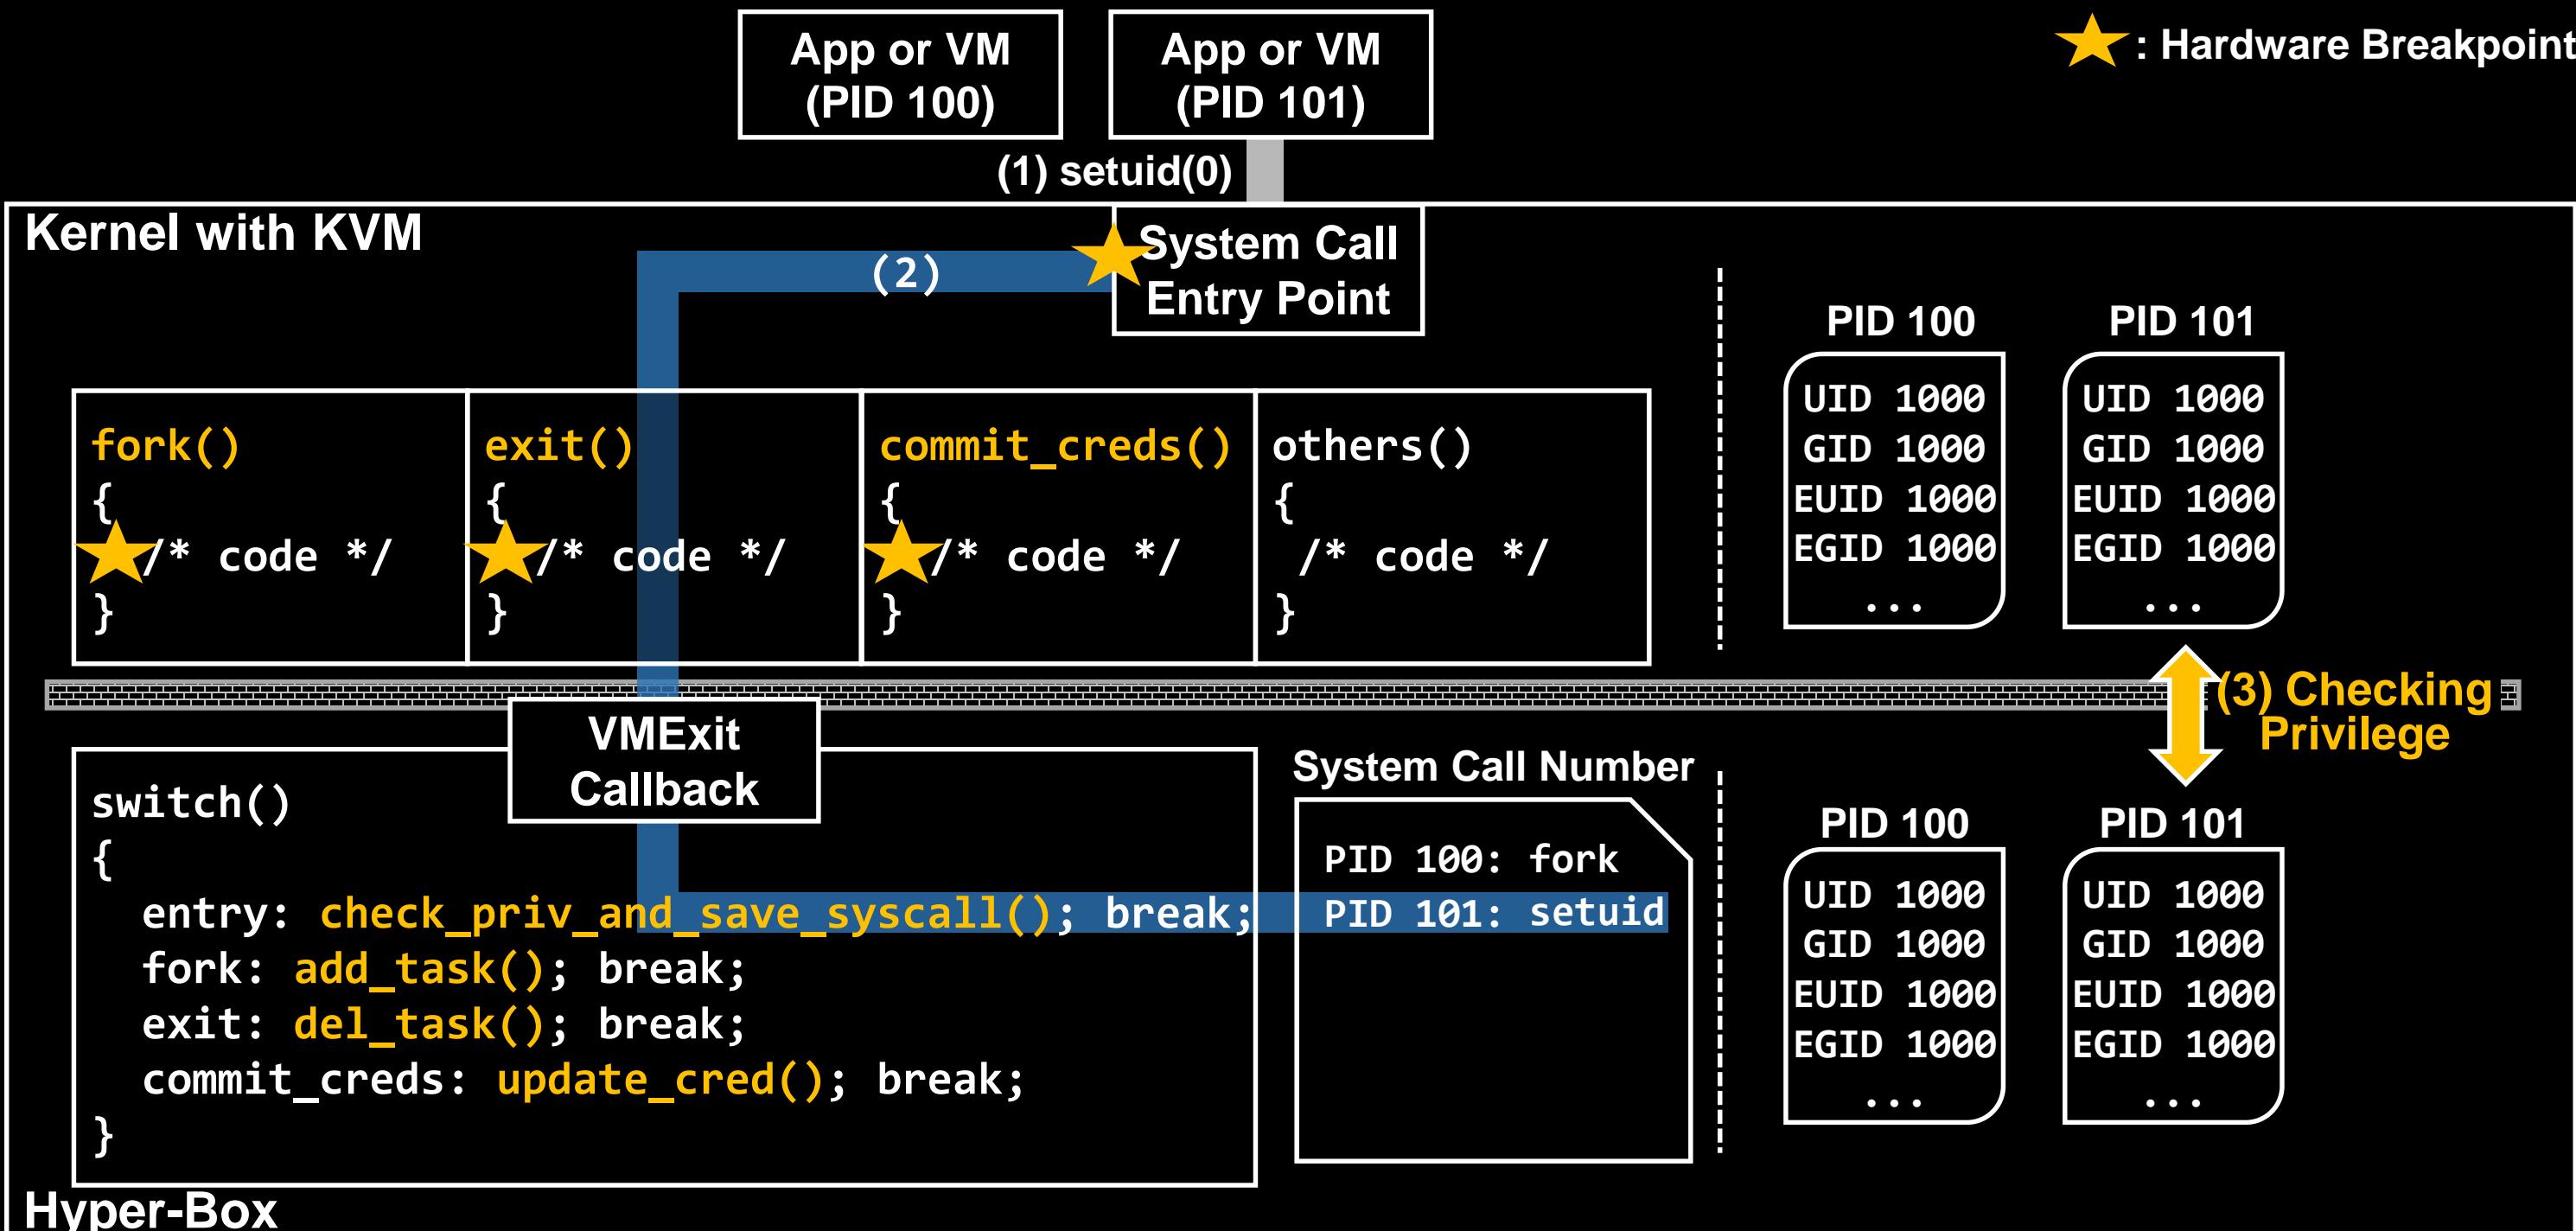
Hyper-Box Design

Preventing Unauthorized Process Creation (4)

- Some processes are not allowed to create a child process
 - QEMU process cannot create it for security reasons
 - Hyper-box has a disallowed list that contains process names
 - It checks the list for every system call and blocks the request
- Some kernel drivers load other modules with a kernel thread
 - They call `call_usermodehelper()`, and the function creates a kernel thread
 - The kernel thread executes modprobe (/usr/bin/kmod) with kernel API
 - It means the system call number is unknown, but modprobe is created!
 - Hyper-box has a workaround feature that allows the kernel thread to execute a process
 - But, the feature is only for logging and not recommended
 - **Load all necessary modules before Alcatraz or include modules to the kernel!**

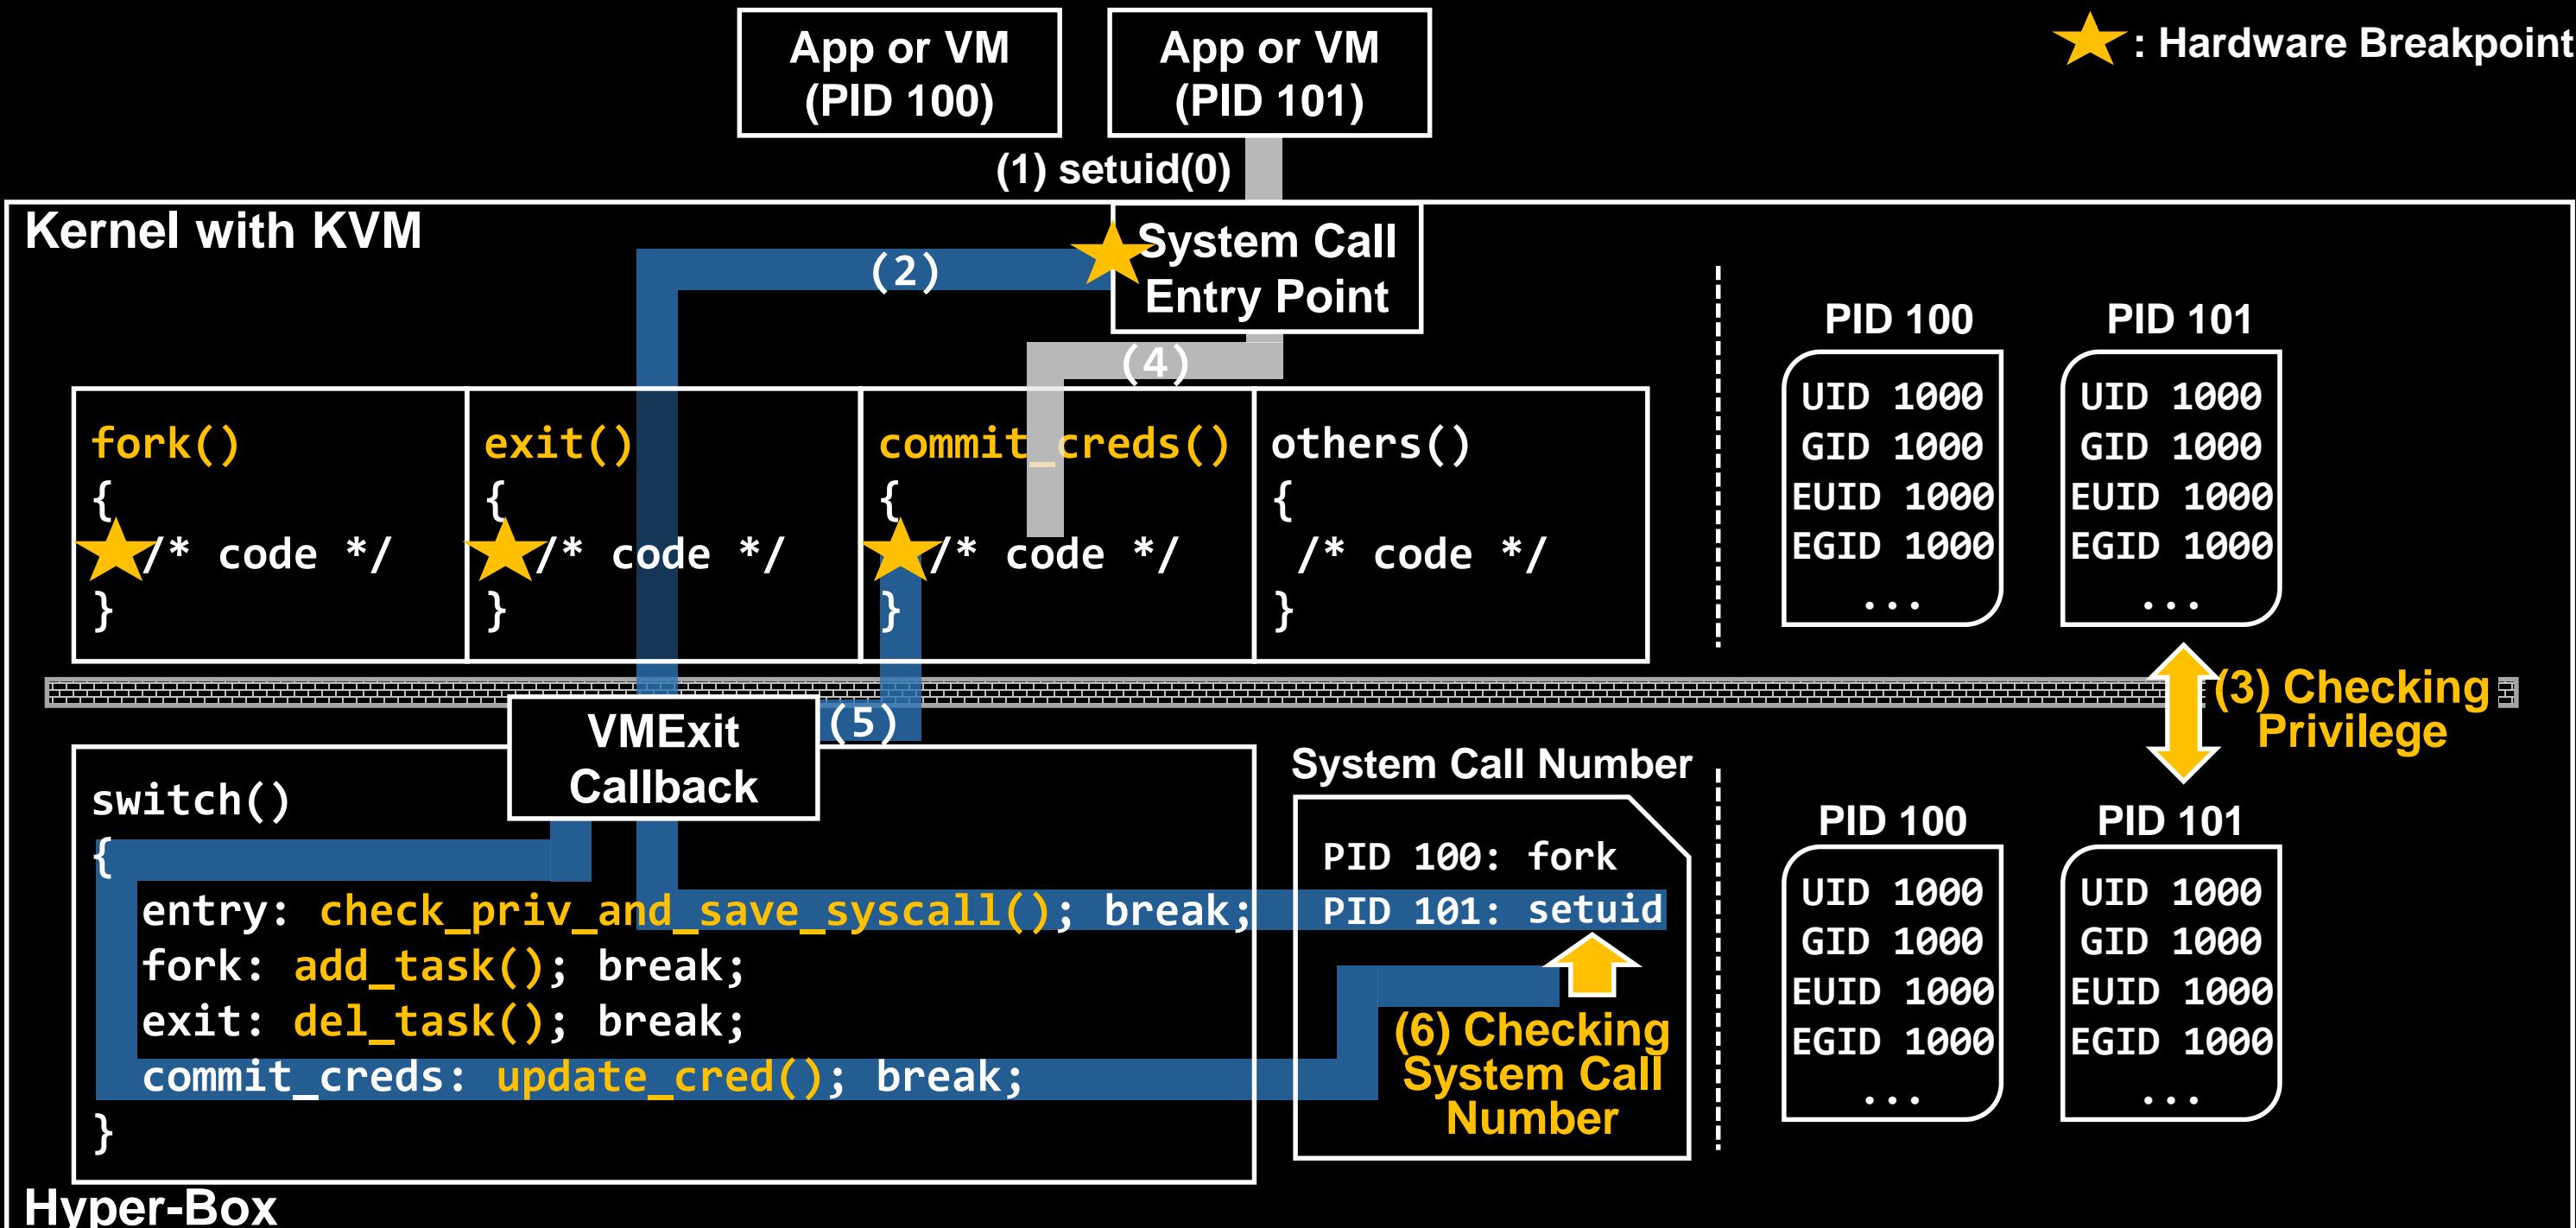
Hyper-Box Design

Preventing Unauthorized Privilege Escalation (1)



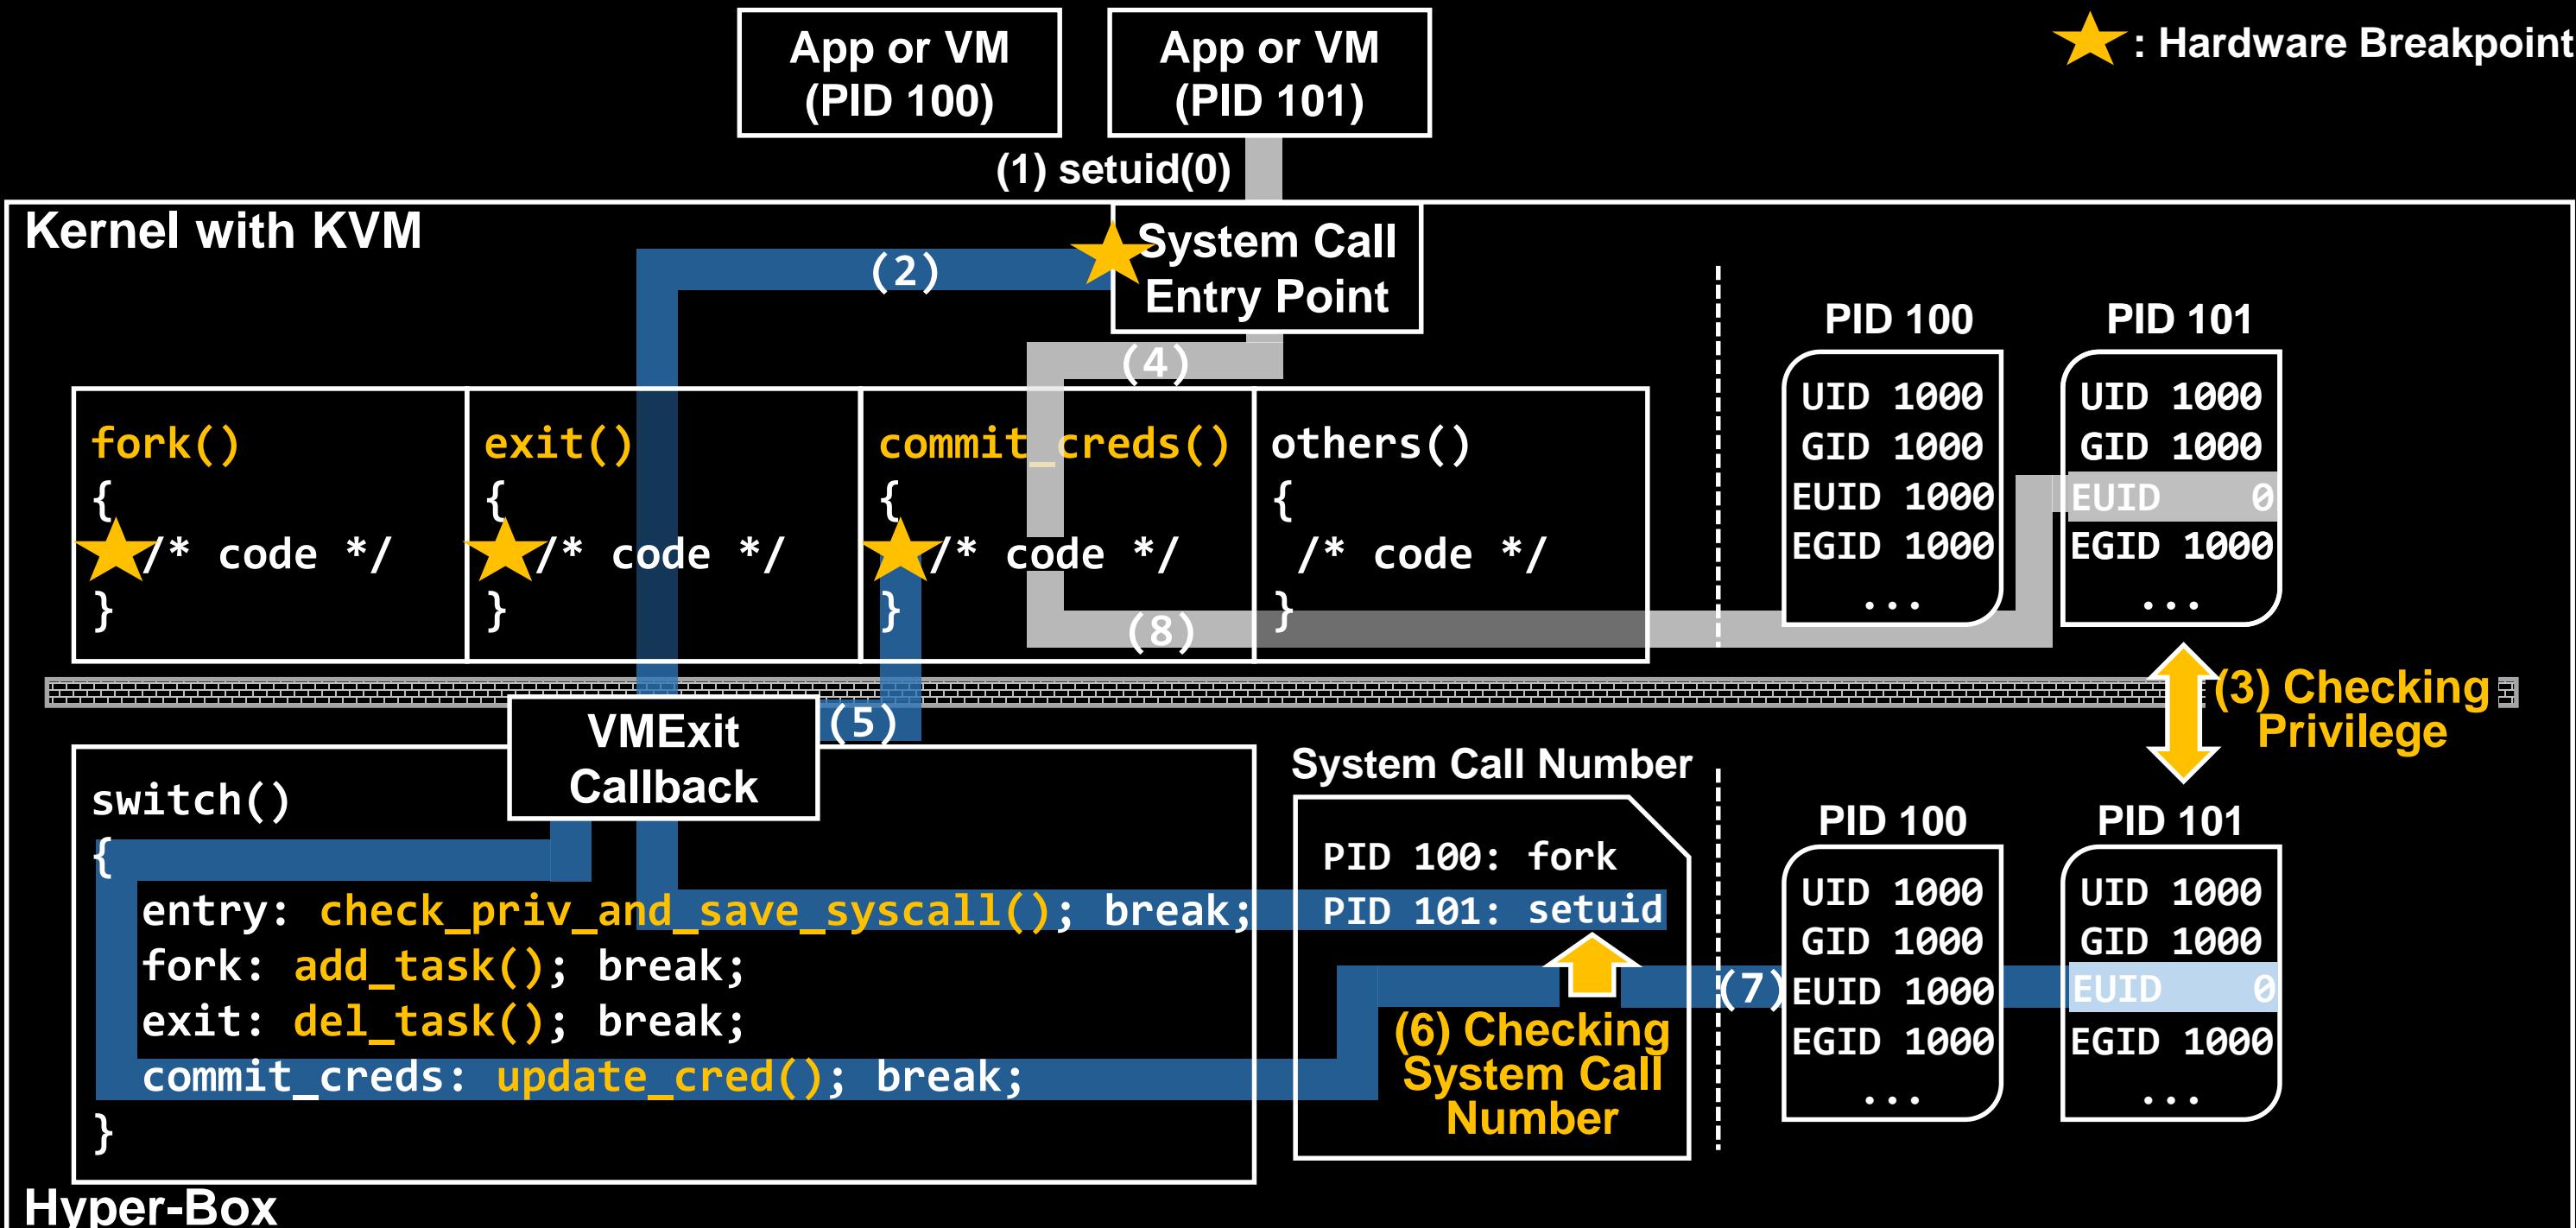
Hyper-Box Design

Preventing Unauthorized Privilege Escalation (1)



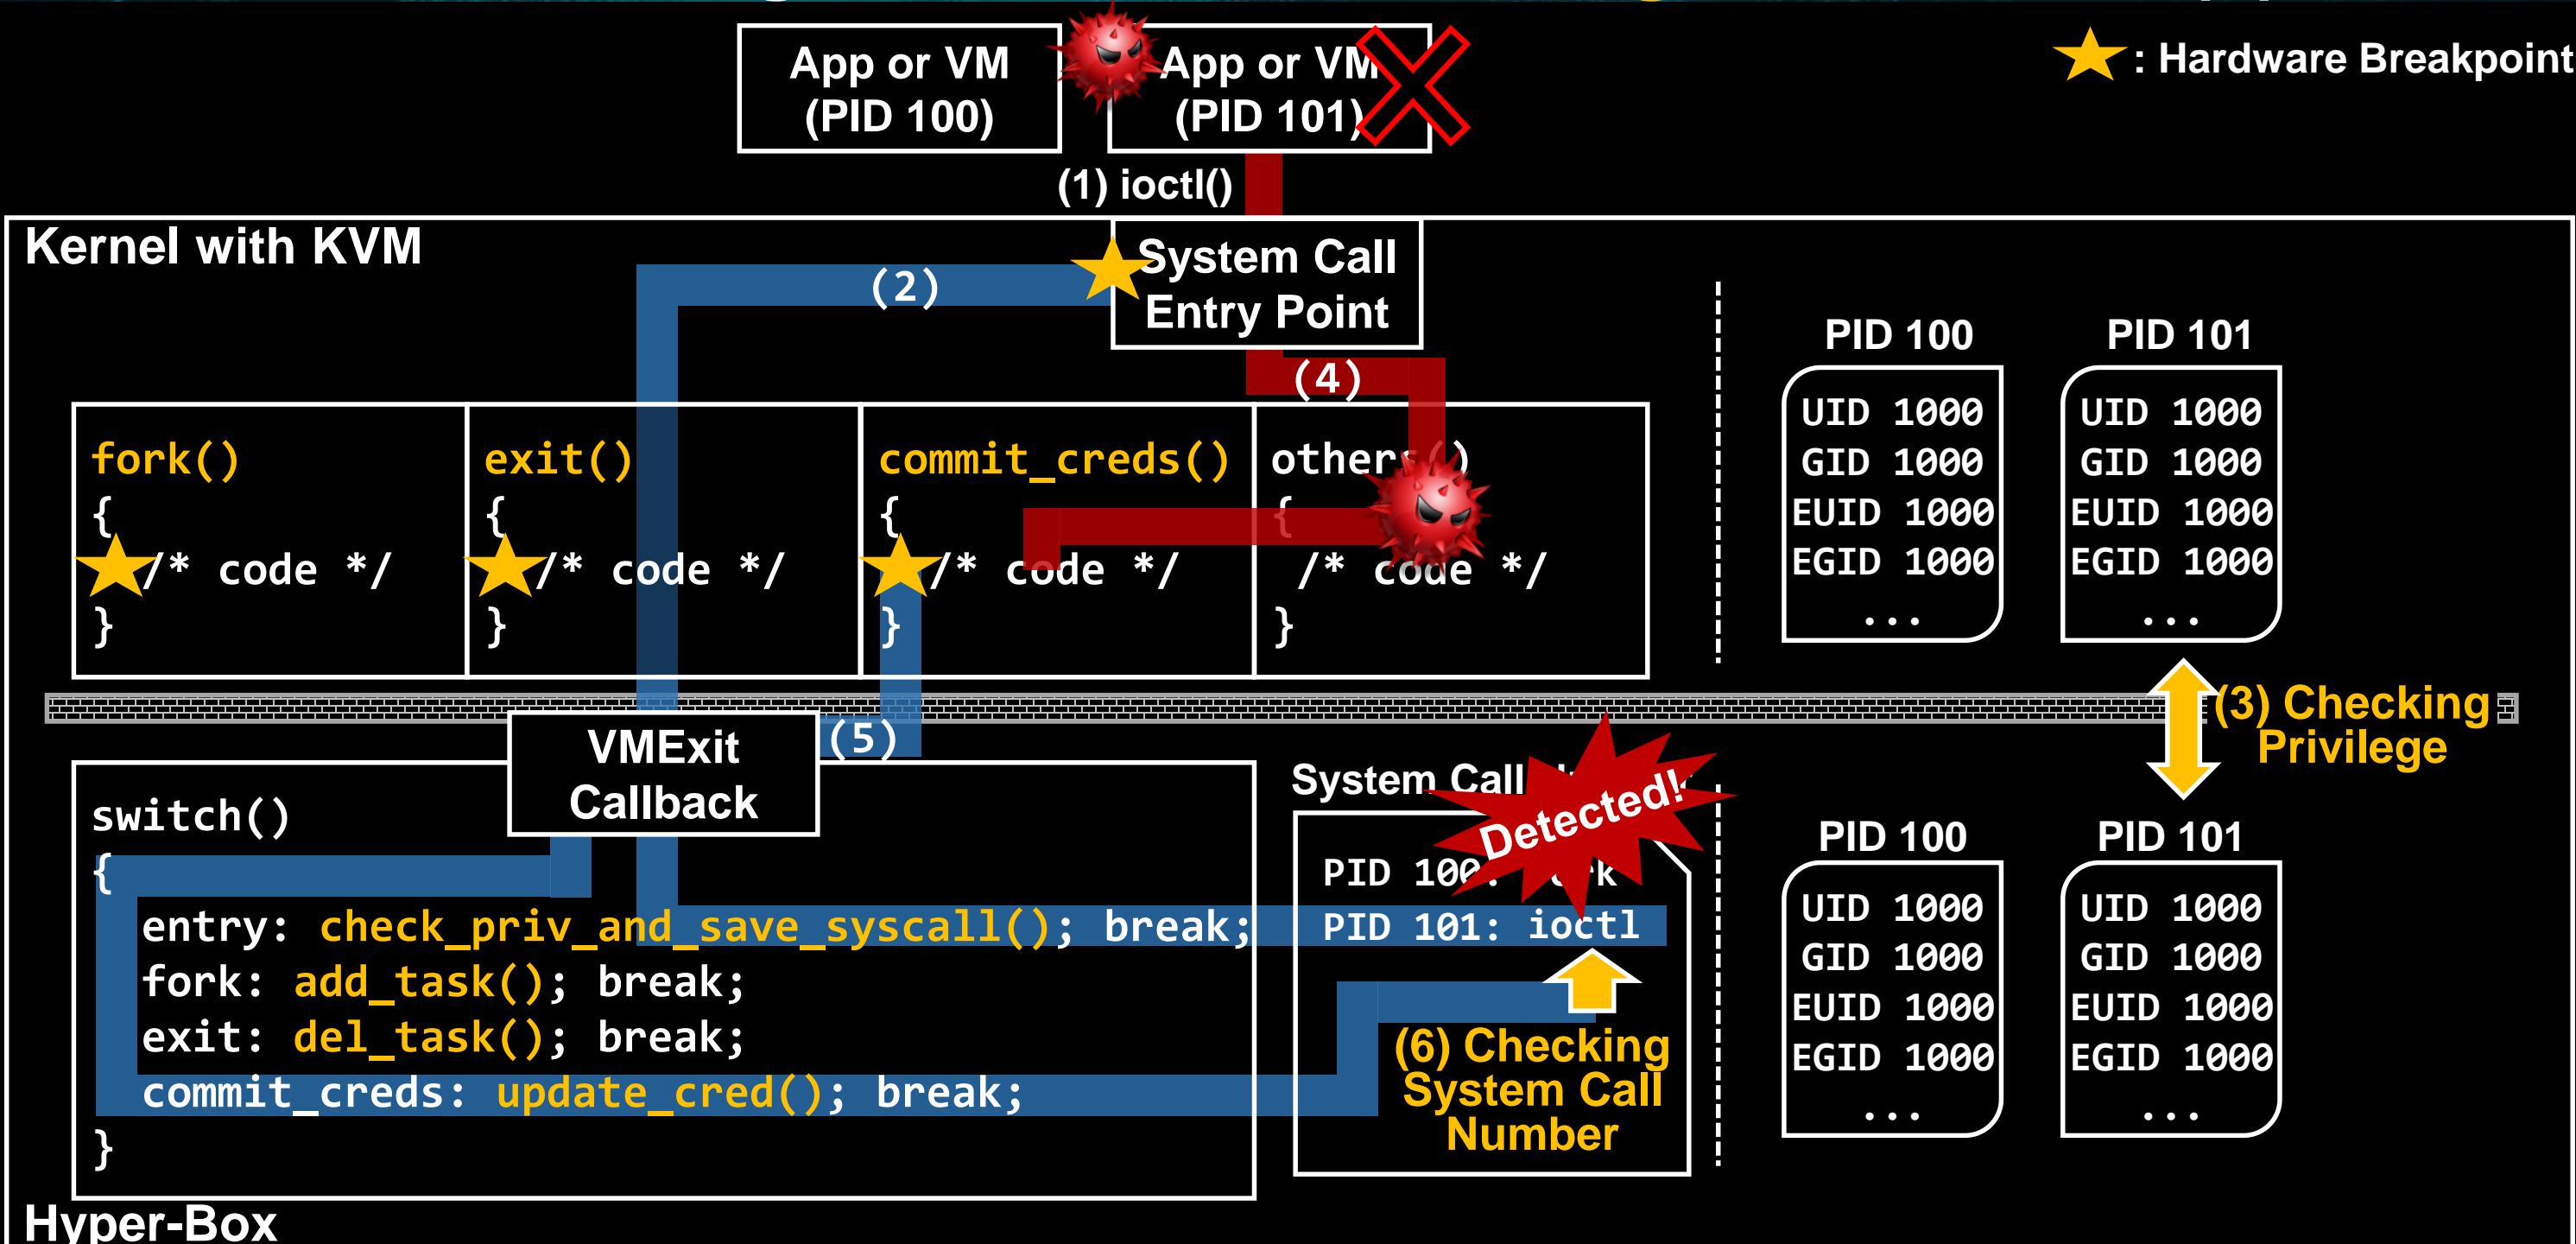
Hyper-Box Design

Preventing Unauthorized Privilege Escalation (1)



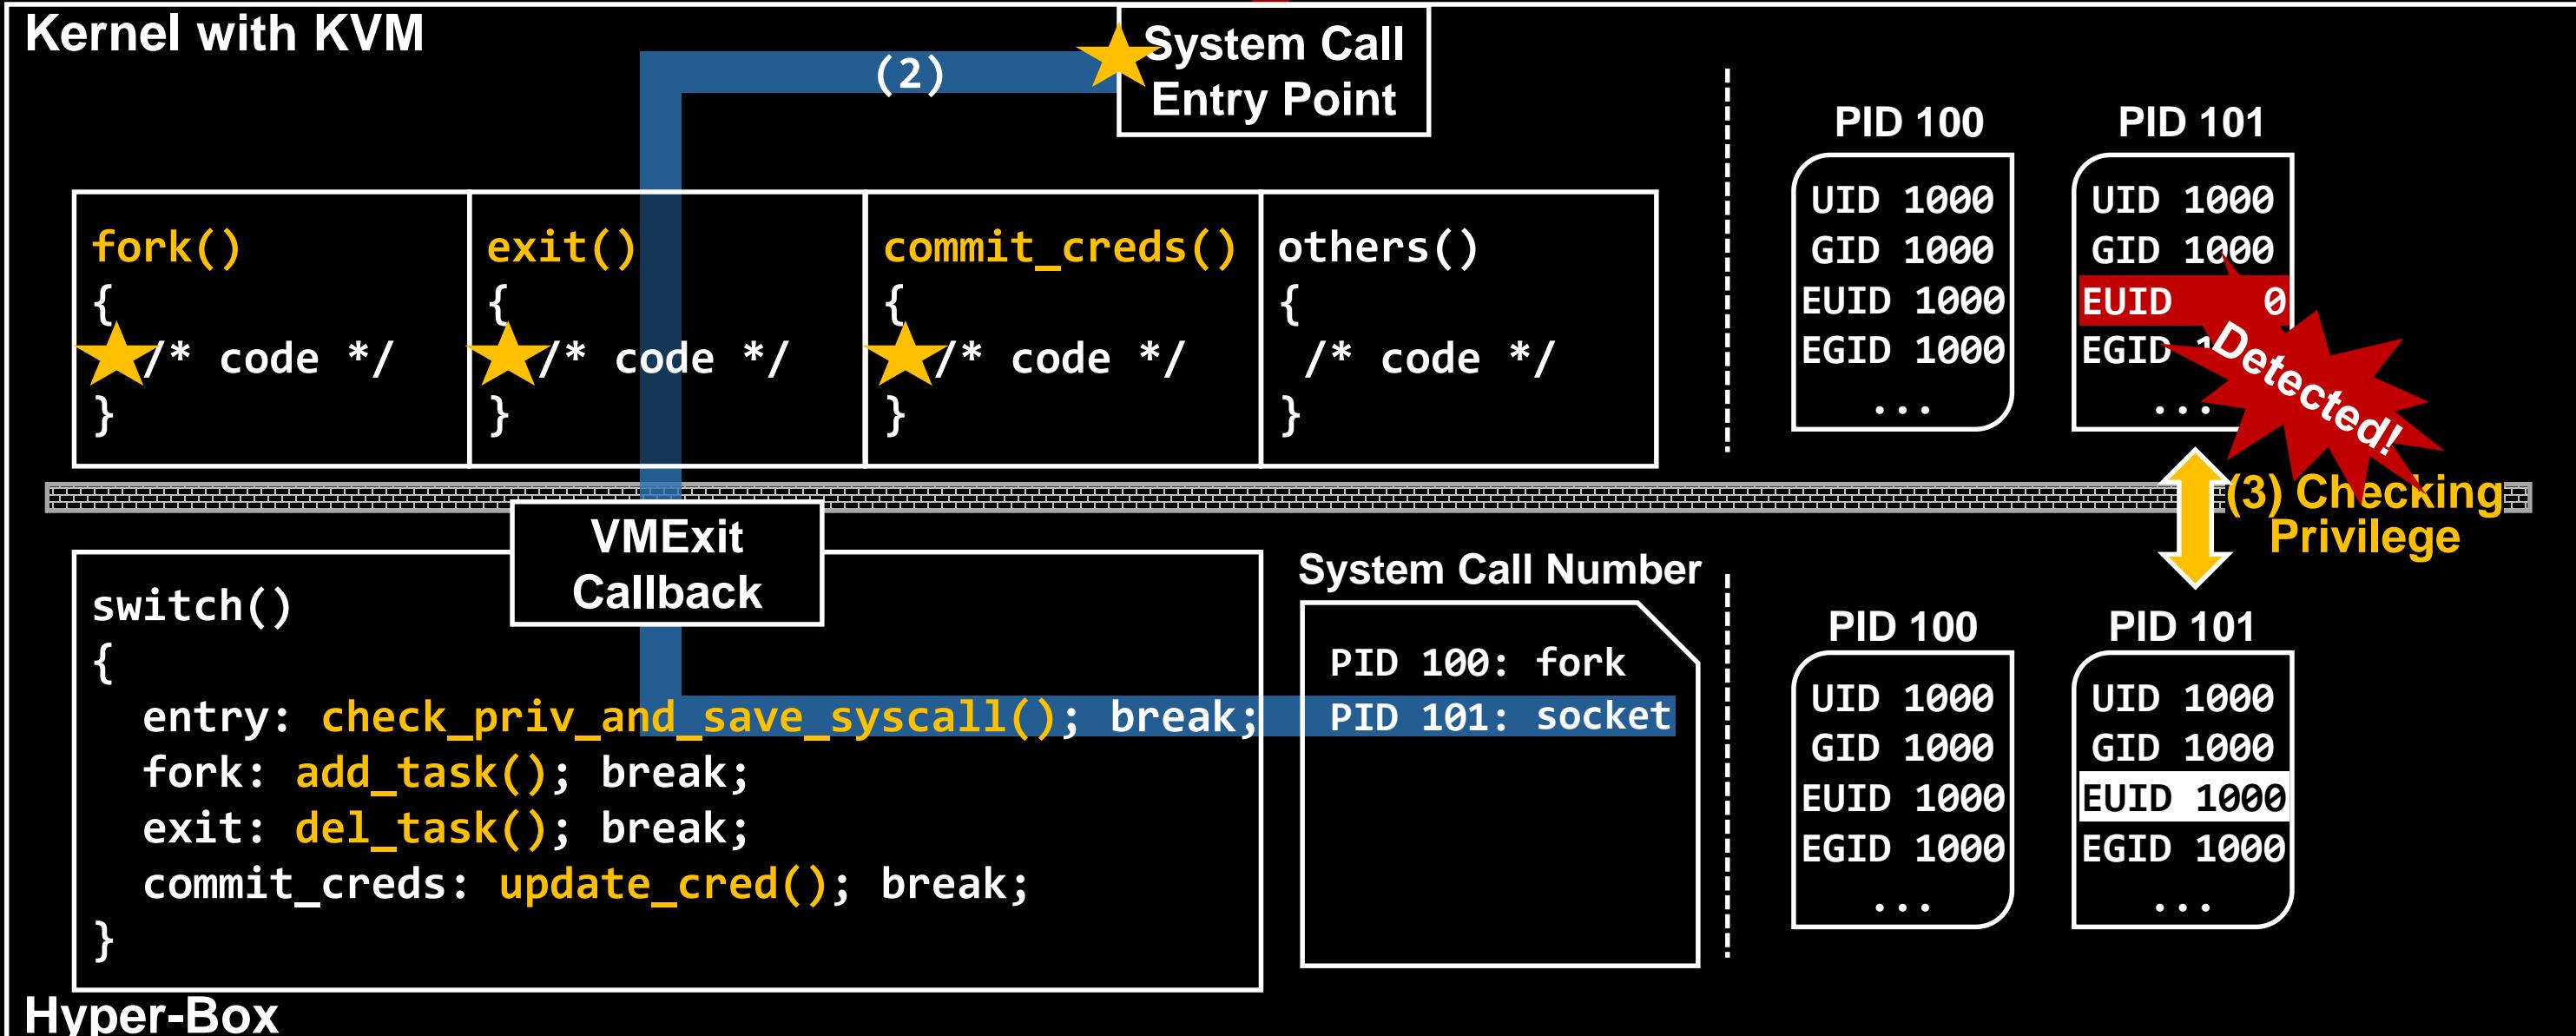
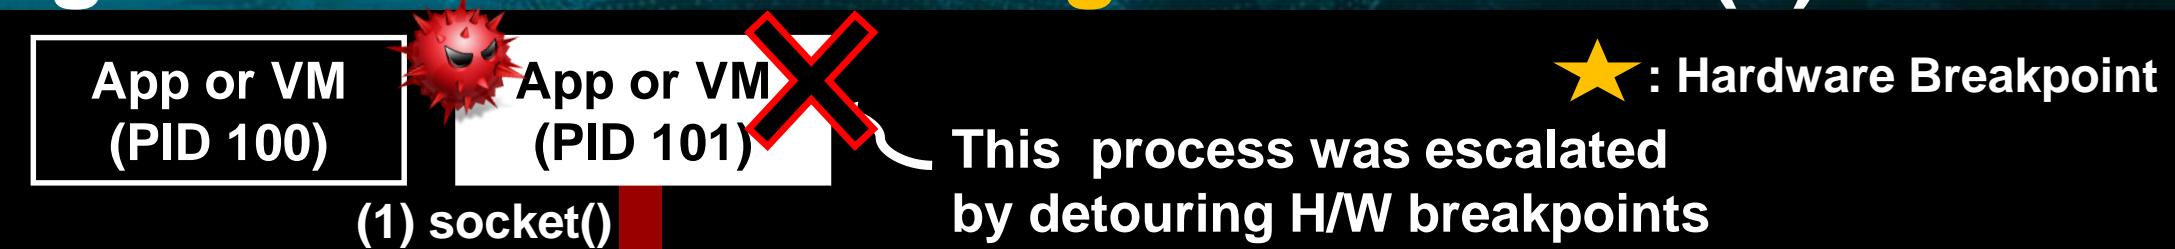
Hyper-Box Design

Preventing Unauthorized Privilege Escalation (2)



Hyper-Box Design

Preventing Unauthorized Privilege Escalation (3)



Hyper-Box Design

Emulating VMX instructions – Only KVM

Kernel with KVM

Ring 0 → Ring -1 (VMXON of KVM)

(1) VMXON, VMXOFF

(3) VMExit

KVM (Guest Hypervisor)
0xAAAABB: VMExitHandler
/* Processing Events */

VM 1

VM 2

VMCS 1
VM_HOST_RIP: 0xAAAABB
VM_HOST_RSP: 0x11111111

VMCS 2
VM_HOST_RIP: 0xAAAABB
VM_HOST_RSP: 0x22222222

(2) VMPTLD, VMPTRST, VMCLEAR, VMREAD, VMWRITE, VMLAUNCH,
VMRESUME, INVEPT, INVVPID, INVPCID, etc.

VMCS 0 (Sandbox)

VM_HOST_RIP : 0xEEEEFFFF
VM_HOST_RSP : 0x0000FFFF
VM_GUEST_RIP: 0x00000000
VM_GUEST_RSP: 0x00000000

VMCS Data List

VMCS 1-RIP: 0xAAAABB
VMCS 1-RSP: 0x11111111
VMCS 2-RIP: 0xAAAABB
VMCS 2-RSP: 0x22222222

Hyper-Box (Host Hypervisor)
0xEEEEFFFF: VMExitHandler
/* Processing VMX instructions */

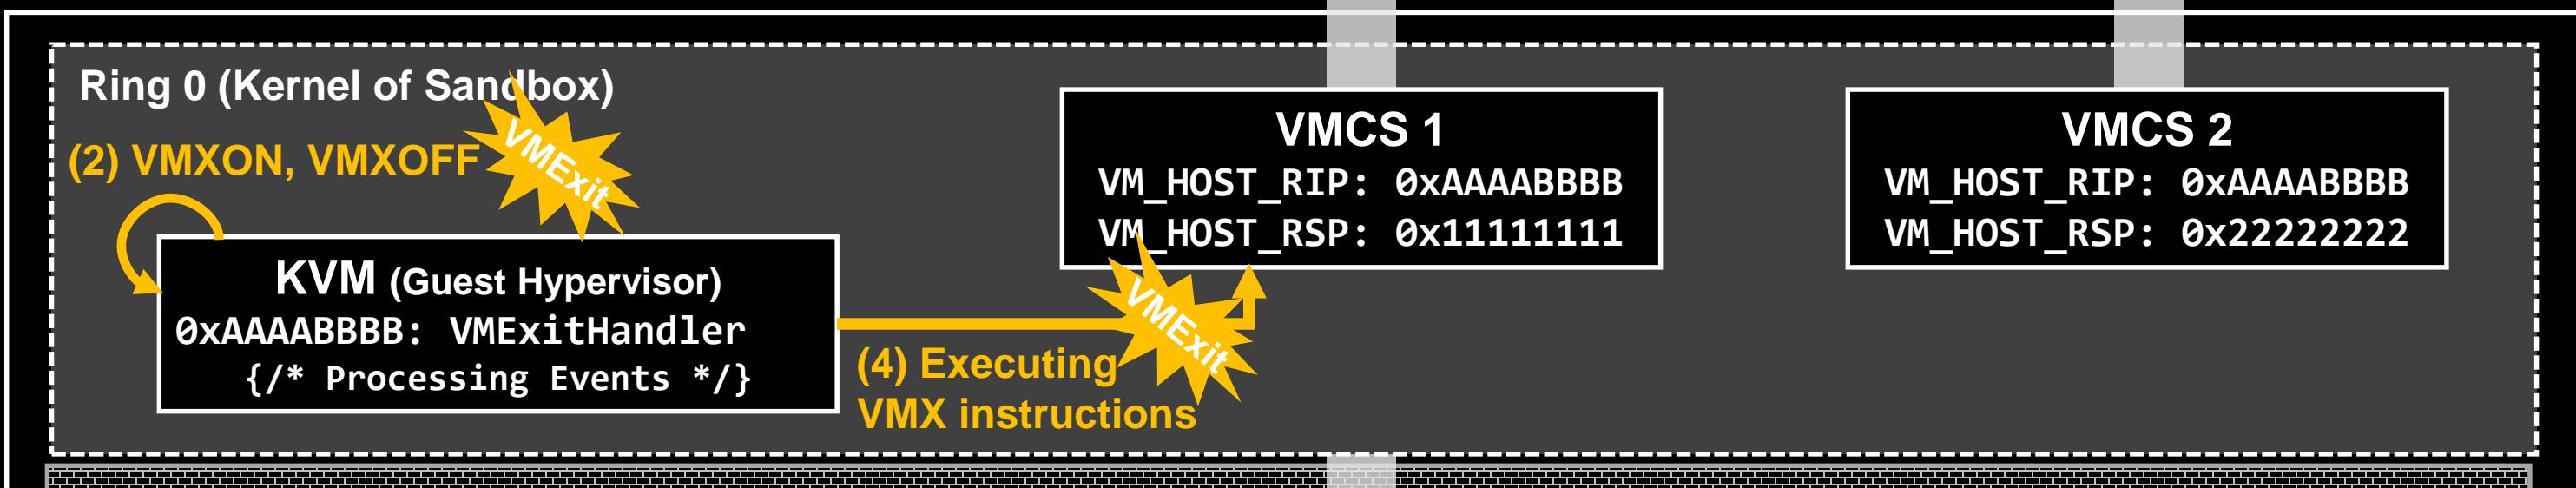
CPU 0

VMCS_PTR: VMCS 1 or 2

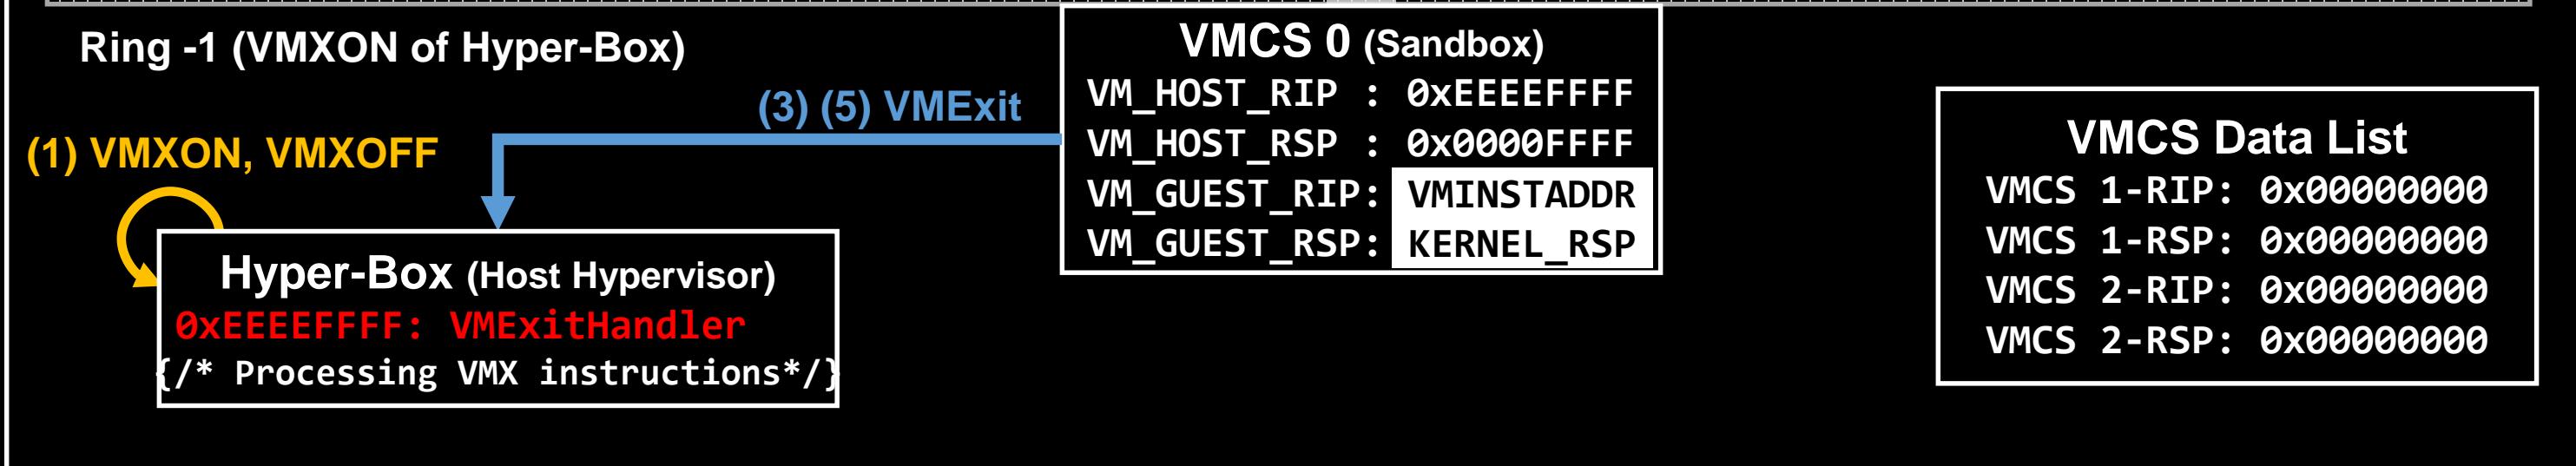
Hyper-Box Design

Emulating VMX instructions – Executing Instructions (1)

Kernel with KVM



Ring -1 (VMXON of Hyper-Box)



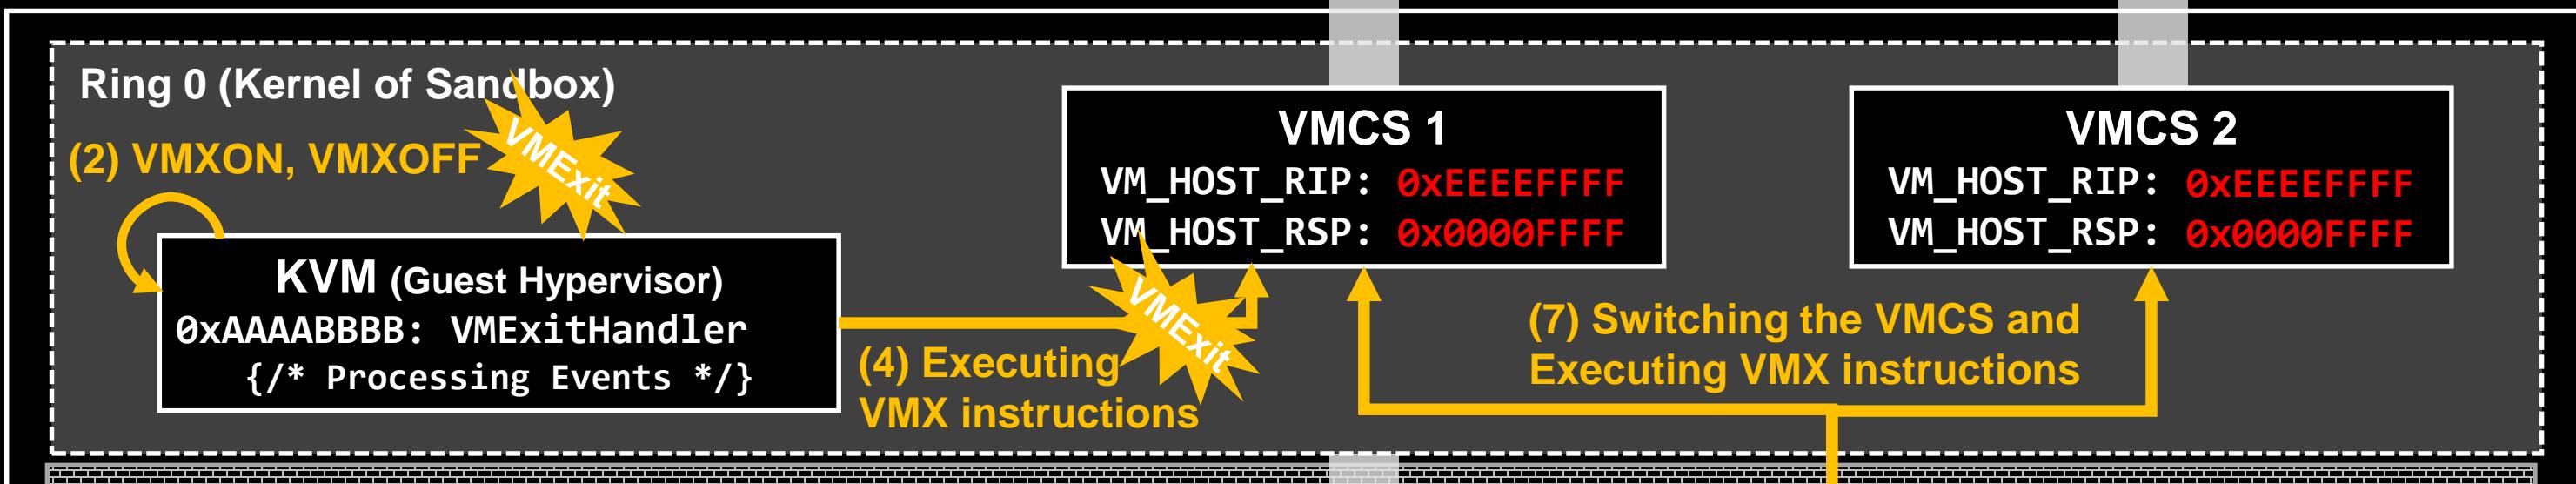
CPU 0

VMCS_PTR: VMCS 0

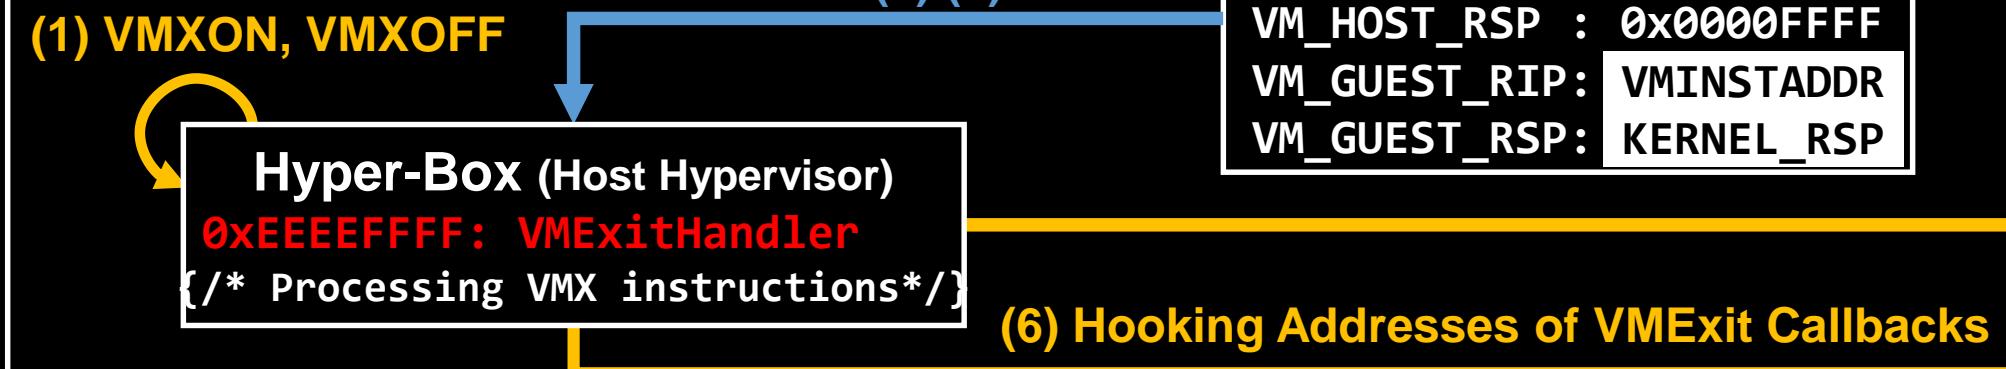
Hyper-Box Design

Emulating VMX instructions – Executing Instructions (2)

Kernel with KVM



Ring -1 (VMXON of Hyper-Box)



CPU 0
VMCS_PTR: VMCS 0 → VMCS 1 or 2

Hyper-Box Design

Emulating VMX instructions – Handling VMExits (1)

Kernel with KVM

Ring 0 (Kernel of Sandbox)

KVM (Guest Hypervisor)
`0xAAAABBBC: VMExitHandler`
`{/* Processing Events */}`

VM 1

VMCS 1

VM_HOST_RIP: `0xEEEEFFFF`
 VM_HOST_RSP: `0x0000FFFF`

VM 2

VMCS 2

VM_HOST_RIP: `0xEEEEFFFF`
 VM_HOST_RSP: `0x0000FFFF`

(1) VMExit

Ring -1 (VMXON of Hyper-Box)

Hyper-Box (Host Hypervisor)
`0xEEEEFFFF: VMExitHandler`
`{/* Processing VMX instructions */}`

VMCS 0 (Sandbox)

VM_HOST_RIP : `0xEEEEFFFF`
 VM_HOST_RSP : `0x0000FFFF`
 VM_GUEST_RIP: `VMINSTADDR`
 VM_GUEST_RSP: `KERNEL_RSP`

VMCS Data List

VMCS 1-RIP: `0xAAAABBBC`
 VMCS 1-RSP: `0x11111111`
 VMCS 2-RIP: `0xAAAABBBC`
 VMCS 2-RSP: `0x22222222`

(2) Retrieving Original RIP (VMExitHandler of KVM) and RSP

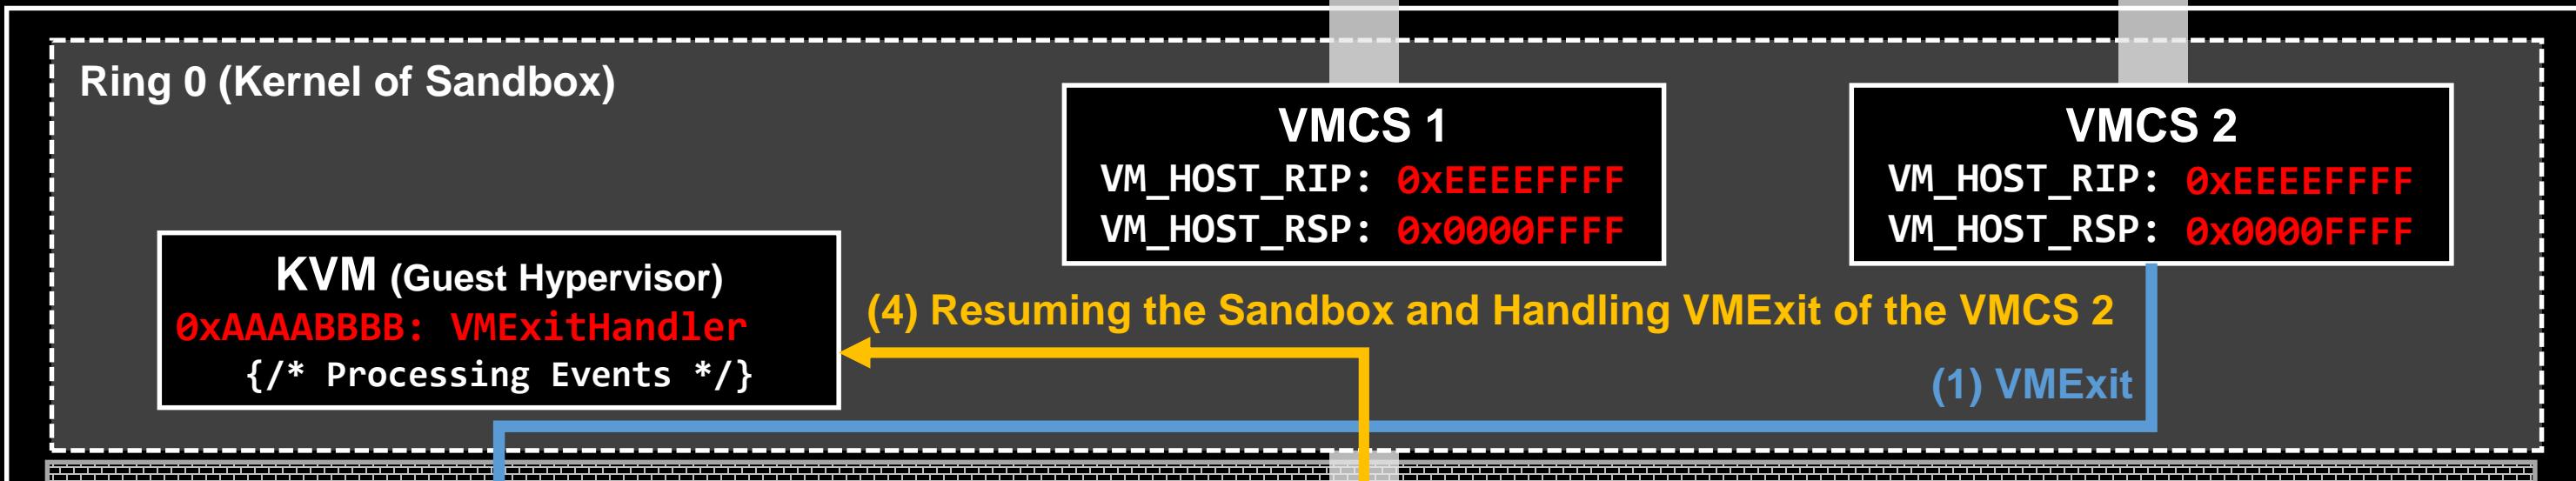
CPU 0

VMCS_PTR:VMCS 2 → VMCS 0

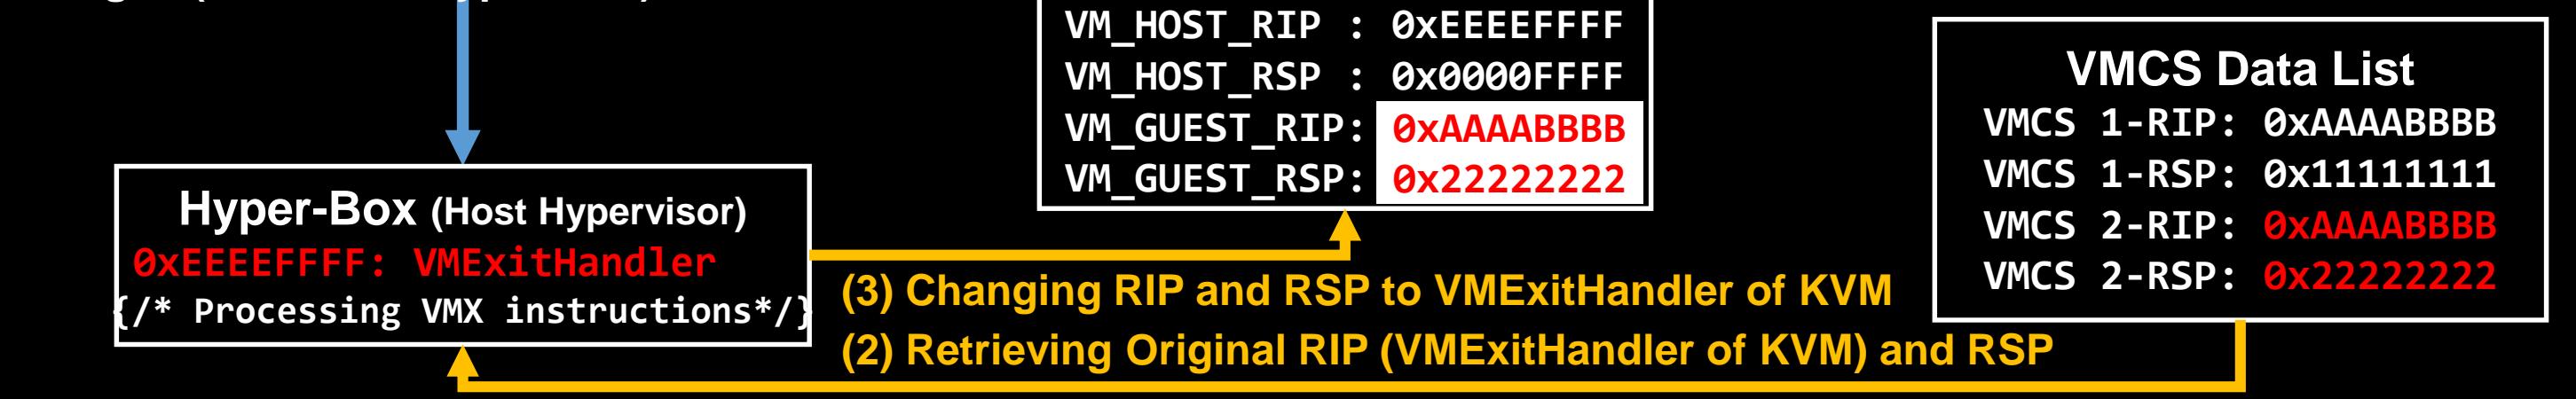
Hyper-Box Design

Emulating VMX instructions – Handling VMExits (2)

Kernel with KVM



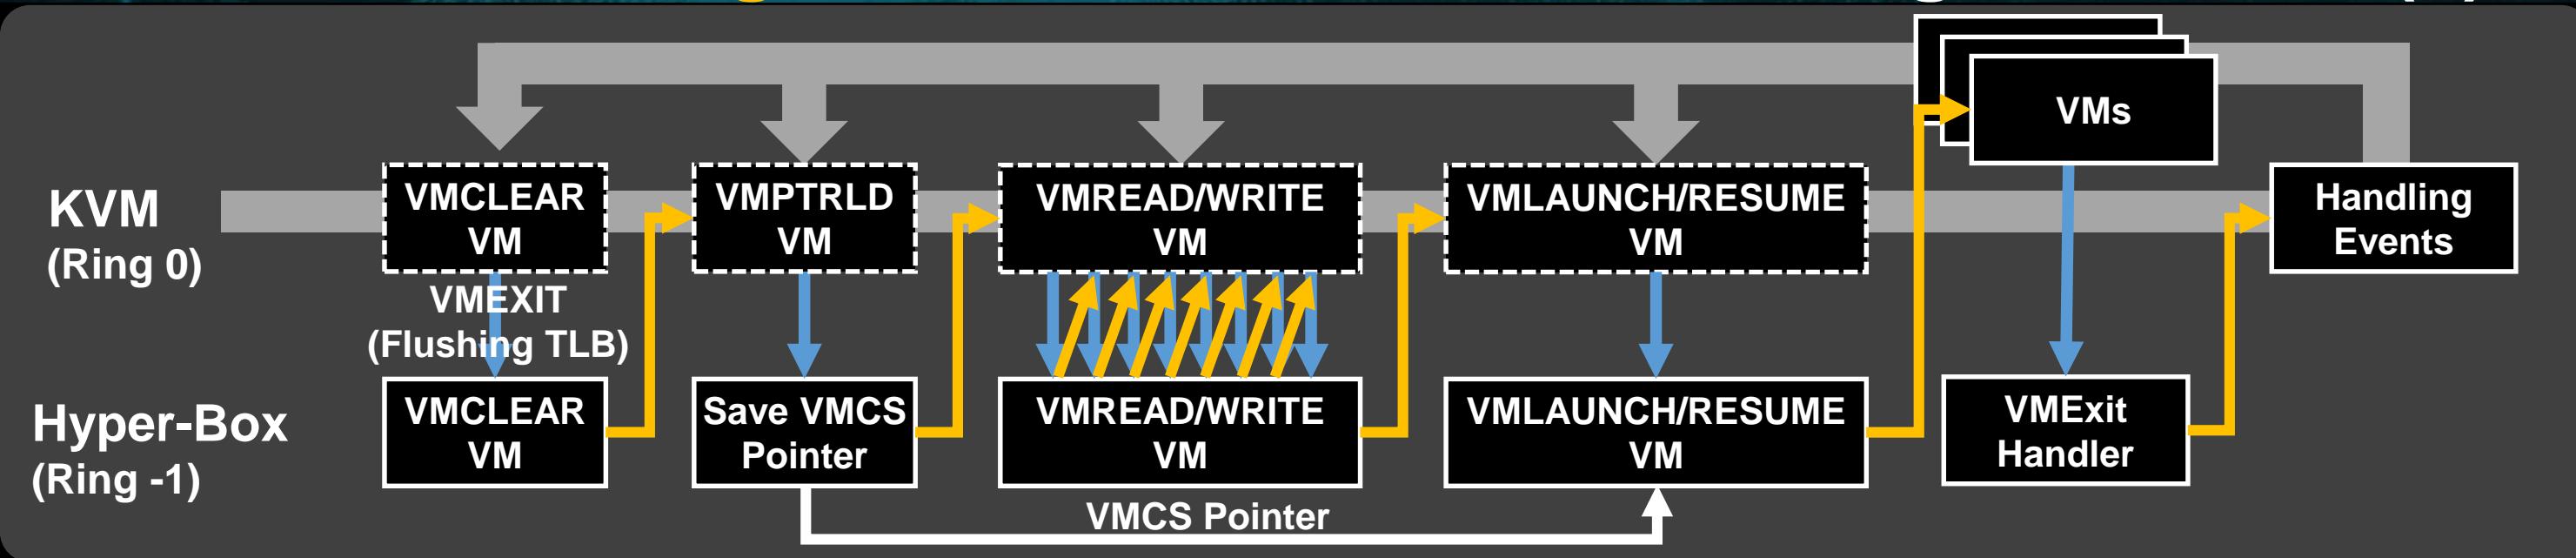
Ring -1 (VMXON of Hyper-Box)



CPU 0
`VMCS_PTR:VMCS 2 → VMCS 0`

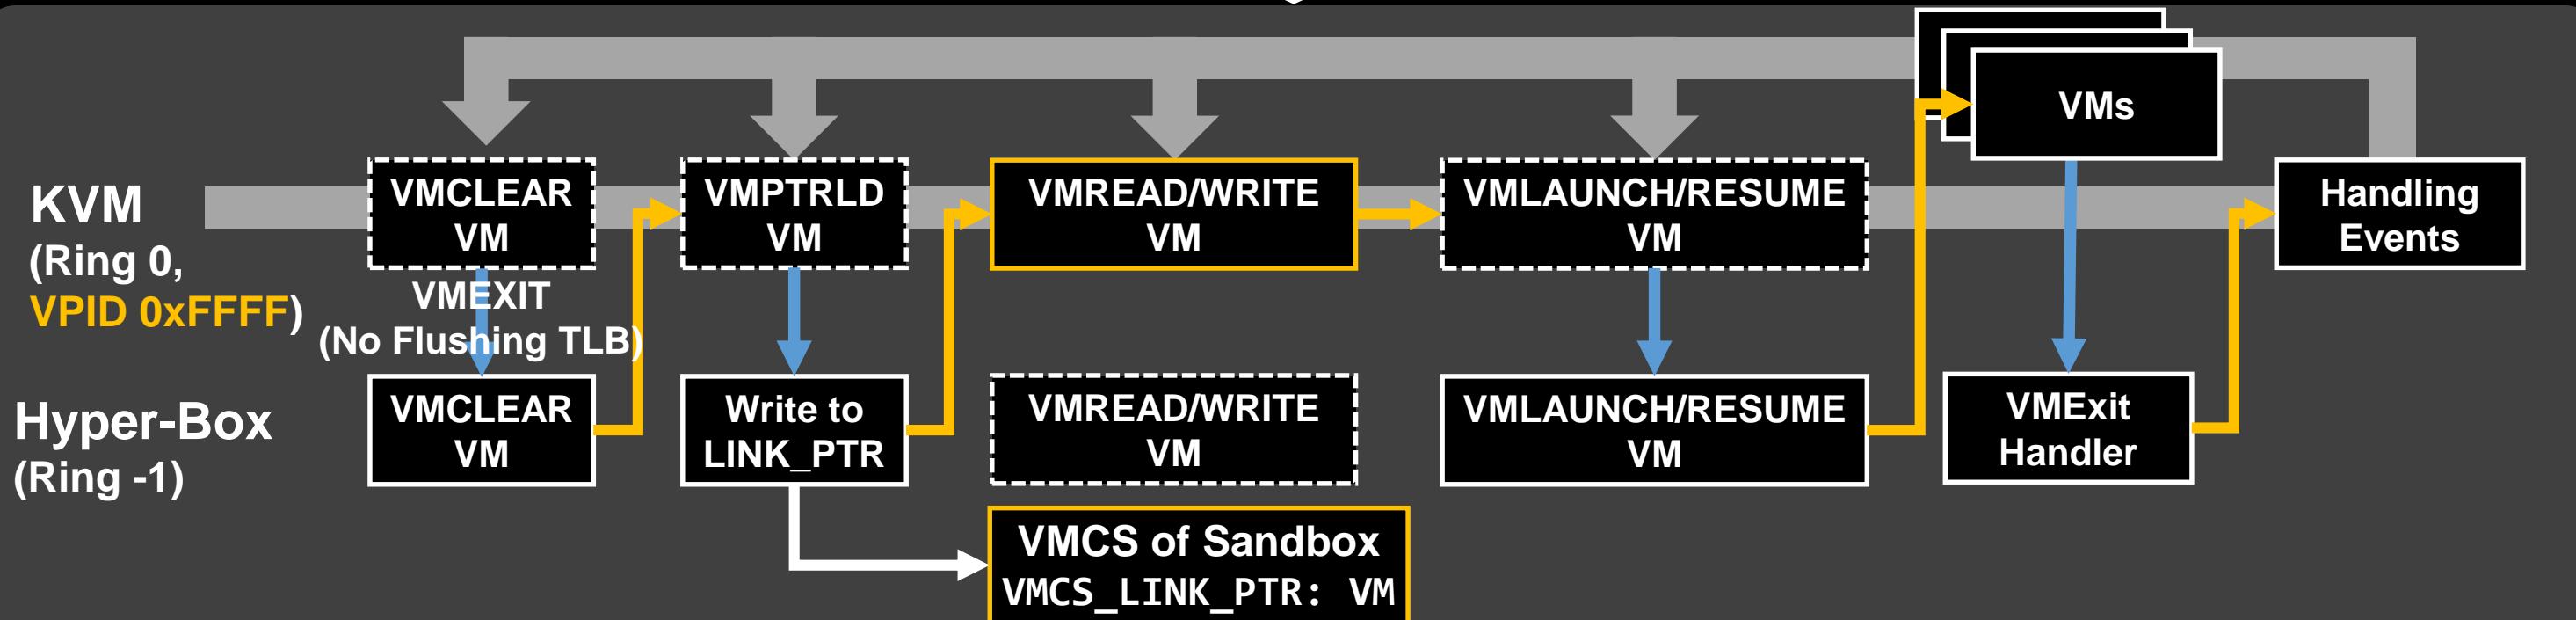
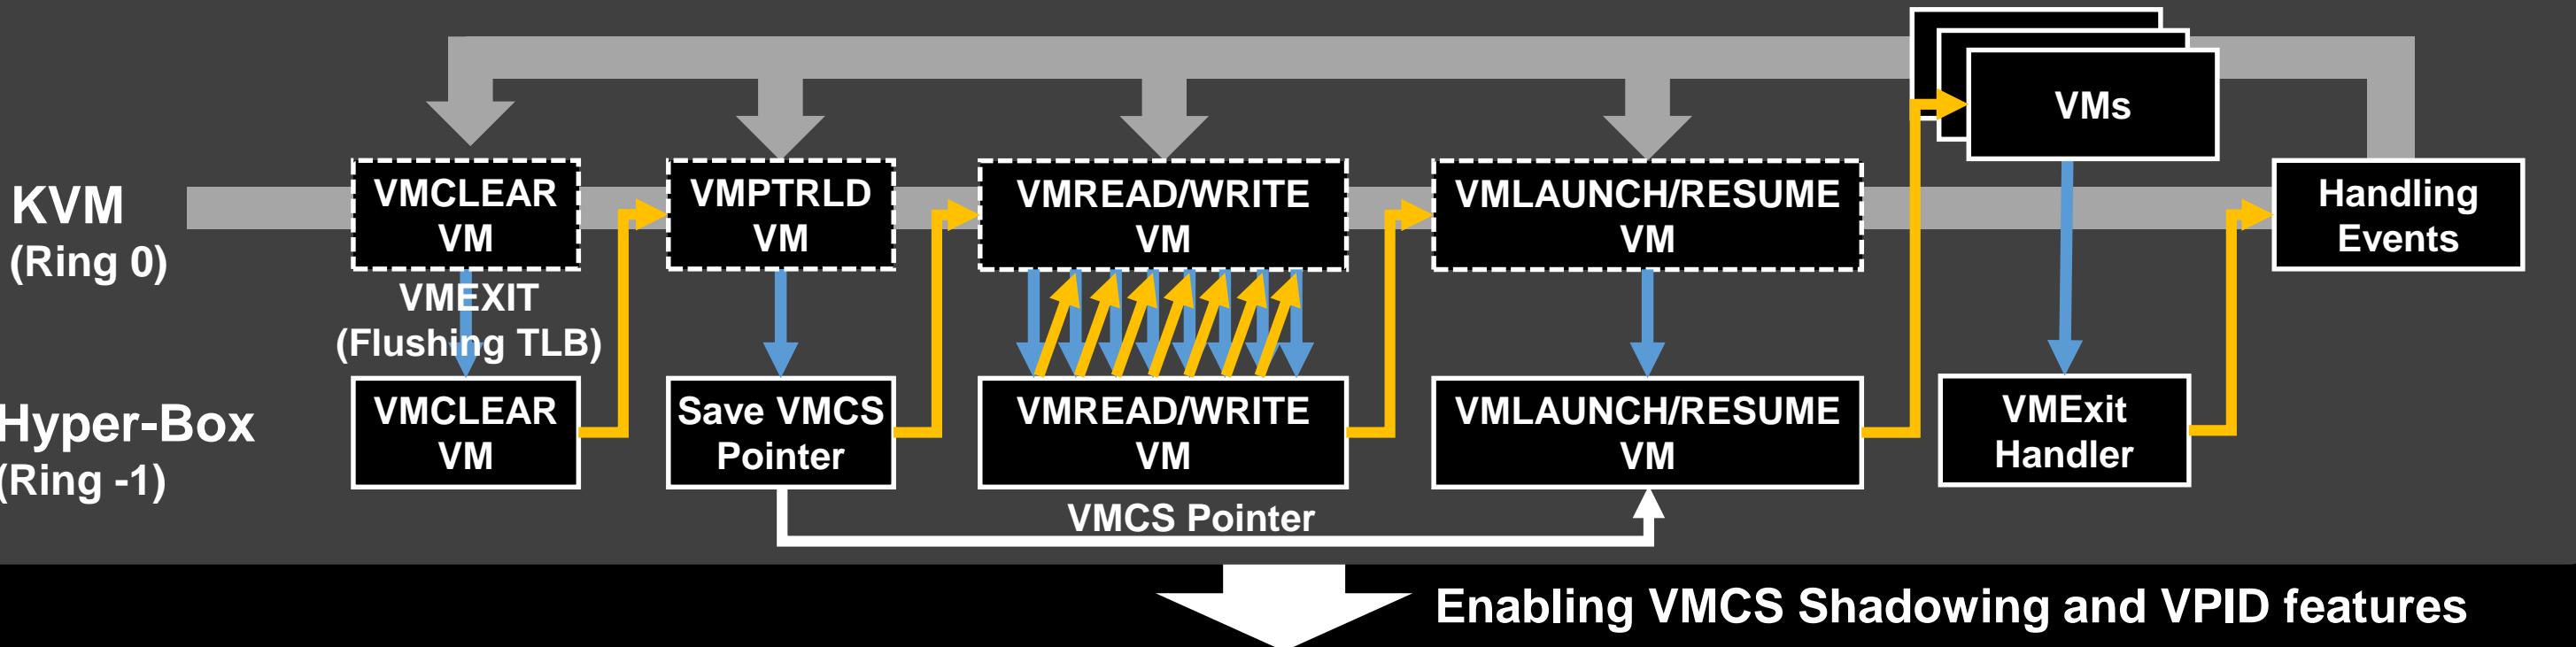
Hyper-Box Design

Emulating VMX instructions – Reducing Overhead (1)



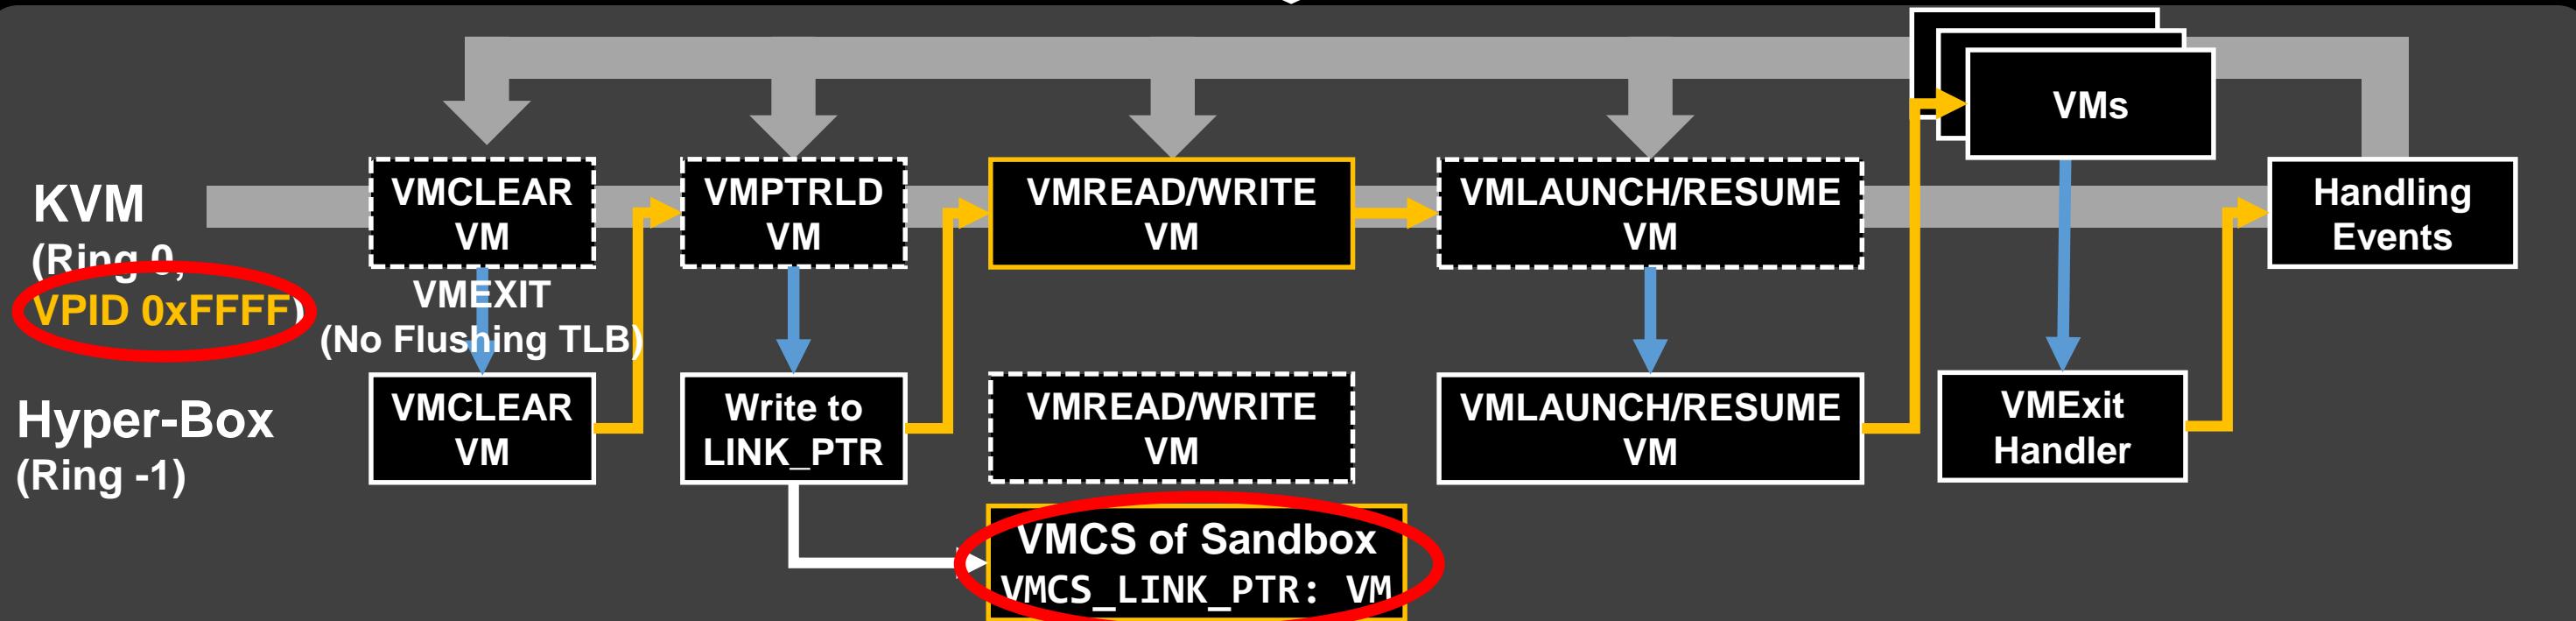
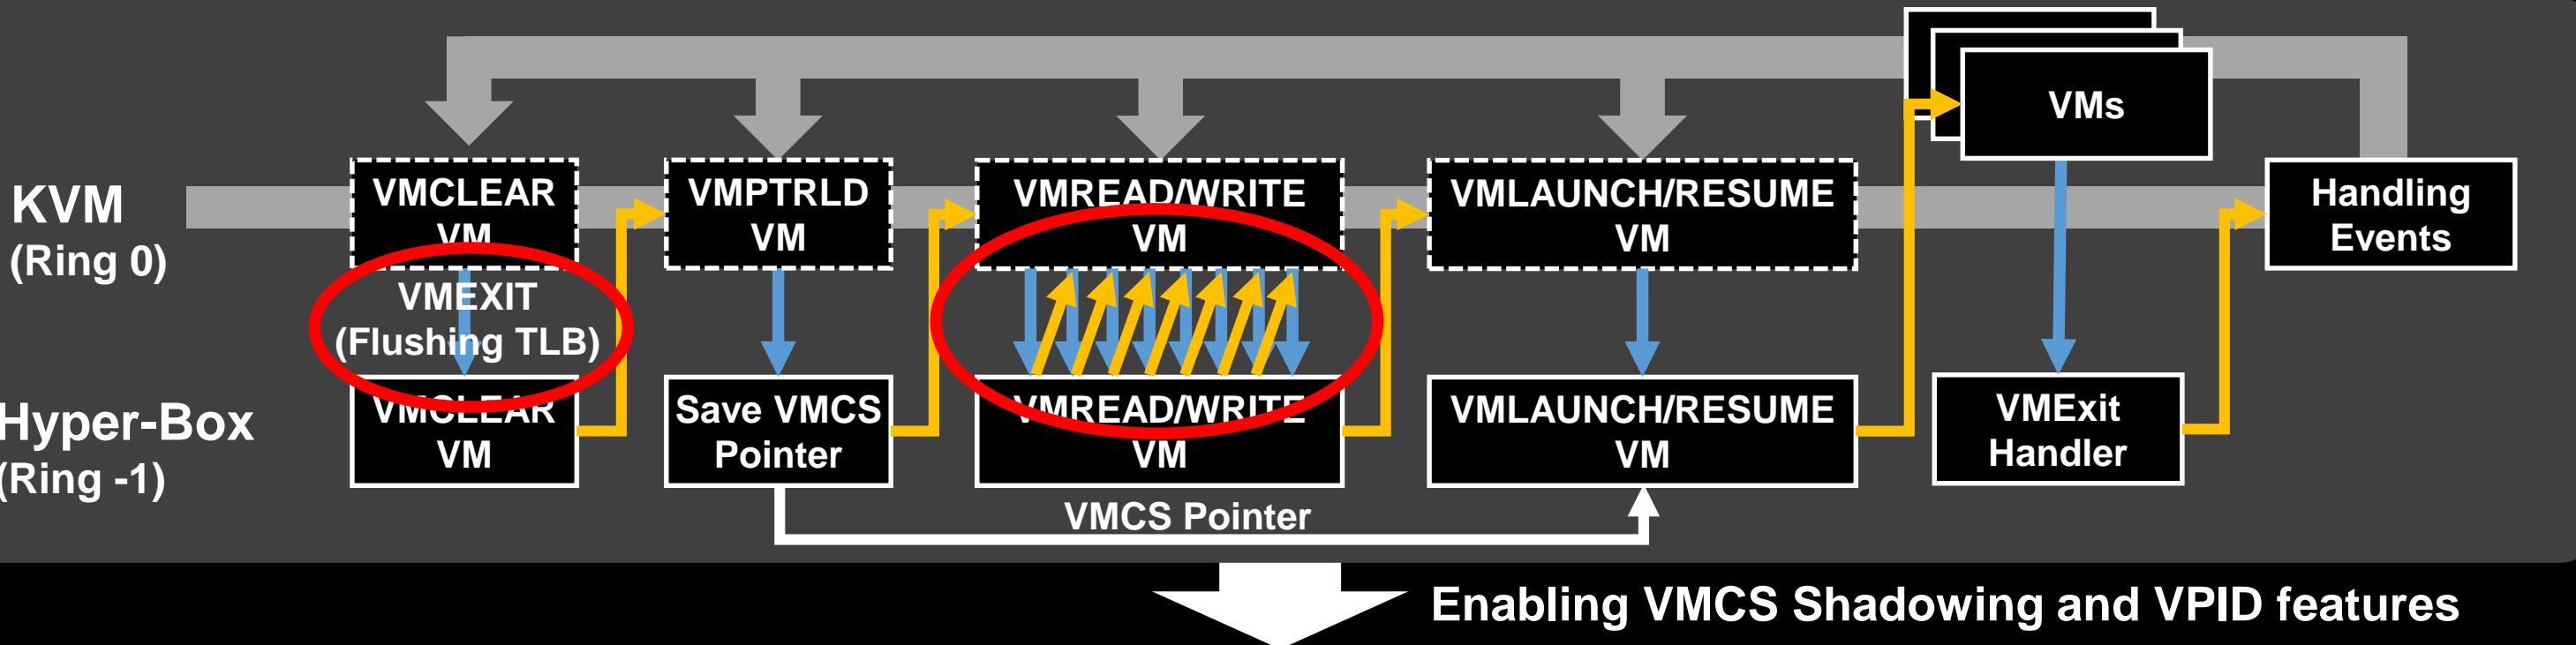
Hyper-Box Design

Emulating VMX instructions – Reducing Overhead (2)



Hyper-Box Design

Emulating VMX instructions – Reducing Overhead (3)



Implementation of Alcatraz

- Hyper-box is a loadable kernel module
 - So, It can be loaded any time BEFORE you start VMs
 - If not, it shows some errors related to the VMXON fail
 - It was made with C and assembly code
- Tailored Linux kernel is just recompiled version
 - Some options of a kernel config file were REMOVED
 - CONFIG_IA32_EMULATION, CONFIG_COMPAT, CONFIG_COMPAT_32, CONFIG_X86_X32 → System call interfaces
 - CONFIG_JUMP_LABEL → Runtime code modifications
 - No more system call interface for legacy!

- Background
- Analysis of Critical Paths of Escapes
- Design and Implementation of Alcatraz
- Evaluation and Demo 
- Conclusion and Black Hat Sound Bytes

Benchmark Machine

- Host Machine (Server)

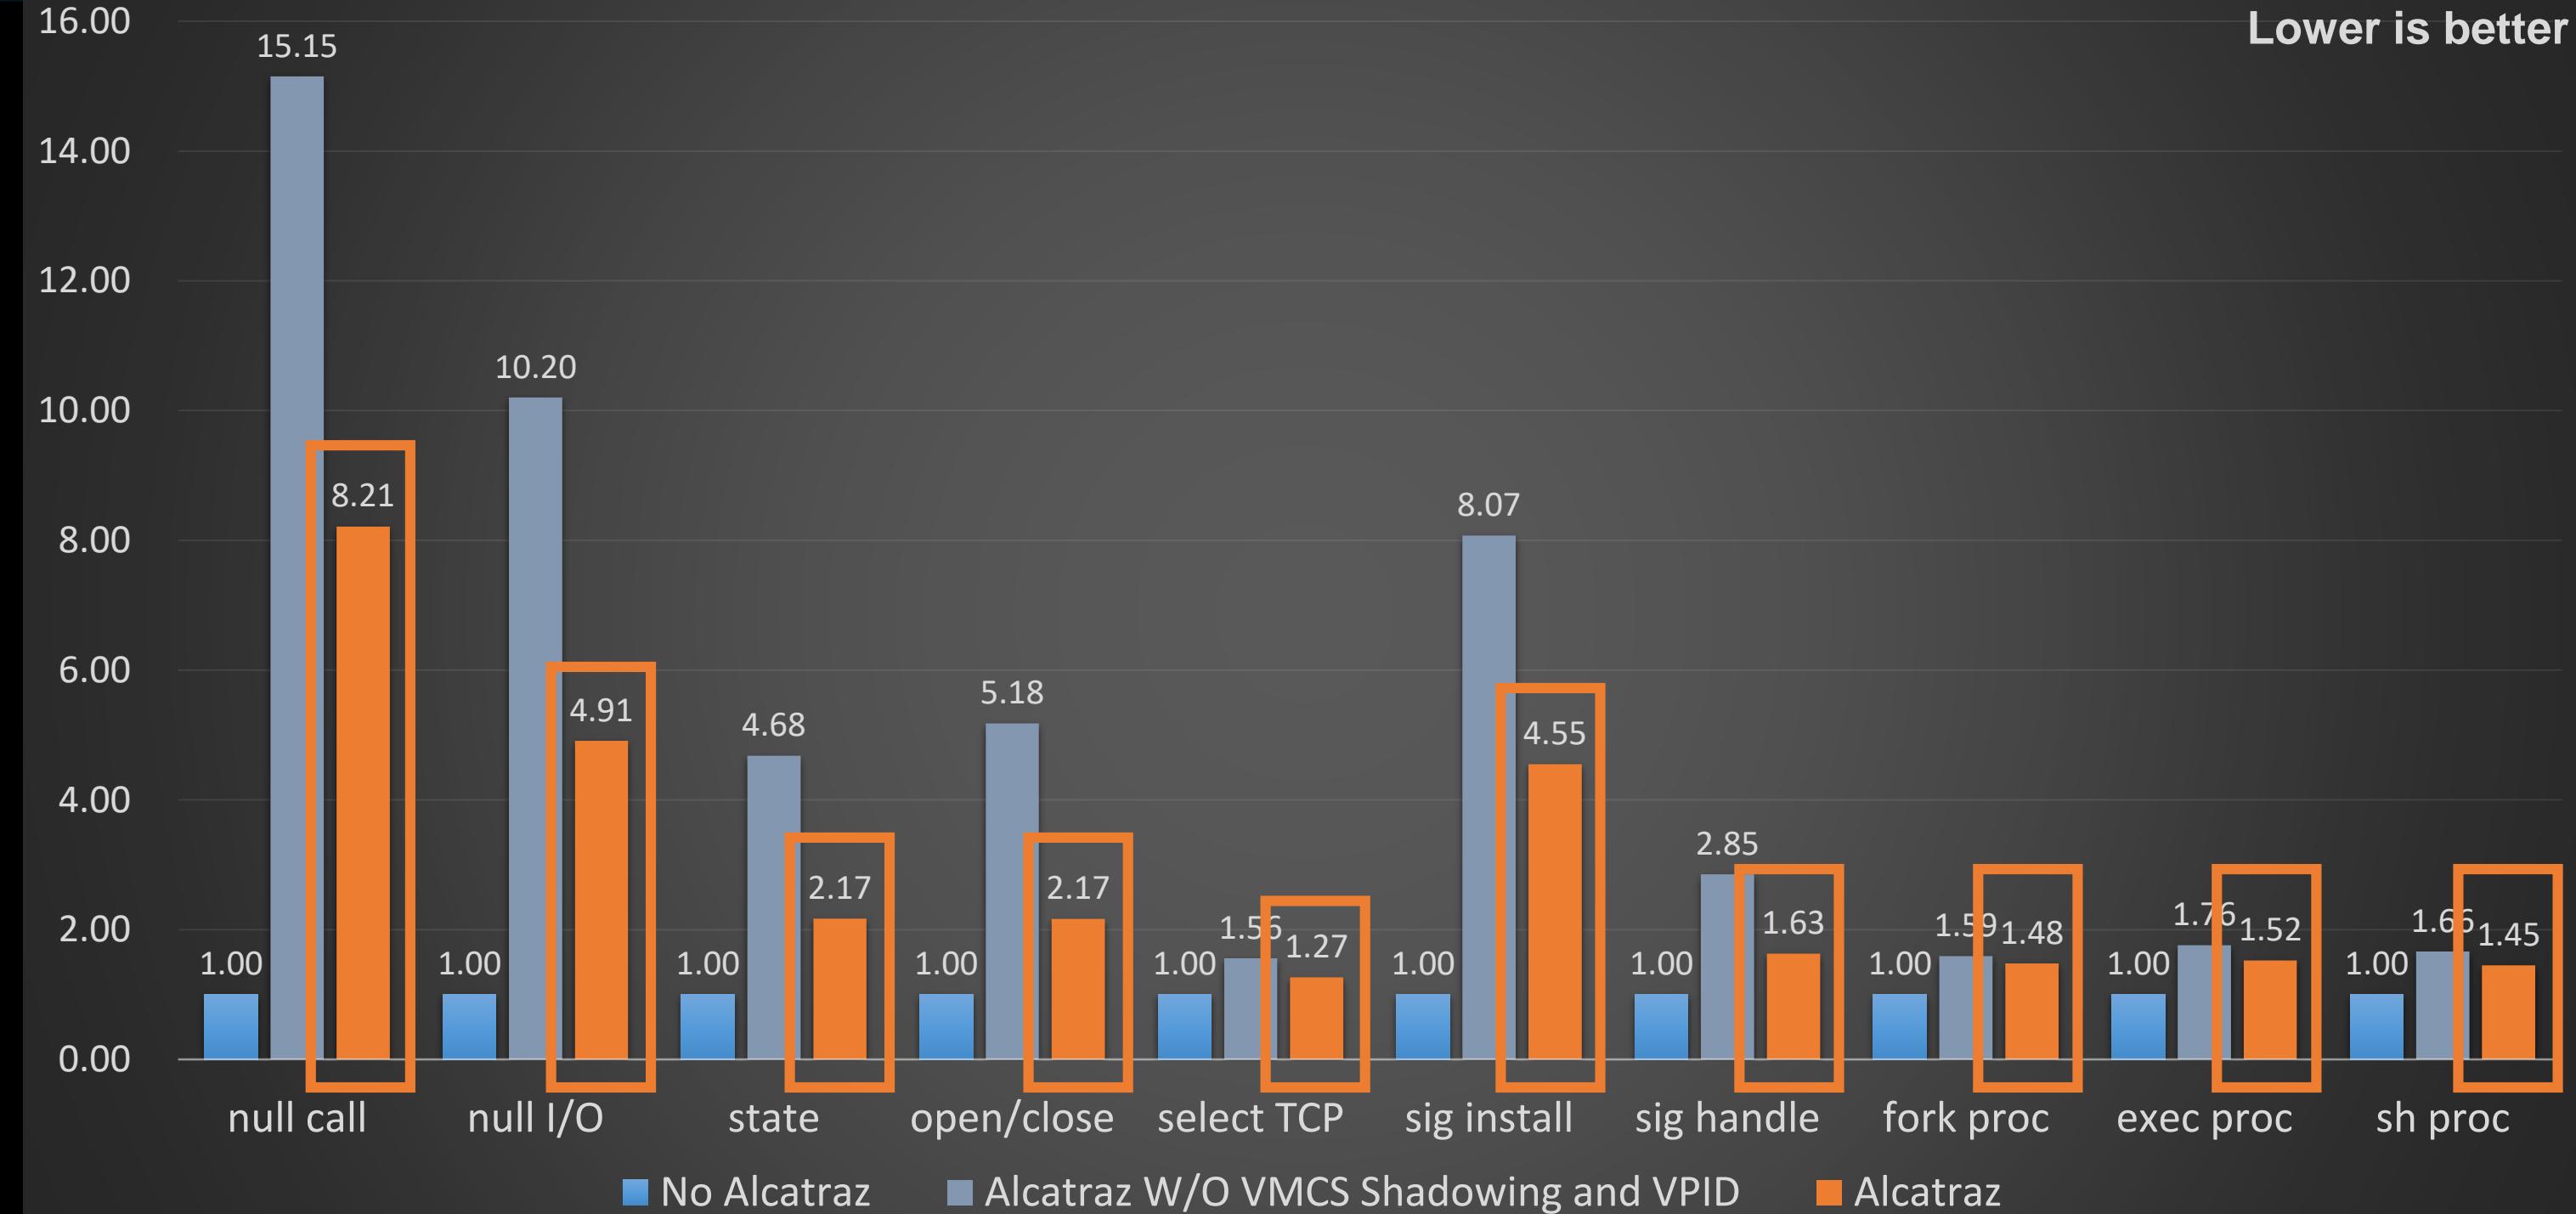
- CPU: Intel Xeon Gold 6242R 2 CPU (40 cores, 80 thread)
- RAM: 512GB
- HDD: NVME 1TB
- OS: Ubuntu 20.04.01
- Kernel: 5.8.0-44-generic Linux kernel

- Guest Machine (QEMU)

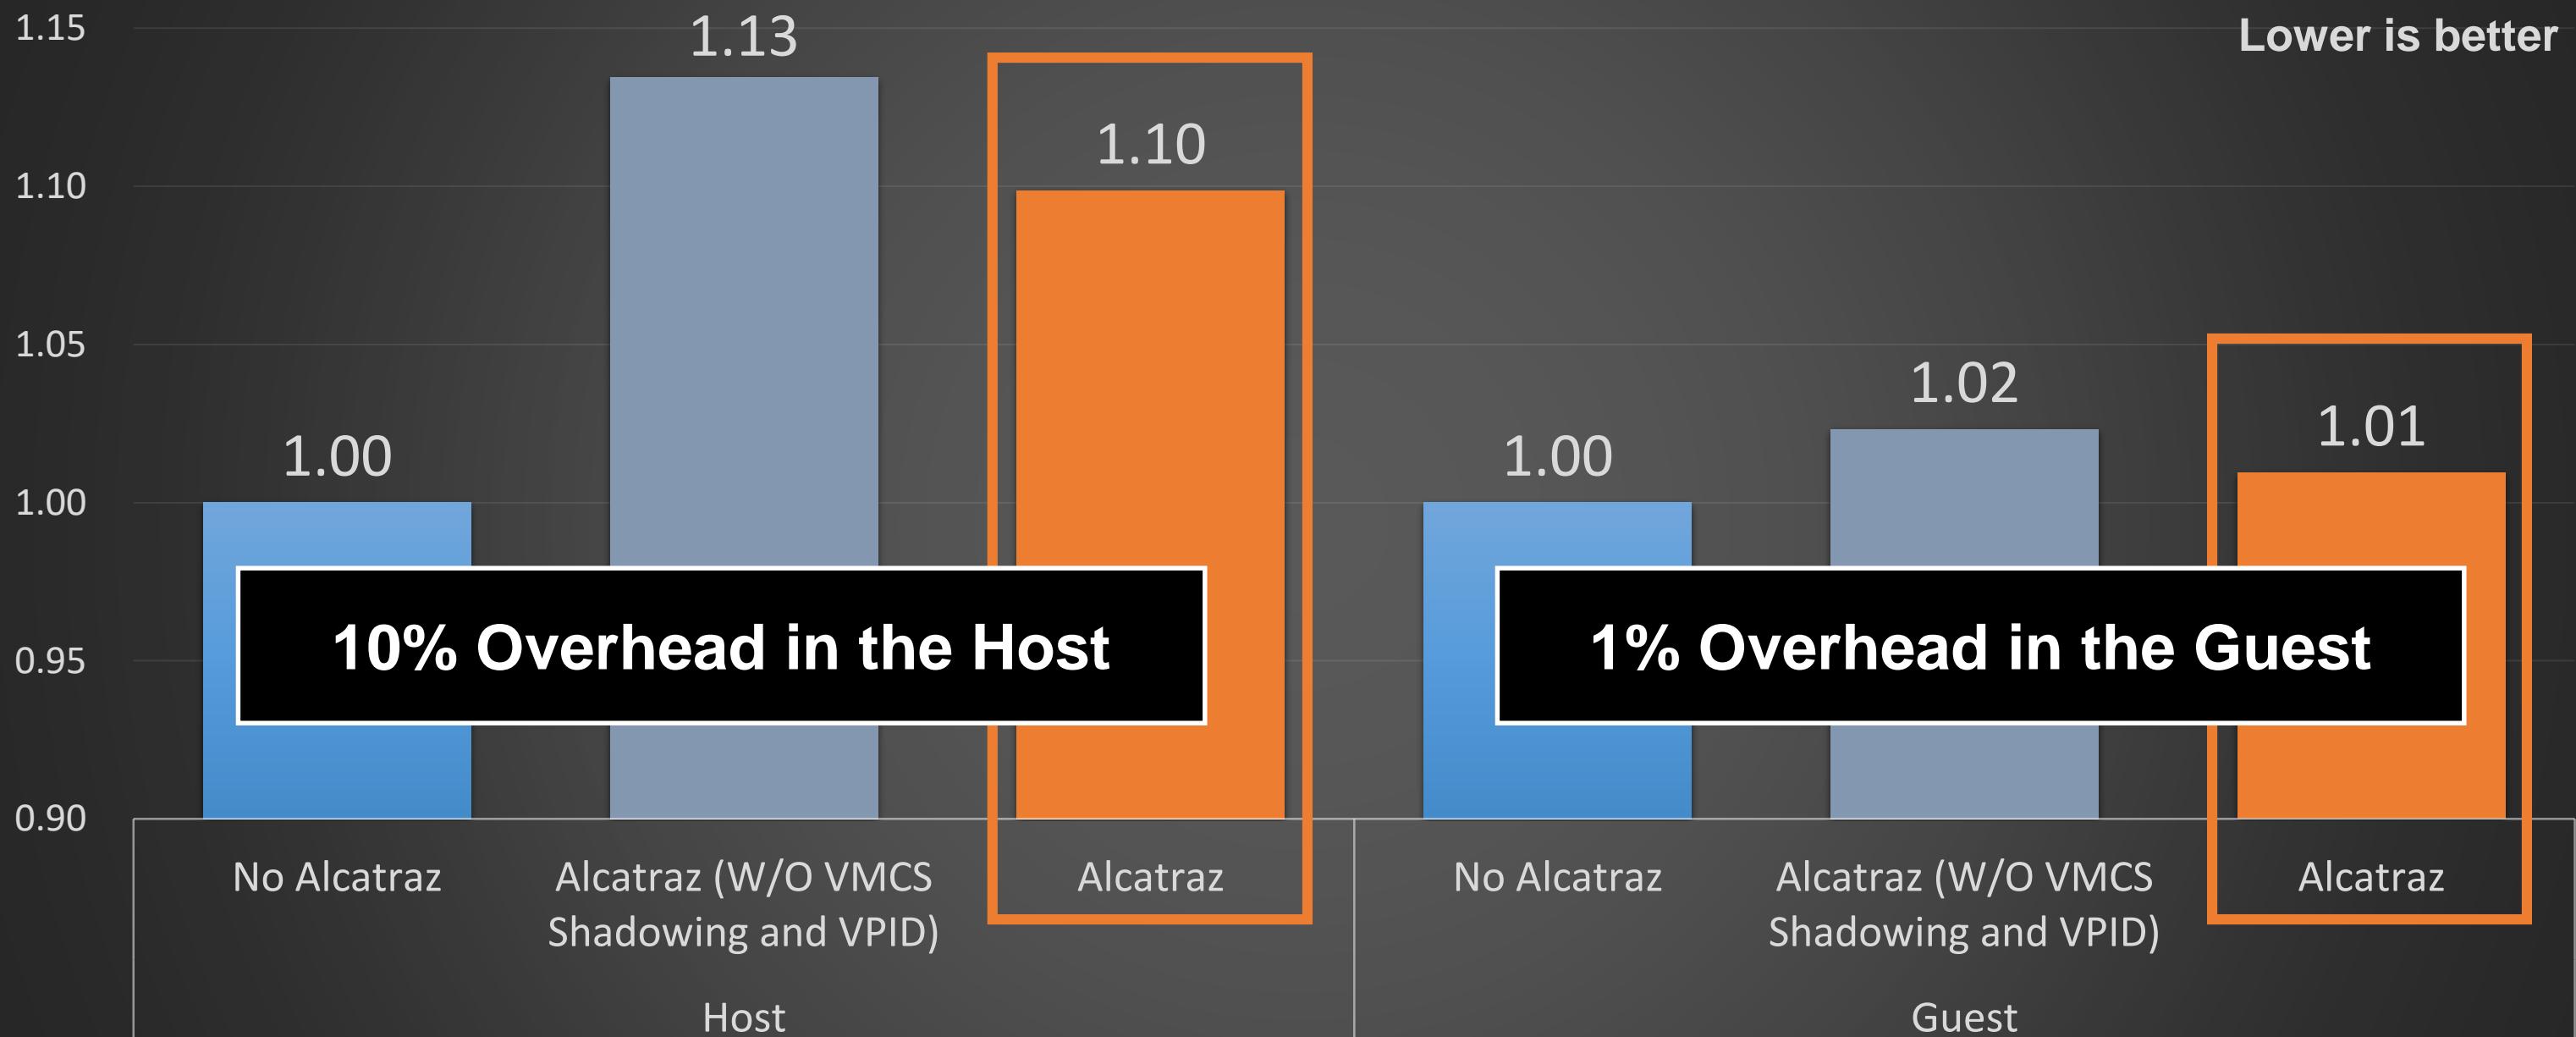
- CPU: 4 cores
- RAM: 16GB
- HDD: 60GB
- OS: Ubuntu 20.04.01
- Kernel: 5.8.0-44-generic Linux kernel

Evaluation – Imbench

Lower is better



Evaluation – Kernel Compilation



Please check the project link: <https://github.com/kkamagui/alcatraz>

Conclusion and Black Hat Sound Bytes

- **Containers, VMs, and microVMs suffer from escapes**
 - The hypervisor can provide strong isolation
 - But, it has excessive privilege and connects with VMs tightly
- **Taking a higher ground is not always the answer**
 - Researchers usually try to solve the problems with the higher privilege
 - The higher privilege we get, the more constraints we have!
- **Alcatraz is a practical hypervisor sandbox to prevent escapes**
 - It downgrades KVM's privilege and makes a sandbox with Intel VT
 - It has approximately 10% overhead in the host and 1% overhead in the guest (kernel compilation benchmark)

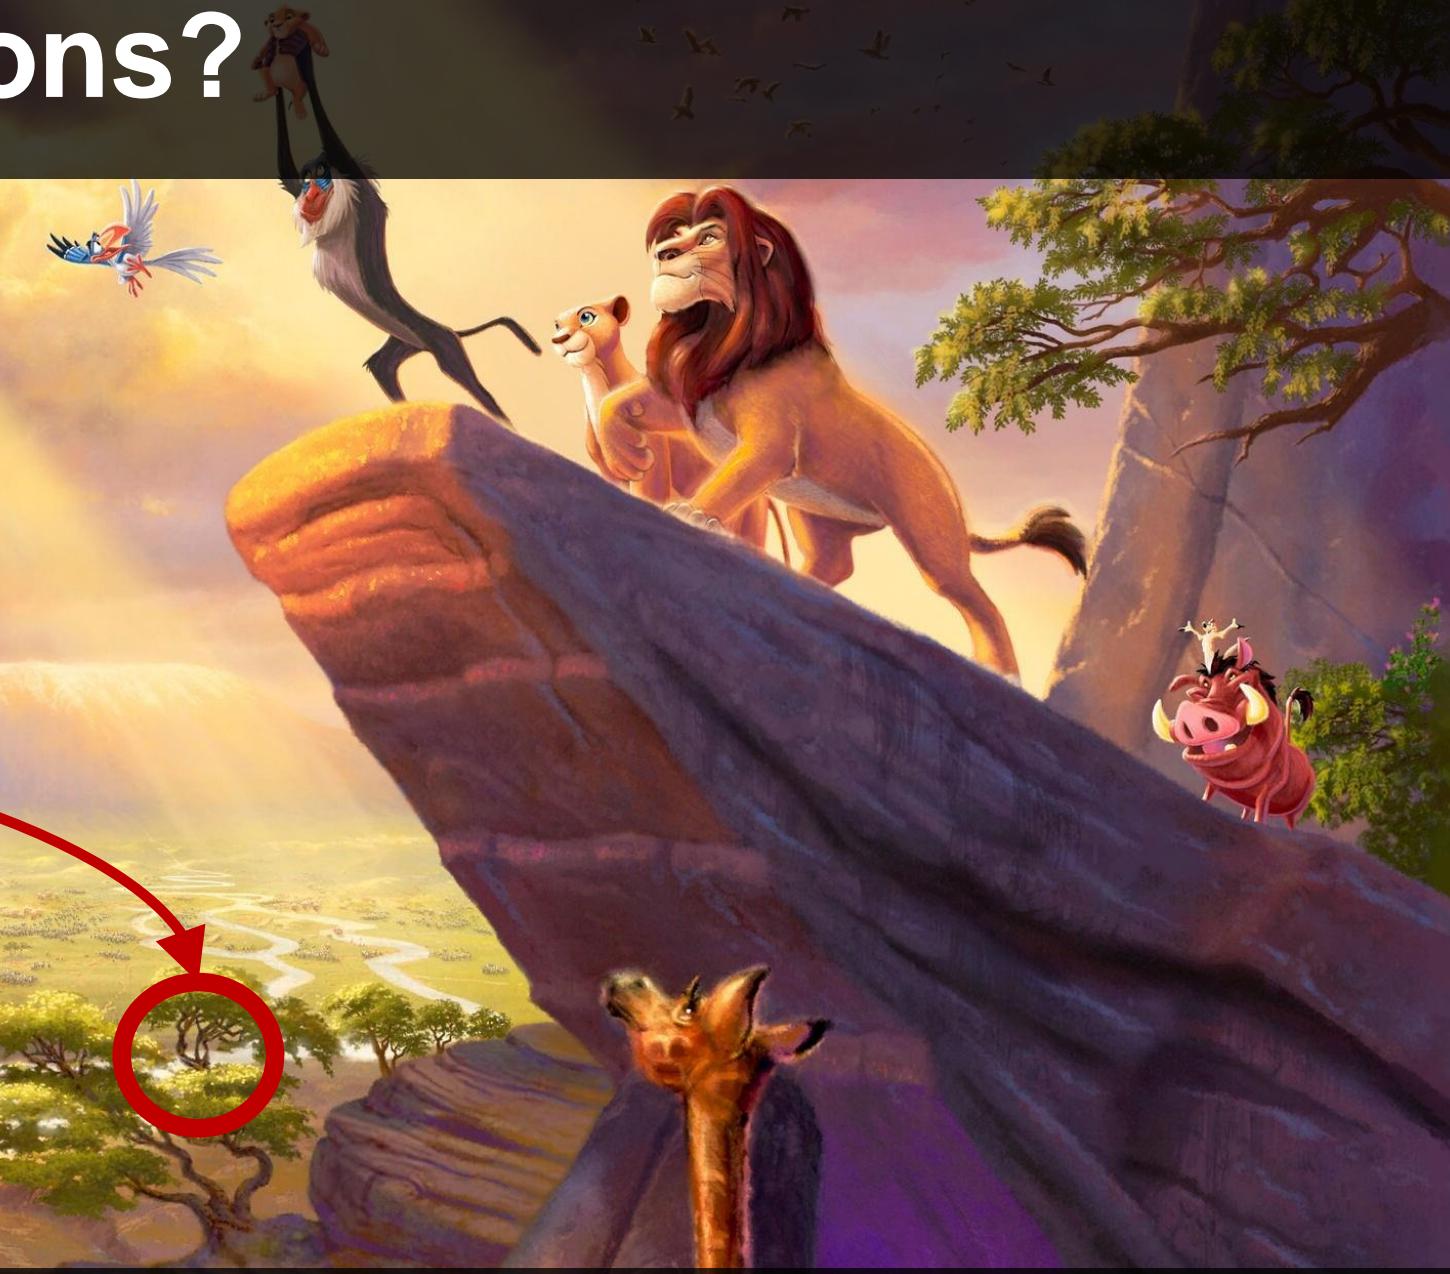
Questions?



**SENIOR
RESEARCHER!**



**BLACK HAT USA
SPEAKER!**



Project : <https://github.com/kkamagui/alcatraz>
Contact: hanseunghun@nsr.re.kr, @kkamagui1

Reference

- Seunghun, H., Junghwan, K., Wook, S., HyoungChun K., and Eungki, P. *Myth and truth about hypervisor-based kernel protector: The reason why you need shadow-box*. Black Hat Asia. 2017.
- Seunghun, H., Jun-Hyeok, P., Wook, S., Junghwan, K., and HyoungChun K. *Shadow-box v2: The practical and omnipotent sandbox for ARM*. Black Hat Asia. 2018.
- Randazzo, A., and Ilenia T. *Kata containers: An emerging architecture for enabling mec services in fast and secure way*. IEEE IOTSMS. 2019.
- Agache, A., et al. *Firecracker: Lightweight virtualization for serverless applications*. NSDI. 2020.
- Young, E., et al. *The true cost of containing: A gVisor case study*. HotCloud. 2019.
- Azab, A., et al. *Hypersentry: enabling stealthy in-context measurement of hypervisor integrity*. ACM CCS. 2010.
- Rutkowska, J., and Rafał W. *Preventing and detecting Xen hypervisor subversions*. Black Hat USA. 2008
- Bulygin, Y., and David S. *Chipset based approach to detect virtualization malware*. Black Hat USA. 2008
- Li, S., John S., and Jason N. *Protecting cloud virtual machines from hypervisor and host operating system exploits*. USENIX Security. 2019.
- Zhijian S., Jian W., and Yue Z. *3d red pill: A guest-to-host escape on QEMU/KVM virtio Devices*. Black Hat Asia. 2020.
- Elhage, N. *Virtunoid: A kvm guest->host privilege escalation exploit*. Black Hat USA. 2011.
- Understanding the VENOM Vulnerability, <https://blog.trendmicro.com/understanding-the-venom-vulnerability/>